

E-Commerce Site (An Online Shopping Site)

Name: Naga Venkata Kanakalakshmi

UNT ID: 11725119

DPDB Project 1

Requirement-1:

Introduction:

In the current world, online shopping has become an integral part of everyone's life. The E site plays the vital role in delivering certified products in the shortest possible time. E-commerce sites provide customers with seamless shopping experience and give them a comfortable lifestyle. It helps the customers to browse the wide range of products by applying different filters, which makes browsing easy. The payment gateway helps the customers to make the safe and Security transactions.

Requirements:

- The goal of the project is to create an e-commerce platform with a strong database schema. It supports a variety of functionalities, including customer signup, order placement, address management, managing products, shopping cart, handling transactions, supplier information, the option to review and the rate, exciting discounts and gift vouchers, an excellent delivery system, and high security using error handling mechanism.
- Overall to implement the e-commerce platform used 17 tables. Which is customer, products, product_category, address, cart, orders, transaction_summary, supplier, reviews, gift_vouchers, delivery_partner, customer_address, customer_gift_voucher, customer_delivery_partner, products_cart, order_products, products_suppliers. These tables include relation tables too. Which are the intermediary tables, links between the two entities.

Assumptions:

- A customer can add multiple products to his/her cart as per choice. Whereas in the same way the same product can be added by multiple customers in the cart. It associates the many-to-many to relation between the products and the cart. As the relation is many-to-many an intermediary table formed to link the products and cart table that is products_cart.
- As mentioned in point (1), in the same way a product can be ordered by many customers and during the product's delivery there can be the same product among the multiple orders. By this the products and orders table form the many-to-many relation. As it's relationship is many-to-many, an intermediary relation table is formed to link the orders and products table that is products_orders.

- There will be multiple suppliers who supply the same product and there are few suppliers where they can supply multiple products. For instance, there will multiple suppliers who supply washing machines like Samsung, Ig, ifb etc. In the same way a supplier can supply multiple products like fridge, washing machine etc. By this the relation associated with the supplier and the product is many-to-many and there is intermediary table to form the link between them that is the products_suppliers
- Multiple products can be placed in the location, and they have a specific address to it. By this it forms a many-to-one relation between the products and the address. When the group of products are stored in a particular location (warehouse) it helps the delivery partner to pick and deliver the products to the customer.
- Group of products have a specific category, in most of the online shopping sites to make searching easier the products are categorized. For instance, dresses, tops, sarees come under clothing section, electronic gadgets like mobiles, power banks have the specific category. So, the relationship between products and the product_category is many-to-one.
- Each product has multiple reviews, which helps customers to purchase the product more easily. The relationship between the products and reviews is many-to-one which says that a product can have zero reviews, one or more reviews.
- Each customer has at most one cart to add the items to purchase. It associates the one-to-one relation between the cart and the customer.
- A customer can store multiple addresses, where he can order the products to multiple locations as per his/her choice. In the idle case, multiple customers can store the same address as when multiple as working in same company they want to order the products to their office location then the group of customers have the same address. The relationship associated between the customer and the address is many-to-many. There is an intermediary table which links customers and address table, which can also store the default address of the customer.
- Each customer can have multiple gift vouchers to redeem, and a single gift voucher can be given to multiple customers with the same coupon code. Many-to-Many relation forms between the customers and the gift vouchers.
- The site provides the leverage to place any number of orders by the customer. This forms the many-to-one relation between the orders and the customer.
- Delivery Partner can deliver any number of products to many customers, it makes the work faster and delivery of the products on time to the customers. A customer may contact multiple delivery partners for multiple orders. Thereby it forms the many-to-many relation between the customer and delivery partner. As it's a many-to-many relationship it forms the intermediary table customer_delivery_partner.
- Each order stores the respective transaction details, this ensures that every order placed within the system is associated with a single recorder which summarizes the transaction details of the specific customer. It helps to easily track the customer transactions. This ensures the one-to-one relation between the orders and the transaction summary.

- Many orders can be placed to a single address, and single order cannot be diverse among the multiple addresses. which ensures a many-to-one relation between the orders and the address.
- A single order is delivered by the specific delivery partner. This ensures the relationship between the orders and the delivery partner is one-to-one.
- On a single order, the customers can provide multiple reviews as per their choice. This forms the many-to-one relationship between the reviews and the order.
- Each supplier has exactly one address to it and each location is dedicated to a supplier. This ensures to form the one-to-one relationship between the suppliers and the address.
- A customer can directly place the order, without adding to cart. Sometimes a customer will add multiple products to cart but he/she won't buy it from cart, he/she can directly place order without using the cart.

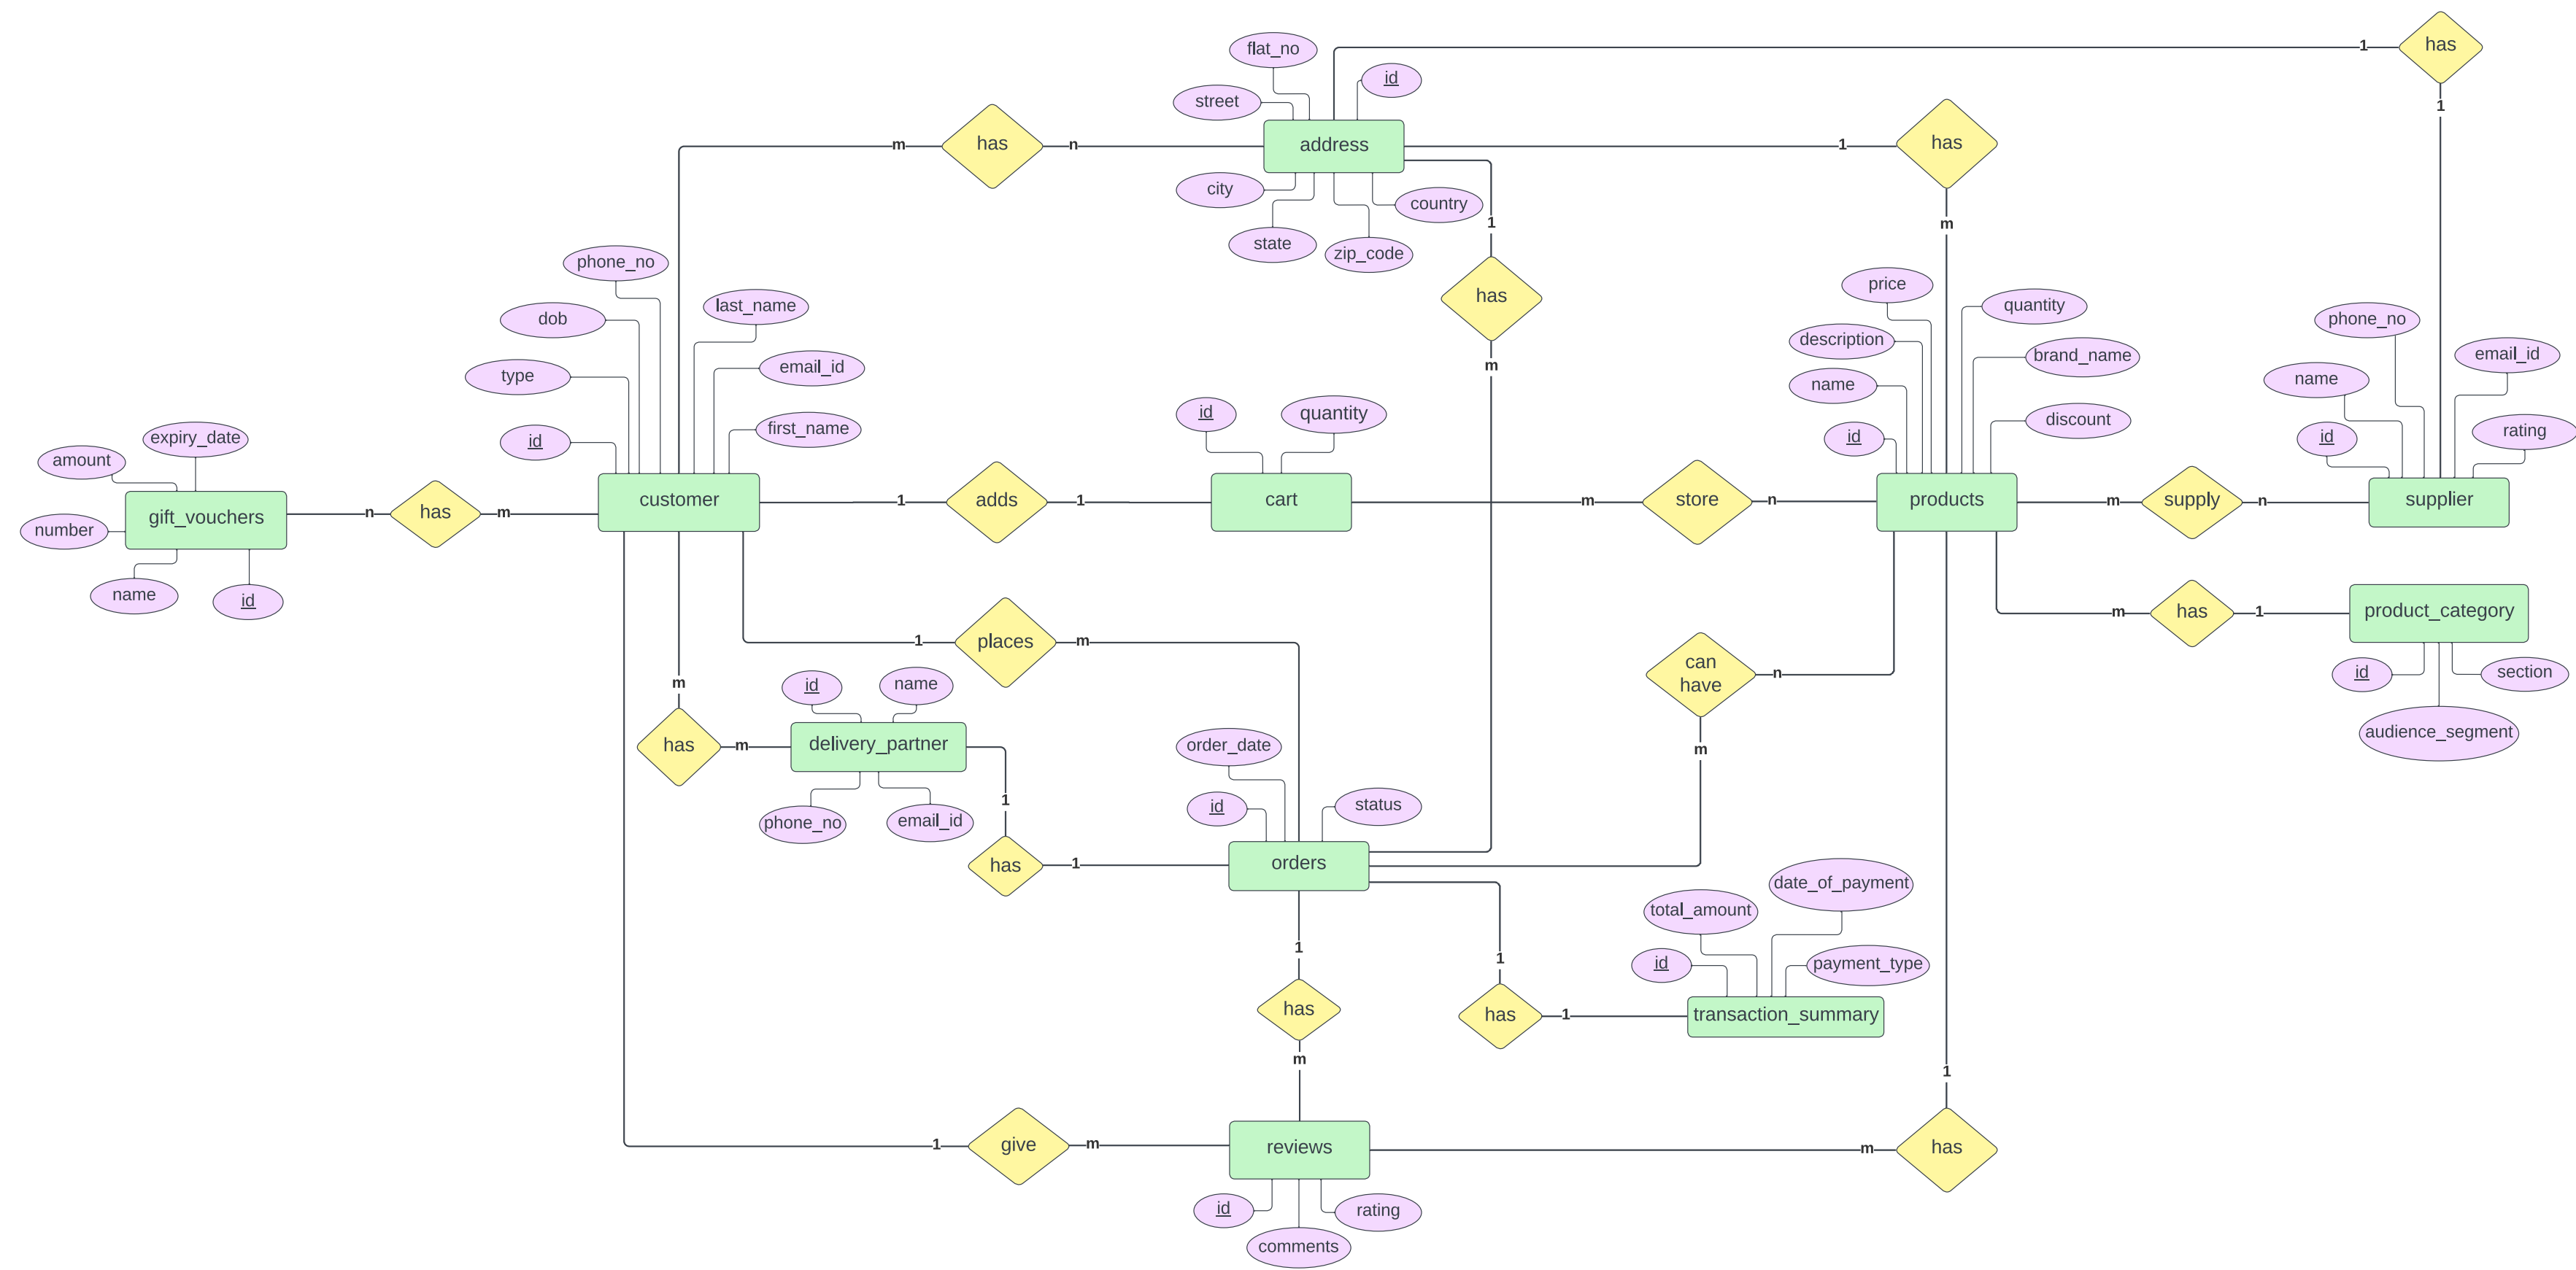
Analysis					
Table	Attributes	Domain Constraints	Check Contraints	Relation	Grant_Access (Customer/Administrator/SuperUser)
products	id	Raw(16)	PRIMARY KEY	products has many-to-many relation with cart product has many-to-many relation with orders products has many-to-many relation with suppliers product has many-to-one relation with address product has many-to-one with product_category product has one-to-many with review	Customer - SELECT Adminstrator - SELECT, INSERT, UPDATE, DELETE, CREATE, ALTER, DROP, EXECUTE Super User - ALL PRIVILEGES
	name	VARCHAR2(50)	dont allow nulls		
	description	VARCHAR2(200)	dont allow nulls		
	price	INTEGER	dont allow nulls		
	quantity	INTEGER	dont allow nulls		
	discount	INTEGER	NA		
	brandname	VARCHAR2(20)	dont allow nulls		
	address_id	Raw(16)	no null values, FOREIGN KEY (address_id) REFERENCES address(id)		
	product_category_id	Raw(16)	dont allow nulls, FOREIGN KEY (product_category_id) REFERENCES product_category(id)		
customer	id	Raw(16)	PRIMARY KEY	customer has one-to-one relation with cart customer has many-to-many relation with address customer has many-to-many relation with giftvoucher customer has one-to-many relation with orders customer has one-to-many relation with reviews customer has many-to-many relation with delivery_partner	Customer - SELECT, INSERT,UPDATE Adminstrator - SELECT, INSERT, UPDATE, DELETE, CREATE, ALTER, DROP, EXECUTE Super User - ALL PRIVILEGES
	first_name	VARCHAR2(20)	dont allow nulls		
	last_name	VARCHAR2(20)	dont allow nulls		
	phone_no	VARCHAR2(20)	dont allow nulls, UNIQUE, CHECK (REGEXP_LIKE (phone, '^ \d{3}\ \d{3}-\d {4}\$'))		
	email_id	VARCHAR2(255)	no null values, UNIQUE, CHECK (REGEXP_LIKE (email, "[A-Za-z0-9. %+-]+@[A-Za-z0-9.-]+\.[A-Za-z]{2,}\$"))		
	dob	DATE	dont allow nulls		
	type - regular/premium/gold/vip	VARCHAR2(20)	dont allow nulls		
cart	id	Raw(16)	PRIMARY KEY	cart has one-to-one relation with customer cart has many-to-many relation with products	Customer - SELECT, INSERT, UPDATE, DELETE Administrator - SELECT, UPDATE, DELETE Super User - ALL PRIVILEGES
	customer_id	Raw(16)	dont allow nulls, FOREIGN KEY (customer_id) REFERENCES customer(id)		
	quantity	INTEGER	dont allow nulls		
orders	id	Raw(16)	PRIMARY KEY	orders has many-to-one relation with customer orders has many-to-many relation with products order has one-to-one relation with transaction_summary orders has many-to-one with address order has one-to-many with review order has one-to-one delivery_partner	Customer - SELECT, INSERT, UPDATE, DELETE Administrator - SELECT, INSERT, UPDATE, DELETE, ALTER, DROP Super User - ALL PRIVILEGES
	status	VARCHAR2(20)	dont allow nulls		
	order_date	DATE	dont allow nulls		
	address_id	Raw(16)	dont allow nulls, FOREIGN KEY (address_id) REFERENCES address(id)		
	customer_id	Raw(16)	dont allow nulls, FOREIGN KEY (customer_id) REFERENCES customer(id)		
transaction_summary	id	Raw(16)	PRIMARY KEY	transaction_summary has one-to-one relation with orders	Customer - SELECT Administrator - SELECT Super User - ALL PRIVILEGES
	total_amount_paid	INTEGER	dont allow nulls		
	payment_type - upi/wallets/creditcard/debitcard/net banking/emi/cod	VARCHAR2(20)	dont allow nulls		
	date_of_payment	DATE	dont allow nulls		

Analysis					
Table	Attributes	Domain Constraints	Check Constraints	Relation	Grant_Access (Customer/Administrator/SuperUser)
	order_id	Raw(16)	dont allow nulls FOREIGN KEY (order_id) REFERENCES orders(id)		
supplier	id	Raw(16)	PRIMARY KEY	suppliers has many-to-many relation with products supplier has one-to-one relation with address	Customer - NA Administrator - SELECT, INSERT, UPDATE, DELETE Super User - ALL PRIVILEGES
	name	VARCHAR2(50)	dont allow nulls		
	phone_no	VARCHAR2(20)	dont allow nulls, UNIQUE, CHECK (REGEXP_LIKE (phone, '^(\d{3})\d{3}-\d{4}\$'))		
	email	VARCHAR2(255)	dont allow nulls, UNIQUE, CHECK (REGEXP_LIKE (email, '^([A-Za-z0-9._%+-]+@[A-Za-z0-9.-]+\.[A-Za-z]{2,}\$'))		
	rating	INTEGER	NA		
	address_id	Raw(16)	dont allow nulls, FOREIGN KEY (address_id) REFERENCES address(id)		
address	id	Raw(16)	PRIMARY KEY	address has many-to-many relation with customer address has one-to-many relation with product address has one-to-many with orders address has one-to-one relation with supplier	Customer - SELECT, INSERT, UPDATE, DELETE Administrator - SELECT, INSERT, UPDATE, DELETE Super User - ALL PRIVILEGES
	flat_no	NUMBER	dont allow nulls		
	street	VARCHAR2(50)	dont allow nulls		
	city	VARCHAR2(50)	dont allow nulls		
	state	VARCHAR2(50)	dont allow nulls		
	country	VARCHAR2(50)	dont allow nulls		
	zip_code	VARCHAR(10)	dont allow nulls		
product_category	id	Raw(16)	PRIMARY KEY	product_category has one-to-many with product	Customer - SELECT Administrator - SELECT, INSERT, UPDATE, DELETE Super User - ALL PRIVILEGES
	section - clothing/ footwear/ accessories /electronics	VARCHAR2(20)	dont allow nulls		
	audience_segment - men/women/kids/girls/boys	VARCHAR2(20)	dont allow nulls		
reviews	id	Raw(16)	PRIMARY KEY	reviews has many-to-one relation with customer review has many-to-one with product review has many-to-one with order	Customer - SELECT, INSERT, UPDATE, DELETE Administrator - SELECT, DELETE Super User - ALL PRIVILEGES
	rating	NUMBER	NA		
	comments	VARCHAR2(150)	NA		
	product_id	Raw(16)	FOREIGN KEY (product_id) REFERENCES products(id)		
	order_id	Raw(16)	FOREIGN KEY (order_id) REFERENCES orders(id)		
	customer_id	Raw(16)	FOREIGN KEY (customer_id) REFERENCES customer(id)		
gift_vouchers	id	Raw(16)	PRIMARY KEY	gift_voucher has many-to-many relation with customer	Customer - SELECT Administrator - SELECT, INSERT, UPDATE, DELETE Super User - ALL PRIVILEGES
	name	VARCHAR2(50)	NA		
	number	NUMBER	NA		
	amount	INTEGER	NA		
	expiry_date	DATE	NA		
	id	Raw(16)	PRIMARY KEY		
	name	VARCHAR2(50)	dont allow nulls		

Analysis					
Table	Attributes	Domain Constraints	Check Constraints	Relation	Grant_Access (Customer/Administrator/SuperUser)
delivery_partner	phone_no	VARCHAR2(20)	dont allow nulls, CHECK (REGEXP_LIKE (phone, '^ \d{3}\d{3}-\d {4}\$'))	delivery_partner has many-to-many relation with customer delivery_partner has one-to-one orders	Customer - NA Administrator - SELECT, INSERT, UPDATE, DELETE Super User - ALL PRIVILEGES
	email	VARCHAR2(255)	dont allow nulls, CHECK (REGEXP_LIKE (email, '^ [A-Za-z0-9. _%+ -] +@[A-Za-z0-9. -] + \. [A-Za-z] {2,}\$'))		
	order_id	Raw(16)	dont allow nulls FOREIGN KEY (order_id) REFERENCES orders(id)		
customer_address (relation table)	customer_id	Raw(16)	PRIMARY KEY dont allow nulls, FOREIGN KEY (customer_id) REFERENCES customer(id)	NA	NA
	address_id	Raw(16)	PRIMARY KEY dont allow nulls, FOREIGN KEY (address_id) REFERENCES address(id)		
	default_address - yes/no	BOOLEAN	dont allow nulls		
customer_gift_voucher (relation table)	customer_id	Raw(16)	PRIMARY KEY dont allow nulls, FOREIGN KEY (customer_id) REFERENCES customer(id)	NA	
	gift_voucher_id	Raw(16)	PRIMARY KEY		
	count	NUMBER	dont allow nulls		
customer_delivery_partner (relation table)	customer_id	Raw(16)	PRIMARY KEY dont allow nulls, FOREIGN KEY (customer_id) REFERENCES customer(id)	NA	NA
	delivery_partner_id	Raw(16)	PRIMARY KEY dont allow nulls, FOREIGN KEY (delivery_partner_id) REFERENCES delivery_partner(id)		
products_cart (relation table)	product_id	Raw(16)	PRIMARY KEY FOREIGN KEY (product_id) REFERENCES products(id)	NA	NA
	cart_id	Raw(16)	PRIMARY KEY		
	count	NUMBER	dont allow nulls		
orders_products (relation table)	orders_id	Raw(16)	PRIMARY KEY dont allow nulls FOREIGN KEY (order_id) REFERENCES orders(id)	NA	NA
	products_id	Raw(16)	PRIMARY KEY FOREIGN KEY (product_id) REFERENCES products(id)		
	quantity	INTEGER	dont allow nulls		
products_suppliers	product_id	Raw(16)	PRIMARY KEY FOREIGN KEY (product_id) REFERENCES products(id)		

Analysis					
Table products_suppliers (relation table)	Attributes	Domain Constraints	Check Constraints	Relation	Grant_Access (Customer/Administrator/SuperUser)
	supplier_id	Raw(16)	PRIMARY KEY FOREIGN KEY (supplier_id) REFERENCES supplier(id)	NA	NA
	quantity	INTEGER	dont allow nulls		

E-R Diagram



ERD Transformations:

1. products(id, name, description, price, quantity, brand_name, discount, address_id,
product_category_id)
2. customer(id, first_name, last_name, email_id, phone_no, dob, type)
3. cart(id, quantity, customer_id)
4. orders(id, order_date, status, address_id, customer_id)
5. transaction_summary(id, total_amount, date_of_payment, payment_type, order_id)
6. supplier(id, name, phone_no, email_id, rating, address_id)
7. address(id, flat_no, street, city, state, country, zip_code)
8. product_category(id, section, audience_segment)
9. reviews(id, rating, comments, product_id, order_id, customer_id)
10. gift_vouchers(id, name, number, amount, expiry_date)
11. delivery_partner(id, name, phone, email_id, order_id)
12. customer_address(customer_id, address_id, default_address)
13. customer_gift_voucher (customer_id, gift_voucher_id, count)
14. customer_delivery_partner(customer_id, delivery_partner_id)
15. product_cart(product_id, cart_id, count)
16. orders_products(order_id, product_id, quantity)
17. products_suppliers(product_id, supplier_id, quantity)

Requirement-2:

I've set up Docker to work on DPDB Project 1. Here is a screenshot of the Docker Dashboard.



Fig-1: Docker Dashboard

Once Docker is installed, download the postgres:alpine image from the terminal. The command used to pull the postgres image is “**docker pull postgres:alpine**” As alpine Linux is a lightweight, small size, and secure. used Postgres with alpine.

To verify the list of instances running in Docker, the command used is **“docker ps”**. We can see the container details, such as the container ID, status, ports, etc.




Fig-4: Check the list of Docker instances.

Below is the screenshot, which displays the list of Docker containers in the Docker Dashboard.




Fig-5: Docker Dashboard with the kl-dpdb-node1 instance.

Once the container is ready, we need to execute the container to start working on the respective machine. The command used to execute the container is
“docker exec -it kl-dpdb-node1 bash”




Fig-6: Execute running docker container 'kl-dpdb-node1'.

Now to work with the postgres, we need to connect to postgres from the terminal in docker. To connect to postgres from docker container the command used is
“psql -U postgres”




Fig-7: Connect to Postgres.

Create an ecommerce database to start working on project requirements. Command used to create ecommerce database is

"create database kl_dpdb_ecommerce_database;"




Fig-8: Create ' kl_dpdb_ecommerce_database' database in postgres.

Once the database is created, we see the list of databases using “\l” command. In the screenshot below, we can see the created database i.e., kl_dpdb_ecommerce_database

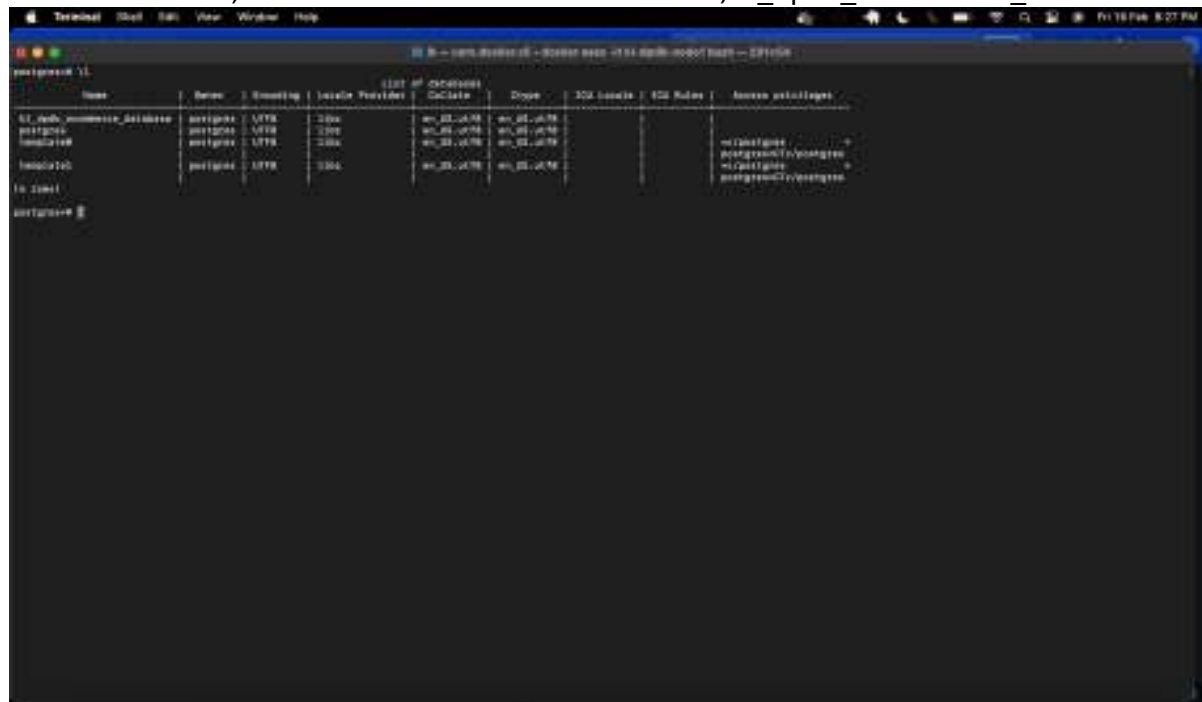


Fig-9: Check the list of databases.

To make use of created database, The command used is “\c kl_dpdb_ecommerce_database”

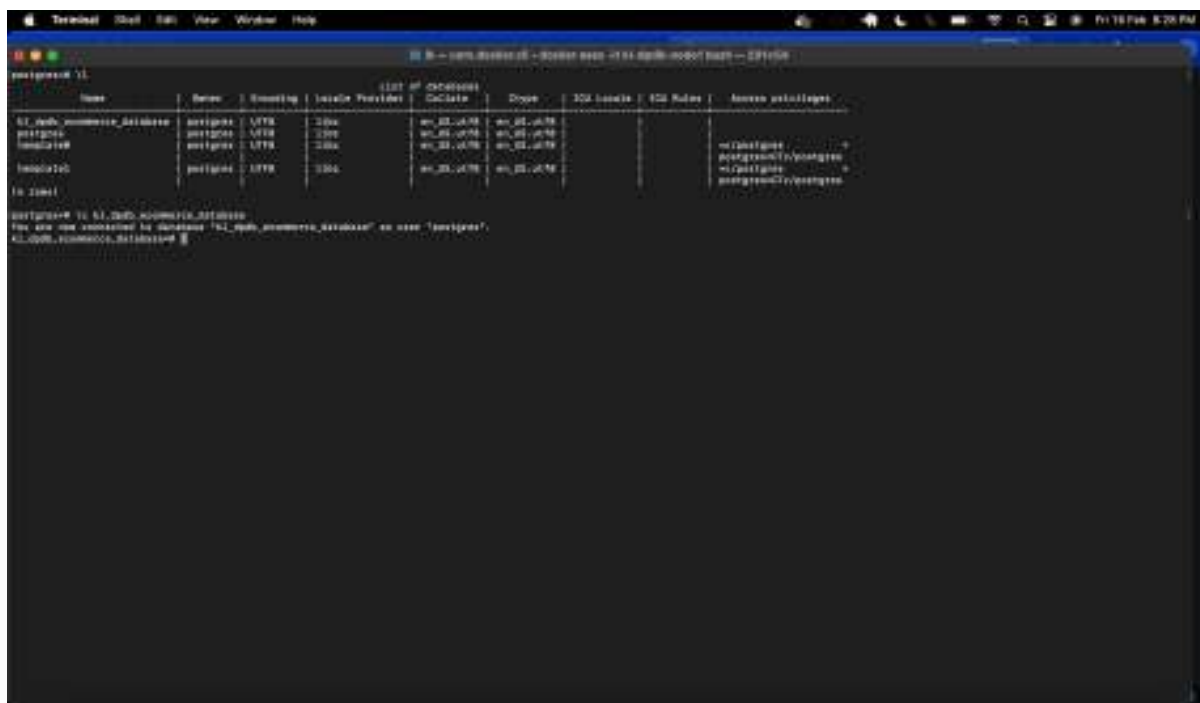


Fig-10: Use kl_dpdb_ecommerce_database.

Creation of Tables in kl_dpdb_ecommerce_database:

To generate the primary key automatically in RAW(16) format, we need to create the extension for (uuid - ossp). The command used to create the extension is:

```
-- uuid-oss: the extension used to create primary key with UUID.  
CREATE EXTENSION IF NOT EXISTS "uuid-oss";
```

The uuid extension helps to generate the random 16-digit number. Which helps have the less collision.

address table:

query:

```
-- Create the address table
```

```
CREATE TABLE address (  
id UUID PRIMARY KEY DEFAULT uuid_generate_v4(),  
flat_no INTEGER NOT NULL,  
street VARCHAR(50) NOT NULL,  
city VARCHAR(50) NOT NULL,  
state VARCHAR(50) NOT NULL,  
country VARCHAR(50) NOT NULL,  
zip_code VARCHAR(10) NOT NULL  
);
```




Fig-11: Create address table and display the schema of it.

Explanation: The Address table is used to store the address details of customers, products, and supplier. The address table has a unique identifier with the attribute flat_no, street, city, state, country, and the zip-code. The address table is one of the crucial tables in the e-commerce site which helps the delivery partners where to pick and deliver the products.

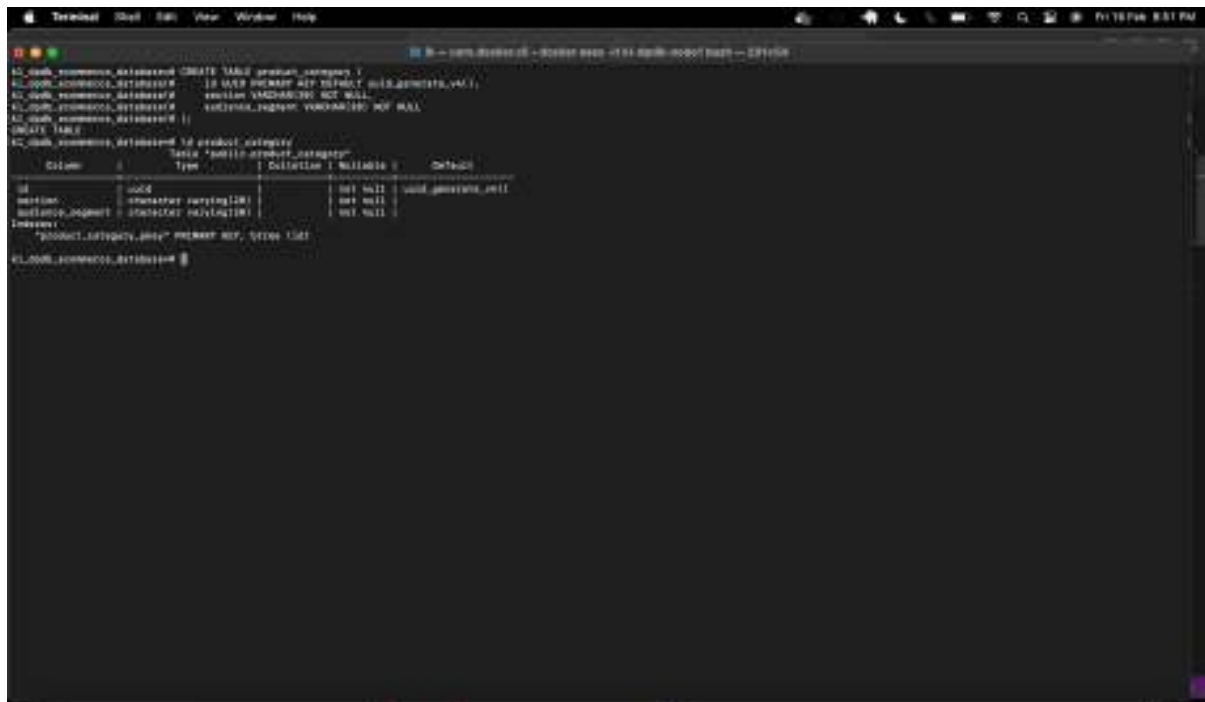
A customer can have many addresses of his/her choice. The customer can order the products from any address.

product_category table:

query:

-- create product_category table

```
CREATE TABLE product_category (  
  id UUID PRIMARY KEY DEFAULT uuid_generate_v4(),  
  section VARCHAR(20) NOT NULL,  
  audience_segment VARCHAR(20) NOT NULL  
);
```



```
kl_@db: ~$ psql -h localhost -U postgres -d ecommerce -c 'CREATE TABLE product_category (  
id UUID PRIMARY KEY DEFAULT uuid_generate_v4(),  
section VARCHAR(20) NOT NULL,  
audience_segment VARCHAR(20) NOT NULL  
);'  
CREATE TABLE  
kl_@db: ~$ \d product_category  
Table "public.product_category"  
  Column          | Type          | Collation | Nullable | Default  
id                 | uuid         |           | not null | uuid_generate_v4()  
section            | varchar(20)  |           | not null |  
audience_segment | varchar(20)  |           | not null |  
Indexes:  
  "product_category_pkey" PRIMARY KEY, btree (id)  
kl_@db: ~$
```

Fig-12: Create the product_category table and display the schema of it.

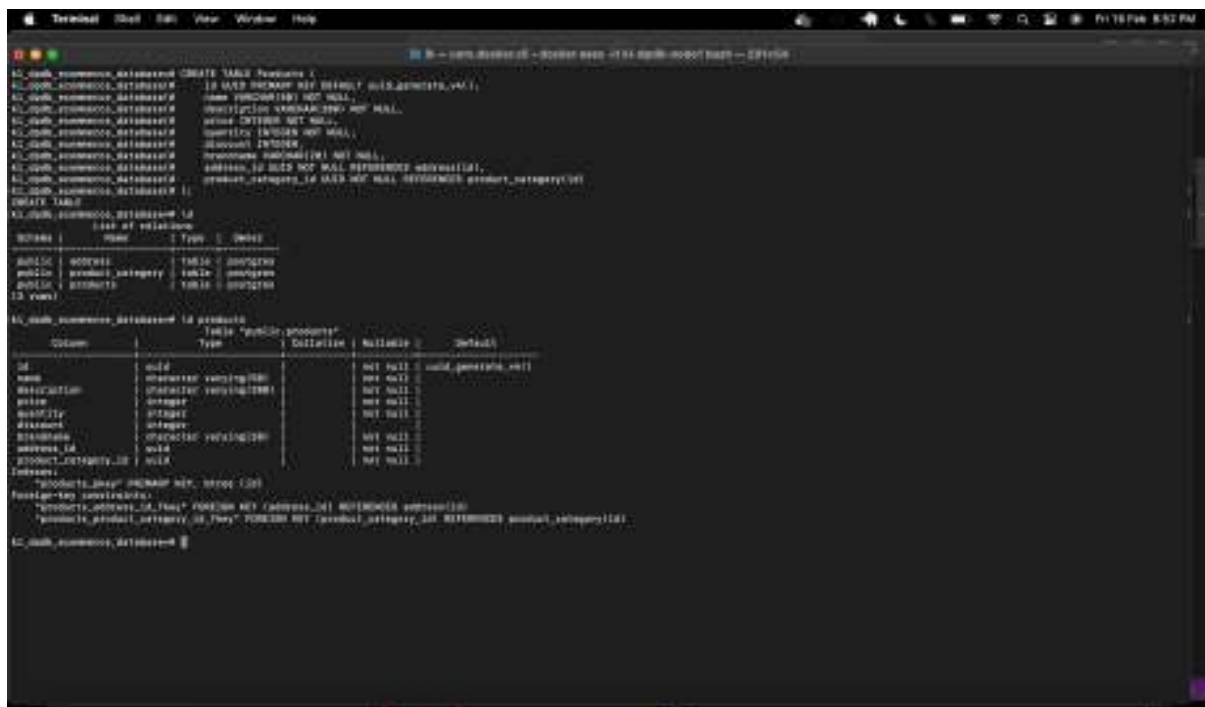
Explanation: The Products_Category table helps the customer to filter the products based on their choice. It helps the customer to search for products easily. There are different types of categories such as clothing, toys, accessories, etc. The product_category has a unique identifier, with attributes section and audience segment. The section talks about, in which section the products are in like clothing, accessories, Footwear, etc. and the audience_segment talks about to whom that product is for like kids, women, adults, girls, etc.

products table:

query:

-- create Products table

```
CREATE TABLE products (  
  id UUID PRIMARY KEY DEFAULT uuid_generate_v4(),  
  name VARCHAR(50) NOT NULL,  
  description VARCHAR(200) NOT NULL,  
  price INTEGER NOT NULL,  
  quantity INTEGER NOT NULL,  
  discount INTEGER,  
  brandname VARCHAR(20) NOT NULL,  
  address_id UUID NOT NULL REFERENCES address(id),  
  product_category_id UUID NOT NULL REFERENCES product_category(id)  
);
```



```
15.ash@commerce-database:~$ CREATE TABLE products (  
16.ash@commerce-database:~$ id UUID PRIMARY KEY DEFAULT uuid_generate_v4(),  
17.ash@commerce-database:~$ name VARCHAR(50) NOT NULL,  
18.ash@commerce-database:~$ description VARCHAR(200) NOT NULL,  
19.ash@commerce-database:~$ price INTEGER NOT NULL,  
20.ash@commerce-database:~$ quantity INTEGER NOT NULL,  
21.ash@commerce-database:~$ discount INTEGER,  
22.ash@commerce-database:~$ brandname VARCHAR(20) NOT NULL,  
23.ash@commerce-database:~$ address_id UUID NOT NULL REFERENCES address(id),  
24.ash@commerce-database:~$ product_category_id UUID NOT NULL REFERENCES product_category(id)  
25.ash@commerce-database:~$ );  
26.ash@commerce-database:~$  
27.ash@commerce-database:~$ \d products  
Table "public.products"  
  Column          | Type          | Collation | Nullable | Default  
-----+-----+-----+-----+-----  
id                 | uuid         |           | not null | uuid_generate_v4()  
name              | varchar(50)  |           | not null |  
description        | varchar(200) |           | not null |  
price             | integer      |           | not null |  
quantity          | integer      |           | not null |  
discount          | integer      |           |          |  
brandname         | varchar(20)  |           | not null |  
address_id        | uuid         |           | not null |  
product_category_id | uuid         |           | not null |  
Indexes:  
  "products_pkey" PRIMARY KEY (id)  
Foreign-key constraints:  
  "products_address_id_fkey" FOREIGN KEY (address_id) REFERENCES address(id)  
  "products_product_category_id_fkey" FOREIGN KEY (product_category_id) REFERENCES product_category(id)  
16.ash@commerce-database:~$
```

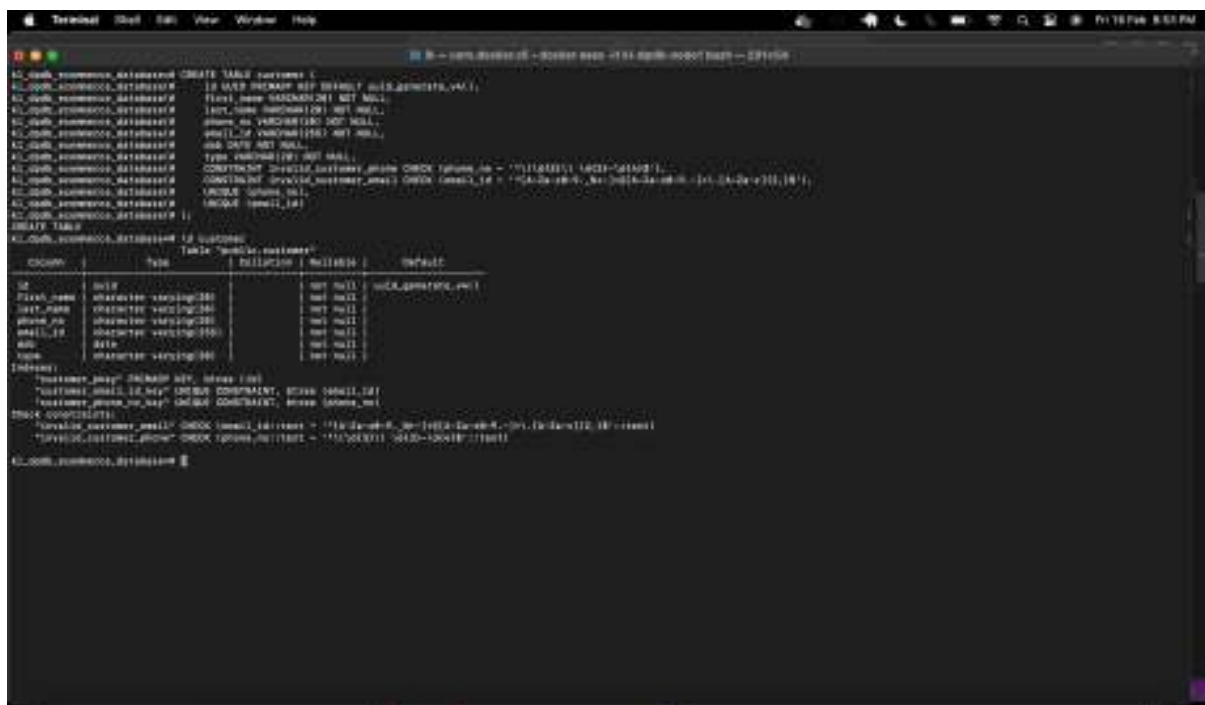
Fig-13: Create the products table and display the schema of it.

Explanation: The Products contain the list of products/items available. The products have a unique identifier with attributes name, description, price, quantity, discount, brandname, address_id and product_category_id. The address_id and the product_category_id are the foreign keys which refers to address and products_category table.

customer table:

query:

```
-- create customer table
CREATE TABLE customer (
id UUID PRIMARY KEY DEFAULT uuid_generate_v4(),
first_name VARCHAR(20) NOT NULL,
last_name VARCHAR(20) NOT NULL,
phone_no VARCHAR(20) NOT NULL,
email_id VARCHAR(255) NOT NULL,
dob DATE NOT NULL,
type VARCHAR(20) NOT NULL,
CONSTRAINT invalid_customer_phone CHECK (phone_no ~ '^\(\\d{3}\\) \\d{3}-\\d{4}$'),
CONSTRAINT invalid_customer_email CHECK (email_id ~ '^[A-Za-z0-9._%+-]+@[A-Za-z0-9.-
]+\.[A-Za-z]{2,}$'),
UNIQUE (phone_no),
UNIQUE (email_id)
);
```

The image shows a terminal window with a dark background. At the top, there's a title bar for a terminal application. The main area displays the SQL code for creating a table named 'customer'. The code includes column definitions for 'id' (UUID, primary key), 'first_name', 'last_name', 'phone_no', 'email_id', 'dob', and 'type'. It also includes two CHECK constraints for phone and email formats, and two UNIQUE constraints for phone_no and email_id. Below the CREATE TABLE statement, the terminal shows the output of the \d command, which displays the table's schema in a structured format, including column names, data types, and constraints.

```
CREATE TABLE customer (
id UUID PRIMARY KEY DEFAULT uuid_generate_v4(),
first_name VARCHAR(20) NOT NULL,
last_name VARCHAR(20) NOT NULL,
phone_no VARCHAR(20) NOT NULL,
email_id VARCHAR(255) NOT NULL,
dob DATE NOT NULL,
type VARCHAR(20) NOT NULL,
CONSTRAINT invalid_customer_phone CHECK (phone_no ~ '^\(\\d{3}\\) \\d{3}-\\d{4}$'),
CONSTRAINT invalid_customer_email CHECK (email_id ~ '^[A-Za-z0-9._%+-]+@[A-Za-z0-9.-
]+\.[A-Za-z]{2,}$'),
UNIQUE (phone_no),
UNIQUE (email_id)
);
```

Column	Type	Nullable	Default
id	uuid	not null	uuid_generate_v4()
first_name	varchar(20)	not null	
last_name	varchar(20)	not null	
phone_no	varchar(20)	not null	
email_id	varchar(255)	not null	
dob	date	not null	
type	varchar(20)	not null	

Fig-14: Create the customer table and display the schema of it.

Explanation: The customer table contains the list of customer details. This one is the crucial table. The customer table has a unique identifier with attributes first_name, last_name, phone_no, email_id, dob, and type. The phone_no and email_id have unique check constraints which helps to avoid duplicate customer accounts. The dob of the customer helps to provide the discount/gift voucher to a specific user in the birth month. The type of customer helps to identify whether the customer is new to the e-commerce site and based on the purchases the level of the customer increases like premium, gold, VIP etc.

```
CREATE TABLE cart (
id UUID PRIMARY KEY DEFAULT uuid_generate_v4(),
customer_id UUID NOT NULL REFERENCES customer(id),
quantity INTEGER NOT NULL
);
```




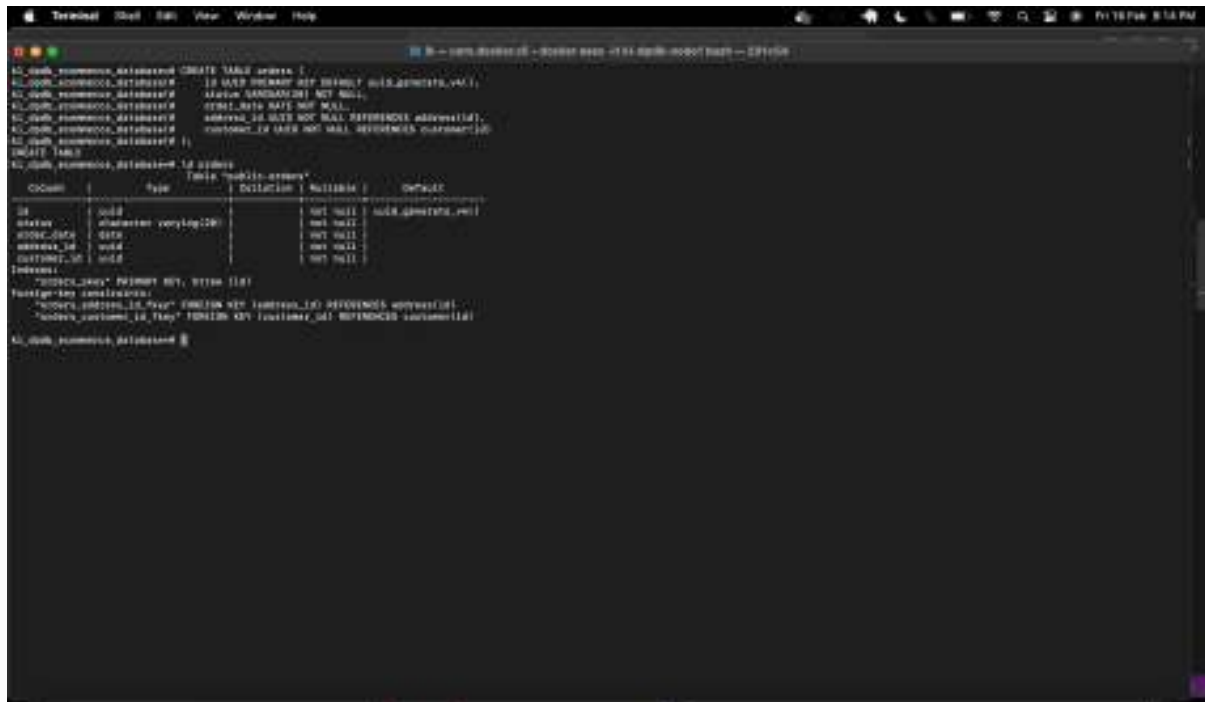
Fig-15: Create the cart table and display the schema of it.

Explanation: The cart table contains the list of products that customers want to buy. Each customer has a cart where they add/store the products. The cart table has a unique identifier with attributes customer_id and quantity. The customer_id is the foreign key reference to the customer table and the quantity talks about how many products the customer added to the cart.

orders table:

query:

```
-- create orders table
CREATE TABLE orders (
  id UUID PRIMARY KEY DEFAULT uuid_generate_v4(),
  status VARCHAR(20) NOT NULL,
  order_date DATE NOT NULL,
  address_id UUID NOT NULL REFERENCES address(id),
  customer_id UUID NOT NULL REFERENCES customer(id)
);
```



```
kl_@db: postgresql=# CREATE TABLE orders (
kl_@db: postgresql=#   id UUID PRIMARY KEY DEFAULT uuid_generate_v4(),
kl_@db: postgresql=#   status VARCHAR(20) NOT NULL,
kl_@db: postgresql=#   order_date DATE NOT NULL,
kl_@db: postgresql=#   address_id UUID NOT NULL REFERENCES address(id),
kl_@db: postgresql=#   customer_id UUID NOT NULL REFERENCES customer(id)
kl_@db: postgresql=# );
CREATE TABLE
kl_@db: postgresql=# \d orders
Table "public.orders"
  Column      | Type          | Collation | Nullable | Default
-----
id            | uuid         |           | not null | uuid_generate_v4()
status       | varchar(20)  |           | not null |
order_date   | date         |           | not null |
address_id    | uuid         |           | not null |
customer_id   | uuid         |           | not null |
Indexes:
    "orders_pkey" PRIMARY KEY, btree (id)
Foreign-key constraints:
    "orders_address_id_fkey" FOREIGN KEY (address_id) REFERENCES address(id)
    "orders_customer_id_fkey" FOREIGN KEY (customer_id) REFERENCES customer(id)
kl_@db: postgresql=#
```

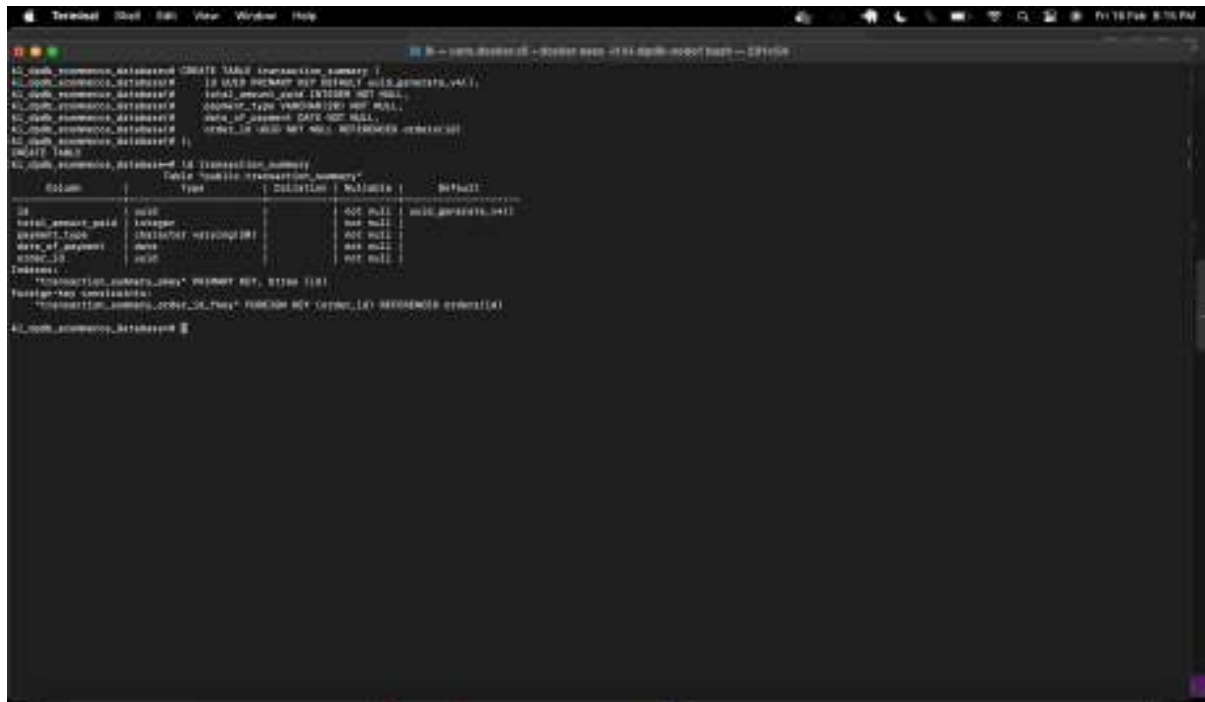
Fig-16: Create the orders table and display the schema of it.

Explanation: when the customer orders any product all the respective details are stored in the orders table. The orders table has a unique identifier with attributes status, order_date, address_id and customer_id. The status talks about the product delivery status such as processing, in transit, delivered etc. order_date talks about the on which date the order has placed by the customer. Here the address_id and customer_id are the foreign keys which refer to address table and customer table. This helps easily to check the orders details if the customer.

transaction_summary table:

query:

```
-- create transaction_summary table
CREATE TABLE transaction_summary (
id UUID PRIMARY KEY DEFAULT uuid_generate_v4(),
total_amount_paid INTEGER NOT NULL,
payment_type VARCHAR(20) NOT NULL,
date_of_payment DATE NOT NULL,
order_id UUID NOT NULL REFERENCES orders(id)
);
```



```
kl_mook_sreenivas@kl_mook_sreenivas:~/workspace$ psql -h localhost -u postgres -d postgres
psql (13.10)
Type "help" for help.

postgres=# CREATE TABLE transaction_summary (
postgres=# id UUID PRIMARY KEY DEFAULT uuid_generate_v4(),
postgres=# total_amount_paid INTEGER NOT NULL,
postgres=# payment_type VARCHAR(20) NOT NULL,
postgres=# date_of_payment DATE NOT NULL,
postgres=# order_id UUID NOT NULL REFERENCES orders(id)
postgres=# );
CREATE TABLE
postgres=# \d transaction_summary
Table "public.transaction_summary"
Column | Type | Collation | Nullable | Default
-----+-----+-----+-----+-----
id      | uuid |           | not null | uuid_generate_v4()
total_amount_paid | integer |           | not null |
payment_type | varchar(20) |           | not null |
date_of_payment | date |           | not null |
order_id | uuid |           | not null |
Indexes:
    "transaction_summary_pkey" PRIMARY KEY, btree (id)
Foreign-key constraints:
    "transaction_summary_order_id_fkey" FOREIGN KEY (order_id) REFERENCES orders(id)
```

Fig-17: Create the transaction_summary table and display the schema of it.

Explanation: The transaction_summary table talks about the transactions done by the customers. The transaction_summary has a unique identifier with attributes total_amount_paid, payment_type, date_of_payment, order_id. The total_amount, talks about the total amount paid to by the customer on specific order. payment_type talks about the through which type of payment did customer placed the order such as Debit card, credit card, etc. The order_id is the foreign key which references to transaction_summary table.

supplier table:

query:

```
-- create supplier table
```

```
CREATE TABLE supplier (  
  id UUID PRIMARY KEY DEFAULT uuid_generate_v4(),  
  name VARCHAR(50) NOT NULL,  
  phone_no VARCHAR(20) NOT NULL UNIQUE CONSTRAINT invalid_supplier_phone_no CHECK  
    (phone_no ~ '^\(\d{3}\) \d{3}-\d{4}$'),  
  email VARCHAR(255) NOT NULL UNIQUE CONSTRAINT invalid_supplier_email CHECK (email ~  
    '^[A-Za-z0-9._%+-]+@[A-Za-z0-9.-]+\.[A-Za-z]{2,}$'),  
  rating INTEGER,  
  address_id UUID NOT NULL REFERENCES address(id)  
);
```

The screenshot shows a PostgreSQL terminal window with the following SQL commands and their output:

```
kl_spm_spm@kl_spm:~$ CREATE TABLE supplier (  
  id UUID PRIMARY KEY DEFAULT uuid_generate_v4(),  
  name VARCHAR(50) NOT NULL,  
  phone_no VARCHAR(20) NOT NULL UNIQUE CONSTRAINT invalid_supplier_phone_no CHECK  
    (phone_no ~ '^\(\d{3}\) \d{3}-\d{4}$'),  
  email VARCHAR(255) NOT NULL UNIQUE CONSTRAINT invalid_supplier_email CHECK (email ~  
    '^[A-Za-z0-9._%+-]+@[A-Za-z0-9.-]+\.[A-Za-z]{2,}$'),  
  rating INTEGER,  
  address_id UUID NOT NULL REFERENCES address(id)  
);  
CREATE TABLE
```

Then, the command `\d supplier` is executed, showing the table schema:

Column	Type	Collation	Nullable	Default
id	uuid		not null	uuid_generate_v4()
name	character varying(50)		not null	
phone_no	character varying(20)		not null	
email	character varying(255)		not null	
rating	integer		not null	
address_id	uuid		not null	

Finally, the command `\d+ supplier` is executed, showing the table constraints:

```
Table "supplier"  
Columns:  
  "id" uuid PRIMARY KEY DEFAULT uuid_generate_v4()  
  "name" character varying(50) NOT NULL  
  "phone_no" character varying(20) NOT NULL  
  "email" character varying(255) NOT NULL  
  "rating" integer  
  "address_id" uuid NOT NULL  
Indexes:  
  "supplier_pkey" PRIMARY KEY, btree (id)  
  "supplier_email_key" UNIQUE CONSTRAINT, btree (email)  
  "supplier_phone_no_key" UNIQUE CONSTRAINT, btree (phone_no)  
Check constraints:  
  "invalid_supplier_email" CHECK (email ~ '^[A-Za-z0-9._%+-]+@[A-Za-z0-9.-]+\.[A-Za-z]{2,}$')  
  "invalid_supplier_phone_no" CHECK (phone_no ~ '^\(\d{3}\) \d{3}-\d{4}$')  
Foreign key constraints:  
  "supplier_address_id_fkey" FOREIGN KEY (address_id) REFERENCES address(id)
```

Fig-18: Create the supplier table and display the schema of it.

Explanation: The supplier table stores the details of the supplier. The supplier table has a unique identifier with attributes name, phone_no, email, rating and address_id. The name identifies the name of the supplier, rating helps to purchase the product easily. The phone_no and email has the check constraints to have security and unique to avoid the duplicate accounts. The address_id is the foreign key which references the address table.

reviews table:

query:

```
-- create reviews table
CREATE TABLE reviews (
id UUID PRIMARY KEY DEFAULT uuid_generate_v4(),
rating INTEGER,
comments VARCHAR(150),
product_id UUID REFERENCES products(id),
order_id UUID REFERENCES orders(id),
customer_id UUID REFERENCES customer(id)
);
```



```
kl_madh_sreenivas@madhmad:~$ CREATE TABLE reviews (
kl_madh_sreenivas@madhmad:~$ \d reviews
Table "public.reviews"
  Column      | Type          | Collation | Nullable | Default
id             | uuid         |           | not null | uuid_generate_v4()
rating        | integer      |           |          |
comments      | varchar(150) |           |          |
product_id    | uuid         |           |          |
order_id      | uuid         |           |          |
customer_id   | uuid         |           |          |
Indexes:
    "reviews_pkey" PRIMARY KEY, btree (id)
Foreign-key constraints:
    "reviews_customer_id_fkey" FOREIGN KEY (customer_id) REFERENCES customer(id)
    "reviews_order_id_fkey" FOREIGN KEY (order_id) REFERENCES orders(id)
    "reviews_product_id_fkey" FOREIGN KEY (product_id) REFERENCES products(id)
kl_madh_sreenivas@madhmad:~$
```

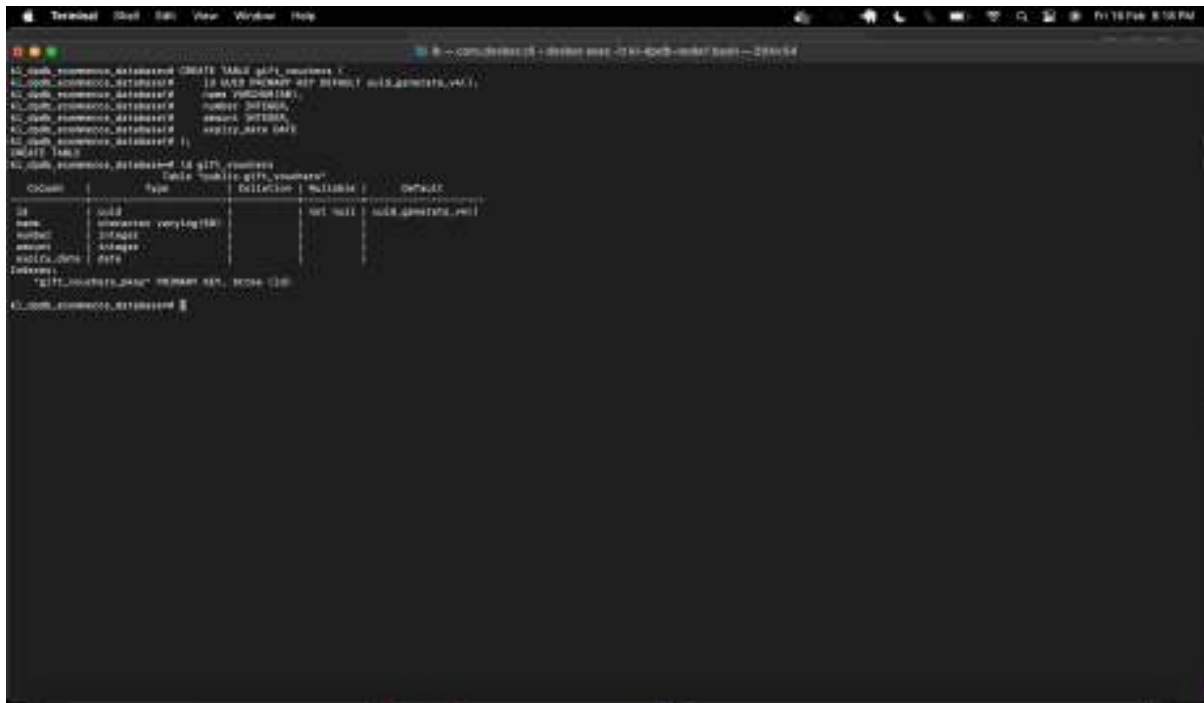
Fig-19: Create the reviews table and display the schema of it.

Explanation: The reviews table plays an important role in knowing more about the product. As the reviews are written by the customers, there will be no chance of false marketing about the products. The reviews table has the unique identifier with attributes rating, comments, product_id, customer_id and order_id. The rating talks about what the rating of the product is on scale of 5 like 1,2,3,4,5 etc. The comments help them to write their opinion about the product how did they felt about it etc. The product_id, customer_id and order_id are the foreign keys which refers to the products, customer and order tables.

gift_vouchers table:

query:

```
-- create gift_vouchers table
CREATE TABLE gift_vouchers (
id UUID PRIMARY KEY DEFAULT uuid_generate_v4(),
name VARCHAR(50),
number INTEGER,
amount INTEGER,
expiry_date DATE
);
```



The screenshot shows a terminal window with a dark background. The top part of the window displays the SQL command to create the `gift_vouchers` table. Below the command, the terminal shows the output of the `\d` command, which displays the table's schema. The schema is presented as a table with columns: Column, Type, Collation, Nullable, and Default. The columns in the table are `id` (UUID, PRIMARY KEY, DEFAULT uuid_generate_v4()), `name` (VARCHAR(50)), `number` (INTEGER), `amount` (INTEGER), and `expiry_date` (DATE). The `id` column is marked as the primary key.

```
CREATE TABLE gift_vouchers (
id UUID PRIMARY KEY DEFAULT uuid_generate_v4(),
name VARCHAR(50),
number INTEGER,
amount INTEGER,
expiry_date DATE
);
```

Column	Type	Collation	Nullable	Default
id	uuid		not null	uuid_generate_v4()
name	varchar(50)			
number	integer			
amount	integer			
expiry_date	date			

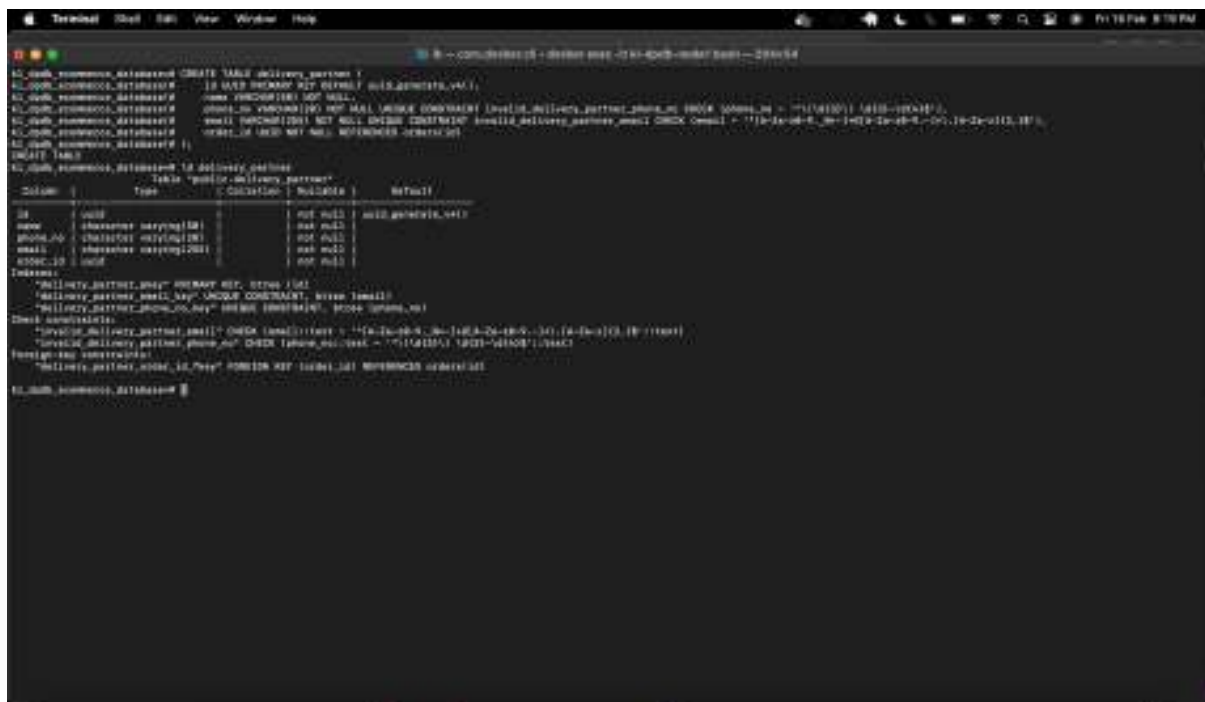
Fig-20: Create the gift_vouchers table and display the schema of it.

Explanation: The gift_vouchers table stores the list of gift voucher details. The gift voucher has a unique identifier and with attributes name, number, amount and expiry_date. The name identifies the name of the gift voucher, number stores the gift voucher number, amount stores the amount allocated to the gift voucher, expiry_date contains date on which day the voucher is going to get expired. The gift voucher is one the marketing techniques used by many of the e-commerce sites.

delivery_partner table:

query:

```
-- create delivery_partner table
CREATE TABLE delivery_partner (
id UUID PRIMARY KEY DEFAULT uuid_generate_v4(),
name VARCHAR(50) NOT NULL,
phone_no VARCHAR (20) NOT NULL CONSTRAINT invalid_delivery_partner_phone_no CHECK
(phone_no ~ '^\(\\d{3}\\) \\d{3}-\\d{4}$'),
email VARCHAR(255) NOT NULL CONSTRAINT invalid_delivery_partner_email CHECK (email ~
'^[A-Za-z0-9._%+-]+@[A-Za-z0-9.-]+\.[A-Za-z]{2,}$'),
order_id UUID NOT NULL REFERENCES orders(id)
);
```



```
Terminal [SQL] [SQL] View Window Help
-- create delivery_partner table
CREATE TABLE delivery_partner (
id UUID PRIMARY KEY DEFAULT uuid_generate_v4(),
name VARCHAR(50) NOT NULL,
phone_no VARCHAR (20) NOT NULL CONSTRAINT invalid_delivery_partner_phone_no CHECK
(phone_no ~ '^\(\\d{3}\\) \\d{3}-\\d{4}$'),
email VARCHAR(255) NOT NULL CONSTRAINT invalid_delivery_partner_email CHECK (email ~
'^[A-Za-z0-9._%+-]+@[A-Za-z0-9.-]+\.[A-Za-z]{2,}$'),
order_id UUID NOT NULL REFERENCES orders(id)
);

CREATE TABLE
Table "public.delivery_partner"
Column | Type | Collation | Nullable | Default
-----+-----+-----+-----+-----
id      | uuid |           | not null | uuid_generate_v4()
name    | varchar(50) |           | not null |
phone_no | varchar(20) |           | not null |
email   | varchar(255) |           | not null |
order_id | uuid |           | not null |
Indexes:
    "delivery_partner_pkey" PRIMARY KEY, btree (id)
    "delivery_partner_email_key" UNIQUE CONSTRAINT, btree (email)
    "delivery_partner_phone_no_key" UNIQUE CONSTRAINT, btree (phone_no)
Check constraints:
    "valid_delivery_partner_email" CHECK (email ~ '^[A-Za-z0-9._%+-]+@[A-Za-z0-9.-]+\.[A-Za-z]{2,}$')
    "valid_delivery_partner_phone_no" CHECK (phone_no ~ '^\(\\d{3}\\) \\d{3}-\\d{4}$')
Foreign-key constraints:
    "delivery_partner_order_id_fkey" FOREIGN KEY (order_id) REFERENCES orders(id)
```

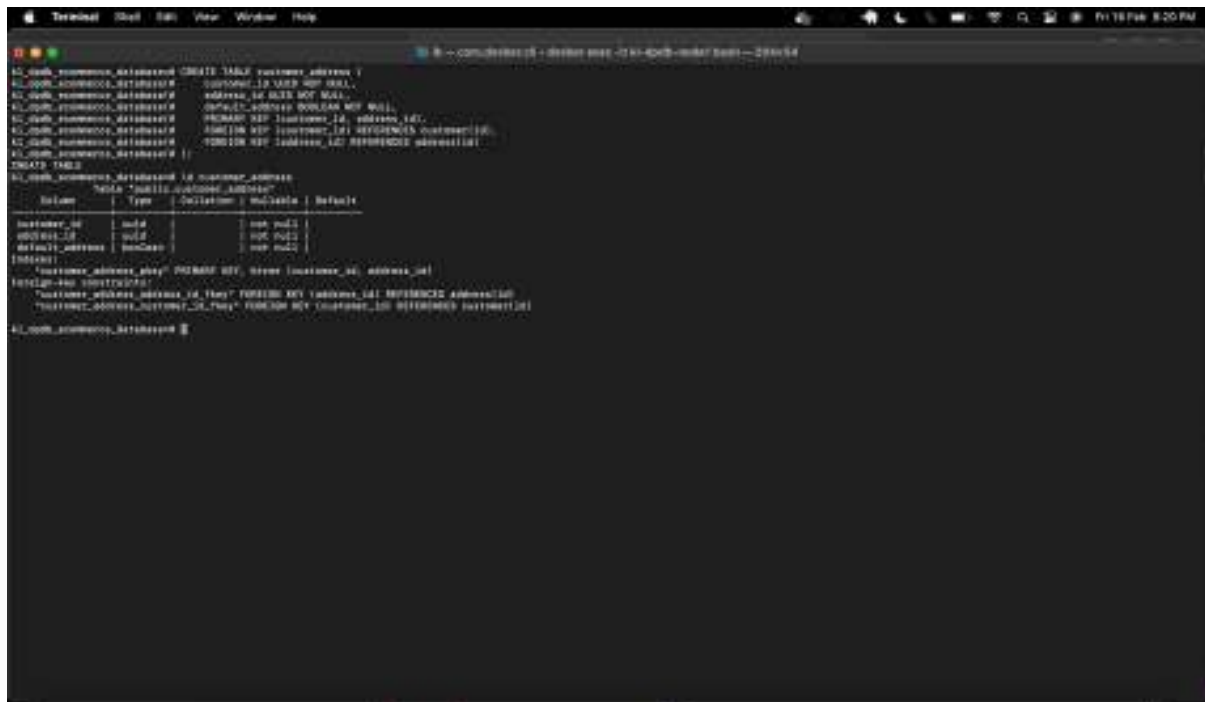
Fig-21: Create the delivery_partner table and display the schema of it.

Explanation: The delivery_partner table stores the delivery_partner details and the order associated to respective delivery partner. The delivery partner has a unique identifier with attributes name, phone_no, email, order_id. The name is the name of the delivery partner, phone_no and email are the details of the delivery partner which helps the customers to contact them. The phone_no and email has check constraints for security reasons. The order_id is the foreign key which references the orders table.

customer_address table:

query:

```
-- create customer_address table
CREATE TABLE customer_address (
customer_id UUID NOT NULL,
address_id UUID NOT NULL,
default_address BOOLEAN NOT NULL,
PRIMARY KEY (customer_id, address_id),
FOREIGN KEY (customer_id) REFERENCES customer(id),
FOREIGN KEY (address_id) REFERENCES address(id)
);
```



```
al_madh@postgres:~/postgres$ CREATE TABLE customer_address (
al_madh@postgres:~/postgres$ customer_id UUID NOT NULL,
al_madh@postgres:~/postgres$ address_id UUID NOT NULL,
al_madh@postgres:~/postgres$ default_address BOOLEAN NOT NULL,
al_madh@postgres:~/postgres$ PRIMARY KEY (customer_id, address_id),
al_madh@postgres:~/postgres$ FOREIGN KEY (customer_id) REFERENCES customer(id),
al_madh@postgres:~/postgres$ FOREIGN KEY (address_id) REFERENCES address(id)
al_madh@postgres:~/postgres$ );
CREATE TABLE
al_madh@postgres:~/postgres$ \d customer_address
Table "public.customer_address"
  Column      | Type          | Collation | Nullable | Default |
--------------+--------------+-----------+----------+-----
customer_id   | uuid         |           | not null |
address_id    | uuid         |           | not null |
default_address | boolean      |           | not null |
Indexes:
    1. "customer_address_pkey" PRIMARY KEY, btree (customer_id, address_id)
Foreign-key constraints:
    "customer_address_address_id_fkey" FOREIGN KEY (address_id) REFERENCES address(id)
    "customer_address_customer_id_fkey" FOREIGN KEY (customer_id) REFERENCES customer(id)
```

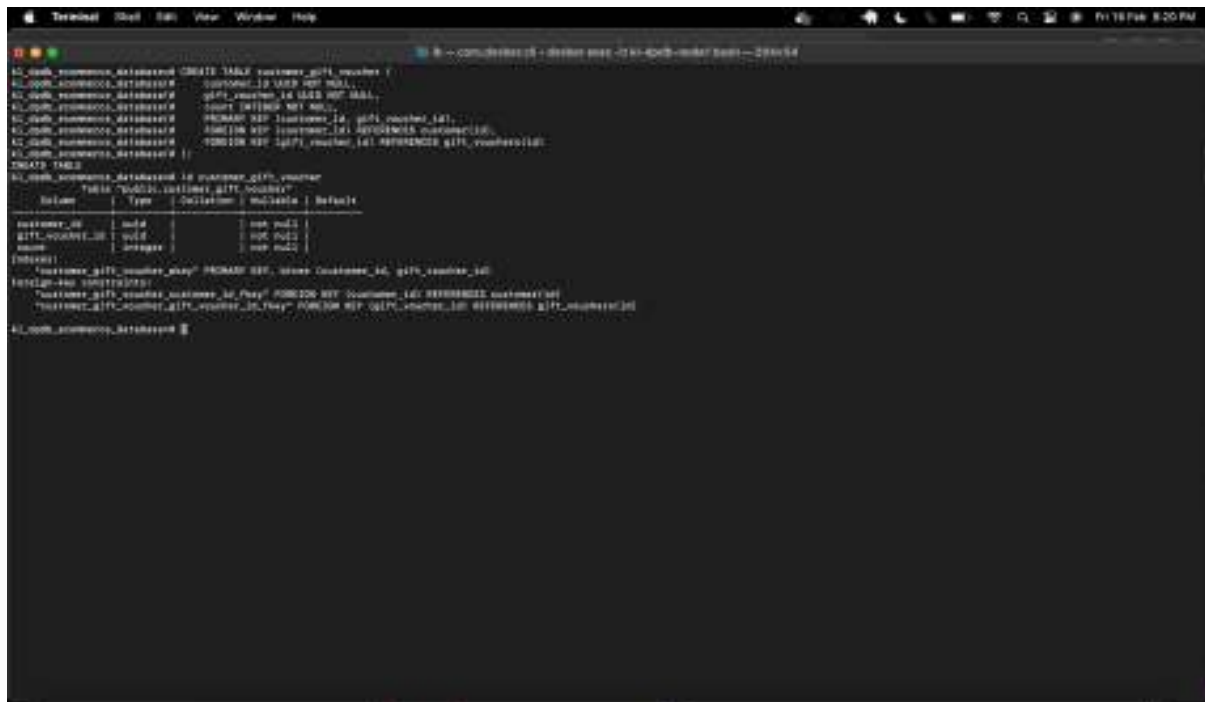
Fig-22: Create the customer_address table and display the schema of it.

Explanation: The relationship between the customer and address is many-to-many. This relation table forms the intermediary link between the customer and address. In this relation table customer_id and address_id is the primary and foreign keys which refer to the customer and address table. There is an attribute, default address – which tells that default address of the customer. This attribute majorly helps when a customer has more than one address.

customer_gift_voucher table:

query:

```
-- create customer_gift_voucher table
CREATE TABLE customer_gift_voucher (
customer_id UUID NOT NULL,
gift_voucher_id UUID NOT NULL,
count INTEGER NOT NULL,
PRIMARY KEY (customer_id, gift_voucher_id),
FOREIGN KEY (customer_id) REFERENCES customer(id),
FOREIGN KEY (gift_voucher_id) REFERENCES gift_vouchers(id)
);
```

A screenshot of a terminal window with a dark background. The terminal shows a series of SQL commands being executed in a PostgreSQL environment. The commands include creating a table named 'customer_gift_voucher' with columns 'customer_id', 'gift_voucher_id', and 'count'. It also shows the creation of primary and foreign keys. After the commands, the user enters '\d customer_gift_voucher' to display the table's schema. The output shows the table structure with data types (UUID, INTEGER), constraints (NOT NULL, PRIMARY KEY, FOREIGN KEY), and a summary of the table's properties.

```
kl_mah_sreenivas@klmahmah:~$ CREATE TABLE customer_gift_voucher (
kl_mah_sreenivas@klmahmah:~$ customer_id UUID NOT NULL,
kl_mah_sreenivas@klmahmah:~$ gift_voucher_id UUID NOT NULL,
kl_mah_sreenivas@klmahmah:~$ count INTEGER NOT NULL,
kl_mah_sreenivas@klmahmah:~$ PRIMARY KEY (customer_id, gift_voucher_id),
kl_mah_sreenivas@klmahmah:~$ FOREIGN KEY (customer_id) REFERENCES customer(id),
kl_mah_sreenivas@klmahmah:~$ FOREIGN KEY (gift_voucher_id) REFERENCES gift_vouchers(id)
kl_mah_sreenivas@klmahmah:~$ );
CREATE TABLE
kl_mah_sreenivas@klmahmah:~$ \d customer_gift_voucher
Table "public.customer_gift_voucher"
  Column      | Type          | Collation | Nullable | Default |
--------------+--------------+-----------+----------+-----
customer_id   | uuid         |           | not null |         |
gift_voucher_id | uuid        |           | not null |         |
count         | integer      |           | not null |         |
Indexes:
    "customer_gift_voucher_pkey" PRIMARY KEY, btree (customer_id, gift_voucher_id)
Foreign-key constraints:
    "customer_gift_voucher_customer_id_fkey" FOREIGN KEY (customer_id) REFERENCES customer(id)
    "customer_gift_voucher_gift_voucher_id_fkey" FOREIGN KEY (gift_voucher_id) REFERENCES gift_vouchers(id)
```

Fig-23: Create the Customer_gift_voucher table and display the schema of it.

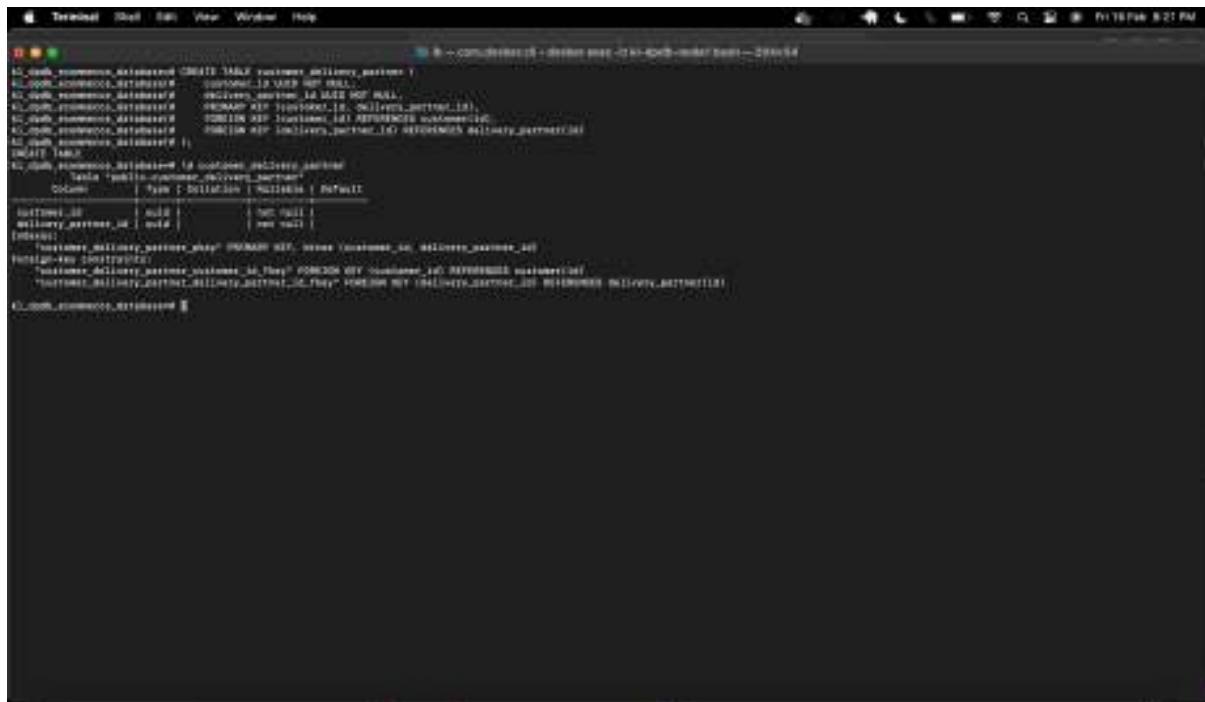
Explanation: The relationship between the customer and the gift voucher is many-to-many. This relation table forms the intermediary link between the customer and the gift voucher. The customer_id and gift_voucher_id are the primary and foreign keys, refers to the customer and gift_vouchers table. There is an attribute count which talks about the number of gift vouchers allocated to the customer.

customer_delivery_partner table:

query

-- create customer_delivery_partner table

```
CREATE TABLE customer_delivery_partner (  
  customer_id UUID NOT NULL,  
  delivery_partner_id UUID NOT NULL,  
  PRIMARY KEY (customer_id, delivery_partner_id),  
  FOREIGN KEY (customer_id) REFERENCES customer(id),  
  FOREIGN KEY (delivery_partner_id) REFERENCES delivery_partner(id)  
);
```



```
kl_@db:~/commerce_database$ CREATE TABLE customer_delivery_partner (  
kl_@db:~/commerce_database$ customer_id UUID NOT NULL,  
kl_@db:~/commerce_database$ delivery_partner_id UUID NOT NULL,  
kl_@db:~/commerce_database$ PRIMARY KEY (customer_id, delivery_partner_id),  
kl_@db:~/commerce_database$ FOREIGN KEY (customer_id) REFERENCES customer(id),  
kl_@db:~/commerce_database$ FOREIGN KEY (delivery_partner_id) REFERENCES delivery_partner(id)  
kl_@db:~/commerce_database$ );  
kl_@db:~/commerce_database$ \d customer_delivery_partner  
Table "public.customer_delivery_partner"  
Column | Type | Collation | Nullable | Default  
-----+-----+-----+-----+-----  
customer_id | uuid | | not null |  
delivery_partner_id | uuid | | not null |  
Indexes:  
1. "customer_delivery_partner_pkey" PRIMARY KEY, btree (customer_id, delivery_partner_id)  
Foreign-key constraints:  
1. "customer_delivery_partner_customer_id_fkey" FOREIGN KEY (customer_id) REFERENCES customer(id)  
2. "customer_delivery_partner_delivery_partner_id_fkey" FOREIGN KEY (delivery_partner_id) REFERENCES delivery_partner(id)  
kl_@db:~/commerce_database$
```

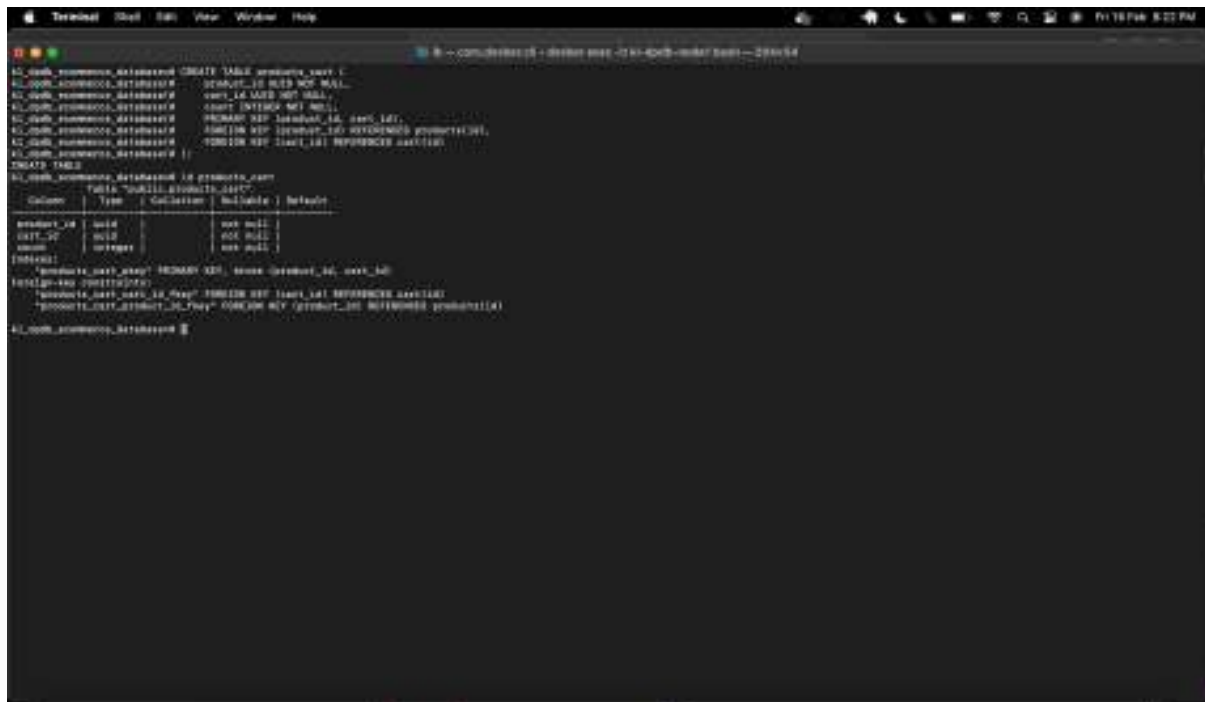
Fig-24: Display and Create the customer_delivery_partner table and display the schema of it.

Explanation: The relationship between the customer and the delivery partner is many-to-many. This relation table forms the intermediary link between customer and the delivery_partner. The table contains customer_id and delivery_partner_id as primary and the foreign keys which refer to customer and delivery_partner tables.

products_cart table:

query:

```
-- create products_cart table
CREATE TABLE products_cart (
product_id UUID NOT NULL,
cart_id UUID NOT NULL,
count INTEGER NOT NULL,
PRIMARY KEY (product_id, cart_id),
FOREIGN KEY (product_id) REFERENCES products(id),
FOREIGN KEY (cart_id) REFERENCES cart(id)
);
```

A screenshot of a terminal window with a dark background. The terminal shows the execution of a SQL query to create a table named 'products_cart'. The query is as follows:

```
CREATE TABLE products_cart (
product_id UUID NOT NULL,
cart_id UUID NOT NULL,
count INTEGER NOT NULL,
PRIMARY KEY (product_id, cart_id),
FOREIGN KEY (product_id) REFERENCES products(id),
FOREIGN KEY (cart_id) REFERENCES cart(id)
);
```


Below the query, the terminal displays the schema of the created table. It shows the table name 'products_cart' and its columns: 'product_id' (UUID, NOT NULL), 'cart_id' (UUID, NOT NULL), and 'count' (INTEGER, NOT NULL). The primary key is listed as '(product_id, cart_id)'. Foreign key constraints are also shown: 'product_id' references 'products(id)' and 'cart_id' references 'cart(id)'. The terminal window has a title bar with 'Terminal' and standard macOS window controls. The top status bar shows the date and time as 'Fri, 18 Feb 8:22 PM'.

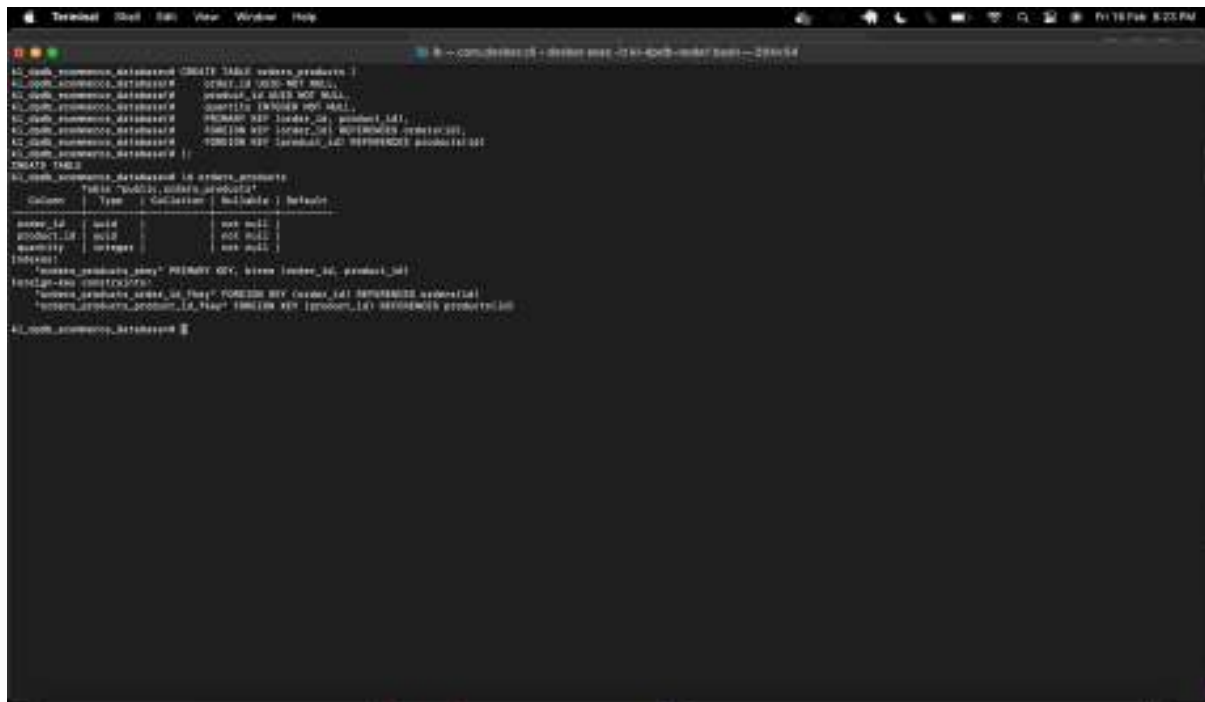
Fig-25: Create the products_cart table and display the schema of it.

Explanation: The relationship between the products and the cart table is many-to-many. The intermediary between these two tables forms the relation table. The table has the primary and foreign keys, product_id and cart_id which refers to the product and the cart table. This contains the attribute as count which helps to identify the count of products in cart. By this we can notify the customers if there is any discount on the respective product.

orders_products:

query:

```
-- create orders_products table
CREATE TABLE orders_products (
order_id UUID NOT NULL,
product_id UUID NOT NULL,
quantity INTEGER NOT NULL,
PRIMARY KEY (order_id, product_id),
FOREIGN KEY (order_id) REFERENCES orders(id),
FOREIGN KEY (product_id) REFERENCES products(id)
);
```



```
kl_mom_snowflake_database> CREATE TABLE orders_products (
kl_mom_snowflake_database> order_id UUID NOT NULL,
kl_mom_snowflake_database> product_id UUID NOT NULL,
kl_mom_snowflake_database> quantity INTEGER NOT NULL,
kl_mom_snowflake_database> PRIMARY KEY (order_id, product_id),
kl_mom_snowflake_database> FOREIGN KEY (order_id) REFERENCES orders(id),
kl_mom_snowflake_database> FOREIGN KEY (product_id) REFERENCES products(id)
kl_mom_snowflake_database> );

TABLES:
kl_mom_snowflake_database> SHOW TABLES;
+-----+
| Table |
+-----+
| orders_products |
+-----+

TABLE: "public.orders_products"
  Column   Type          Collation  Nullable |
+-----+-----+-----+-----+
| order_id | uuid         |            | not null |
| product_id | uuid        |            | not null |
| quantity | integer      |            | not null |
+-----+-----+-----+-----+

INDEXES:
+-----+
| Index          |
+-----+
| "orders_products_pkey" PRIMARY KEY, btree (order_id, product_id)
+-----+

RELATIONS:
+-----+
| Relation          |
+-----+
| "orders_products_order_id_fkey" FOREIGN KEY (order_id) REFERENCES orders(id)
| "orders_products_product_id_fkey" FOREIGN KEY (product_id) REFERENCES products(id)
+-----+
```

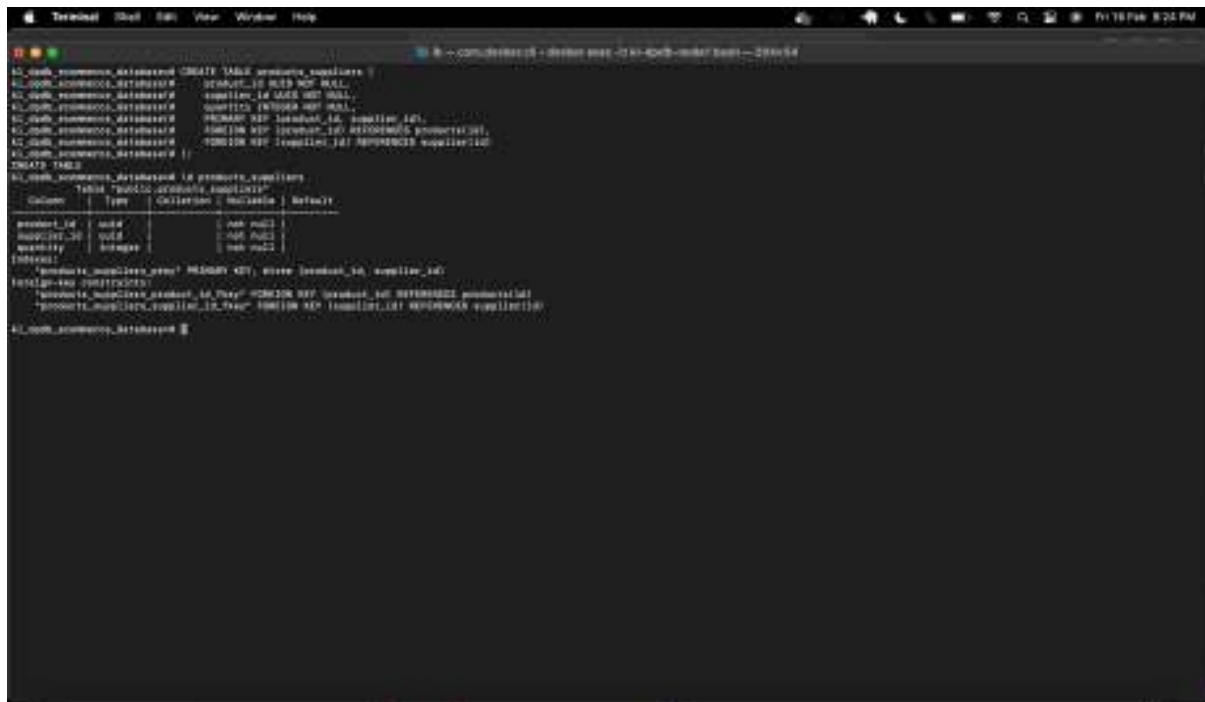
Fig-26: Create and Display the orders_products table and schema of it.

Explanation: The relationship between the orders and the products table is many-to-many. The intermediary link between the orders and the product table forms a relation table. The table has order_id and product_id as primary key and the foreign key, which refers to the order and product table. The table as attribute quantity which stores the quantity of the ordered products by the customer.

products_suppliers table:

query:

```
-- create products_suppliers table
CREATE TABLE products_suppliers (
product_id UUID NOT NULL,
supplier_id UUID NOT NULL,
quantity INTEGER NOT NULL,
PRIMARY KEY (product_id, supplier_id),
FOREIGN KEY (product_id) REFERENCES products(id),
FOREIGN KEY (supplier_id) REFERENCES supplier(id)
);
```



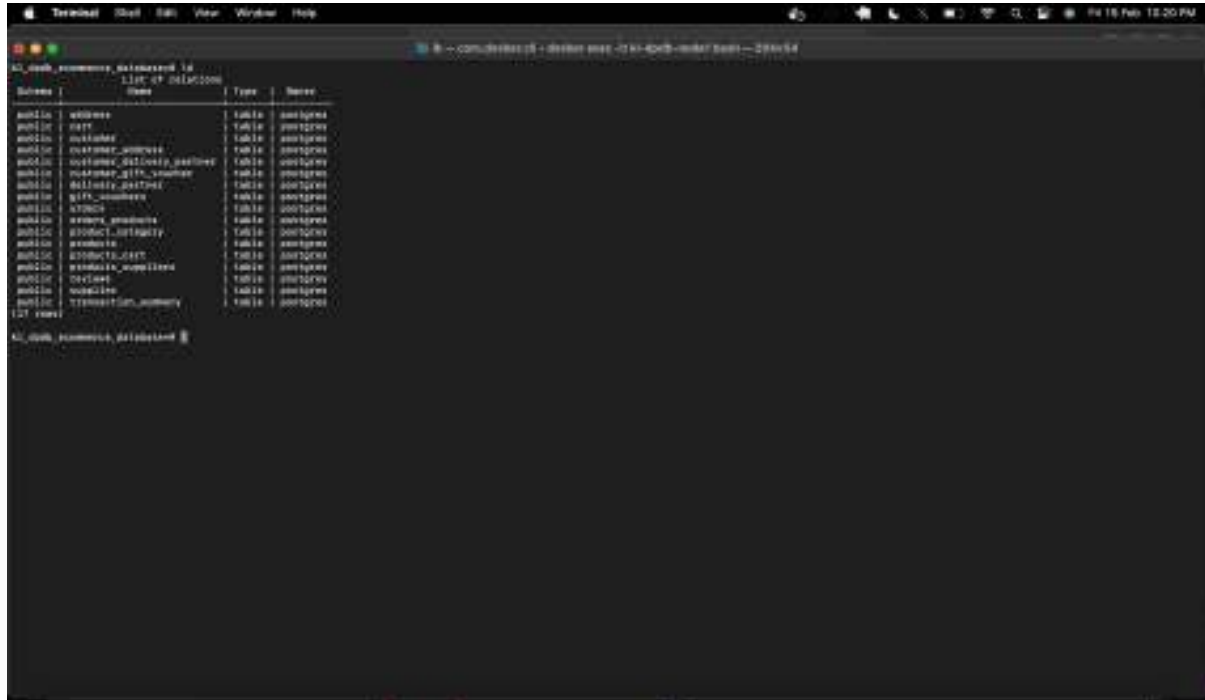
```
kl_mom_sreenivas@postgres:~$ CREATE TABLE products_suppliers (
kl_mom_sreenivas@postgres:~$ product_id UUID NOT NULL
kl_mom_sreenivas@postgres:~$ supplier_id UUID NOT NULL
kl_mom_sreenivas@postgres:~$ quantity INTEGER NOT NULL
kl_mom_sreenivas@postgres:~$ PRIMARY KEY (product_id, supplier_id)
kl_mom_sreenivas@postgres:~$ FOREIGN KEY (product_id) REFERENCES products(id)
kl_mom_sreenivas@postgres:~$ FOREIGN KEY (supplier_id) REFERENCES supplier(id)
kl_mom_sreenivas@postgres:~$ );
CREATE TABLE
kl_mom_sreenivas@postgres:~$ \d products_suppliers
Table "public.products_suppliers"
  Column   | Type          | Collation | Nullable | Default
-----|-----|-----|-----|-----
 product_id | uuid          |           | not null |
supplier_id | uuid          |           | not null |
  quantity  | integer       |           | not null |
Indexes:
  1. "products_suppliers_pkey" PRIMARY KEY, btree (product_id, supplier_id)
Foreign-key constraints:
  "products_suppliers_product_id_fkey" FOREIGN KEY (product_id) REFERENCES products(id)
  "products_suppliers_supplier_id_fkey" FOREIGN KEY (supplier_id) REFERENCES supplier(id)
```

Fig-27: Create and display schema for products_suppliers.

Explanation: The relationship between the products and the suppliers table is many-to-many. The intermediary link between the products and the supplier table forms the relation table. The product_id and the supplier_id are the primary and foreign keys which refer to the product and supplier tables. The table has an attribute quantity which stores the quantity of the products supplied by the supplier. This helps easily to find the products supplied by the supplier.

List of tables created for e-commerce database:

The command used to display the list of tables in a specific database is “\d”. As part of the e-commerce database created 17 tables (11 entities and the 6 relation tables).



```
kl_dadb_ecommerce=# \d
List of schemas
Schema | Name | Type | Encod
-----+-----+-----+-----
public | schema | table | UTF8
public | user | table | UTF8
public | customer | table | UTF8
public | customer_address | table | UTF8
public | customer_address_history | table | UTF8
public | customer_gift_voucher | table | UTF8
public | delivery_order | table | UTF8
public | gift_voucher | table | UTF8
public | inventory | table | UTF8
public | inventory_product | table | UTF8
public | inventory_quantity | table | UTF8
public | product | table | UTF8
public | product_price | table | UTF8
public | product_supplier | table | UTF8
public | review | table | UTF8
public | review_rating | table | UTF8
public | transaction_summary | table | UTF8
17 rows

kl_dadb_ecommerce=#
```

Fig-28: Display the list of tables in kl-dpdb-ecommerce database.

Normalization								
Table	Attribute = Short Forms	Functional Dependencies	Closure	Candidate Key	Prime	Non-Prime	Highest Normal Form	Justification
products	id = A(primary key)	A->B, A->C, A->D, A->E, A->F, A->G, A->H, A->I	A+ = ABCDEFGHI	A	A	B, C, D, E, F, G, H, I	The highest normal form of the products table is "BCNF"	All of the left hand keys of FDs are superkeys, hence the table is in BCNF.
	name = B		B+ = B					
	description = C		C+ = C					
	price = D		D+ = D					
	quantity = E		E+ = E					
	discount = F		F+ = F					
	brandname = G		G+ = G					
	address_id = H		H+ = H					
	product_category_id = I		I+ = I					
customer	id = A(primary key)	A->B, A->C, A->D, A->E, A->F, A->G	A+ = ABCDEFG	A	A	B, C, D, E, F, G	The highest normal form of the customer table is "BCNF"	All of the left hand keys of FDs are superkeys, hence the table is in BCNF.
	first_name = B		B+ = B					
	last_name = C		C+ = C					
	phone_no = D		D+ = D					
	email_id = E		E+ = E					
	dob = F		F+ = F					
	type = G		G+ = G					
cart	id = A	A->B, A->C	A+ = ABC	A	A	B, C	The highest normal form of the cart table is "BCNF"	All of the left hand keys of FDs are superkeys, hence the table is in BCNF.
	customer_id = B		B+ = B					
	quantity = C		C+ = C					
orders	id = A	A->B, A->C, A->D, A->E	A+ = ABCDE	A	A	B, C, D, E	The highest normal form of the orders table is "BCNF"	All of the left hand keys of FDs are superkeys, hence the table is in BCNF.
	status = B		B+ = B					
	order_date = C		C+ = C					
	address_id = D		D+ = D					
	customer_id = E		E+ = E					
transaction_summary	id = A	A->B, A->C, A->D, A->E	A+ = ABCDE	A	A	B, C, D, E	The highest normal form of the transaction_summary table is "BCNF"	All of the left hand keys of FDs are superkeys, hence the table is in BCNF.
	total_amount_paid = B		B+ = B					
	payment_type = C		C+ = C					
	date_of_payment = D		D+ = D					
	order_id = E		E+ = E					
supplier	id = A	A->B, A->C, A->D, A->E, A->F	A+ = ABCDE	A	A	B, C, D, E, F	The highest normal form of the supplier table is "BCNF"	All of the left hand keys of FDs are superkeys, hence the table is in BCNF.
	name = B		B+ = B					
	phone_no = C		C+ = C					
	email = D		D+ = D					
	rating = E		E+ = E					
	address_id = F		F+ = F					
address	id = A	A->B, A->C, A->D, A->E, A->F, A->G	A+ = ABCDEFG	A	A	B, C, D, E, F, G	The highest normal form of the address table is "BCNF"	All of the left hand keys of FDs are superkeys, hence the table is in BCNF.
	flat_no = B		B+ = B					
	street = C		C+ = C					
	city = D		D+ = D					
	state = E		E+ = E					
	country = F		F+ = F					
	zip_code = G		G+ = G					
	id = A		A+ = ABC				The highest normal form of the	All of the left hand keys of FDs are

		Normalization						
Table	Attribute = Short Forms	Functional Dependencies	Closure	Candidate Key	Prime	Non-Prime	Highest Normal Form	Justification
product_category	section = B	A->B, A ->C	B+ = B	A	A	B, C	form of the product_category table is "BCNF"	superkeys, hence the table is in BCNF.
	audience_segment = C		C+ = C					
reviews	id = A	A->B, A ->C, A->D, A->E, A->F	A+ = ABCDEFG	A	A	B, C, D, E, F	The highest normal form of the reviews table is "BCNF"	All of the left hand keys of FDs are superkeys, hence the table is in BCNF.
	rating = B		B+ = B					
	comments = C		C+ = C					
	product_id = D		D+ = D					
	order_id = E		E+ = E					
	customer_id = F		F+ = F					
gift_vouchers	id = A	A->B, A ->C, A->D, A->E	A+ = ABCDE	A	A	B, C, D, E	The highest normal form of the gift_vouchers table is "BCNF"	All of the left hand keys of FDs are superkeys, hence the table is in BCNF.
	name = B		B+ = B					
	number = C		C+ = C					
	amount = D		D+ = D					
	expiry_date = E		E+ = E					
delivery_partner	id = A	A->B, A ->C, A->D, A->E	A+ = ABCDE	A	A	B, C, D, E	The highest normal form of the delivery_partner table is "BCNF"	All of the left hand keys of FDs are superkeys, hence the table is in BCNF.
	name = B		B+ = B					
	phone_no = C		C+ = C					
	email = D		D+ = D					
	order_id = E		E+ = E					
customer_address (relation table)	customer_id = A	AB->C	AB+ = ABC	AB	AB	C	The highest normal form of the customer_address table is "BCNF"	All of the left hand keys of FDs are superkeys, hence the table is in BCNF.
	address_id = B		C+ = C					
	default_address = C							
customer_gift_voucher (relation table)	customer_id = A	AB->C	AB+ = ABC	AB	AB	C	The highest normal form of the customer_gift_voucher table is "BCNF"	All of the left hand keys of FDs are superkeys, hence the table is in BCNF.
	gift_voucher_id = B		C+ = C					
	count = C							
customer_delivery_partner (relation table)	customer_id = A	As the table contains only primary keys there will be no functional dependencies in the table.	NA	AB	AB	NA	The highest normal form of the customer_delivery_part ner table is "BCNF"	All of the left hand keys of FDs are superkeys, hence the table is in BCNF.
	delivery_partner_id = B							
products_cart (relation table)	product_id = A	AB->C	AB+ = ABC	AB	AB	C	The highest normal form of the products_cart table is "BCNF"	All of the left hand keys of FDs are superkeys, hence the table is in BCNF.
	cart_id = B		C+ = C					
	count = C							
orders_products (relation table)	orders_id = A	AB->C	AB+ = ABC	AB	AB	C	The highest normal form of the orders_products table is "BCNF"	All of the left hand keys of FDs are superkeys, hence the table is in BCNF.
	products_id = B		C+ = C					
	quantity = C							
products_suppliers (relation table)	product_id = A	AB->C	AB+ = ABC	AB	AB	C	The highest normal form of the products_suppliers table is "BCNF"	All of the left hand keys of FDs are superkeys, hence the table is in BCNF.
	supplier_id = B		C+ = C					
	quantity = C							

Requirement -3:

There are a number of factors to consider when dealing with databases to ensure optimal performance, including data control, data consistency, the need for ACID features, appropriate fragmentation, etc.

When working with databases, the fragmentation of data is an important technique. Data is kept in an organized manner using fragmentation. We won't be aware of the data's scaling capabilities when we create and save the data in the tables. Fragmentation becomes crucial when we have a circumstance where our memory is full. It is possible to execute fragmentation across the records or columns. Fragmenting data across columns is referred to as vertical fragmentation, and fragmenting data across records is referred to as horizontal fragmentation. The fragmentation can be done on both the records and the columns, called as the hybrid fragmentation.

Completing the horizontal fragmentation on the customer table in accordance with the need. After horizontal fragmentation is complete, the fragmentation's accuracy must be assessed. Completeness, reconstruction, and disjointness can be used to determine whether the fragmentation is valid.

To perform the horizontal fragmentation the table should have number of records. So inserted 1200 records.

query:


```
INSERT INTO customer (id, first_name, last_name, phone_no, email_id, dob, type)
VALUES ('5aec2394-4bb8-4dc2-afbf-590f03cd0414', 'tdLmVeuWfq', 'jHgTMhnFTV', '(787)
510-9892', 'tscep@example.com', '1984-11-17', 'regular');
INSERT INTO customer (id, first_name, last_name, phone_no, email_id, dob, type)
VALUES ('47580a39-1b9d-472b-a2b6-6b1bf46d1b6c', 'sfNcUcAYai', 'LCJqrPCViv', '(463)
662-8591', 'omqwo@example.com', '1964-05-18', 'VIP');
INSERT INTO customer (id, first_name, last_name, phone_no, email_id, dob, type)
VALUES ('c7d39a92-9d22-4ddf-9e0f-3a11860634c2', 'XLLQslUxVc', 'dkNCFcWp0G', '(354)
220-6490', 'dawge@example.com', '1979-12-18', 'platinum');
```

.

.

.

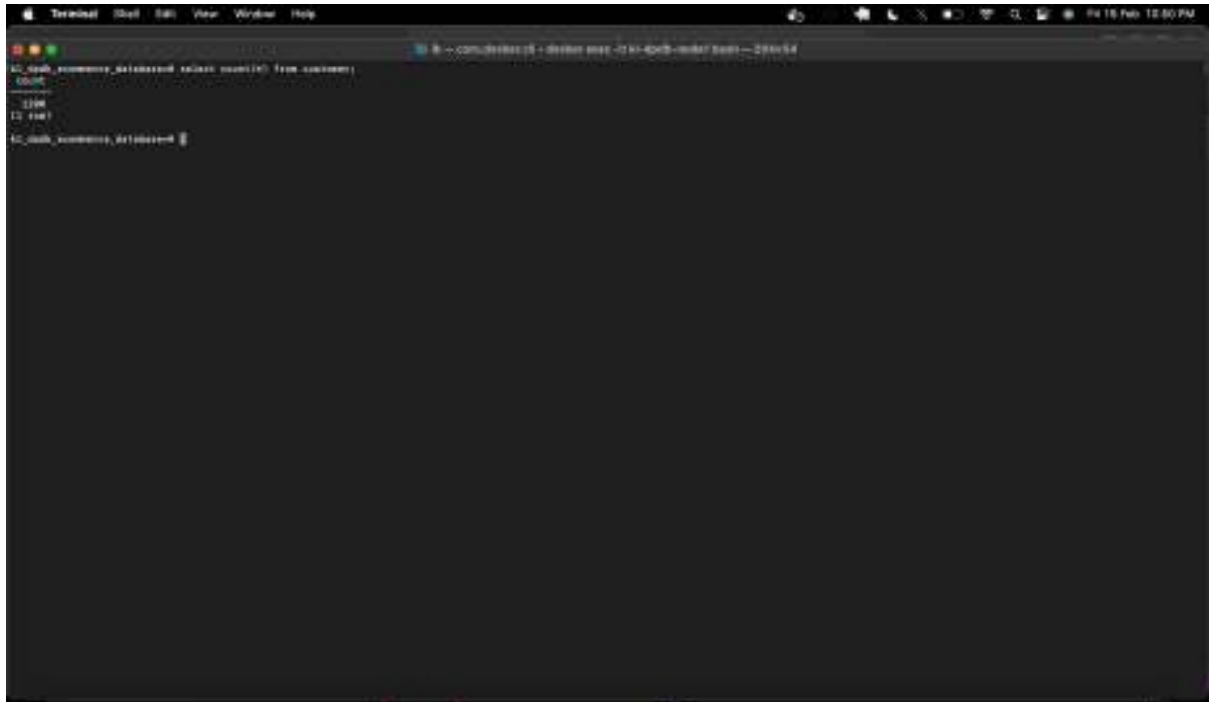
1200 records.



Explanation: Inserted 1200 records to perform the horizontal fragmentation. If we have a greater number of records, we can clearly visualize the performance of implemented horizontal fragmentation.

query:

```
SELECT count(*) from customer;
```

A screenshot of a terminal window with a dark background. The terminal title bar shows 'Terminal' and standard window controls. The prompt is 'root@kali:~#'. The command entered is 'mysql -u root -p'. The prompt changes to 'mysql>'. The SQL query 'SELECT count(*) from customer;' is entered. The output shows '1200' on a new line. The prompt returns to 'mysql>'.

```
root@kali:~# mysql -u root -p
mysql> SELECT count(*) from customer;
1200
mysql>
```

Fig-30: Display the count of records in the customer table.

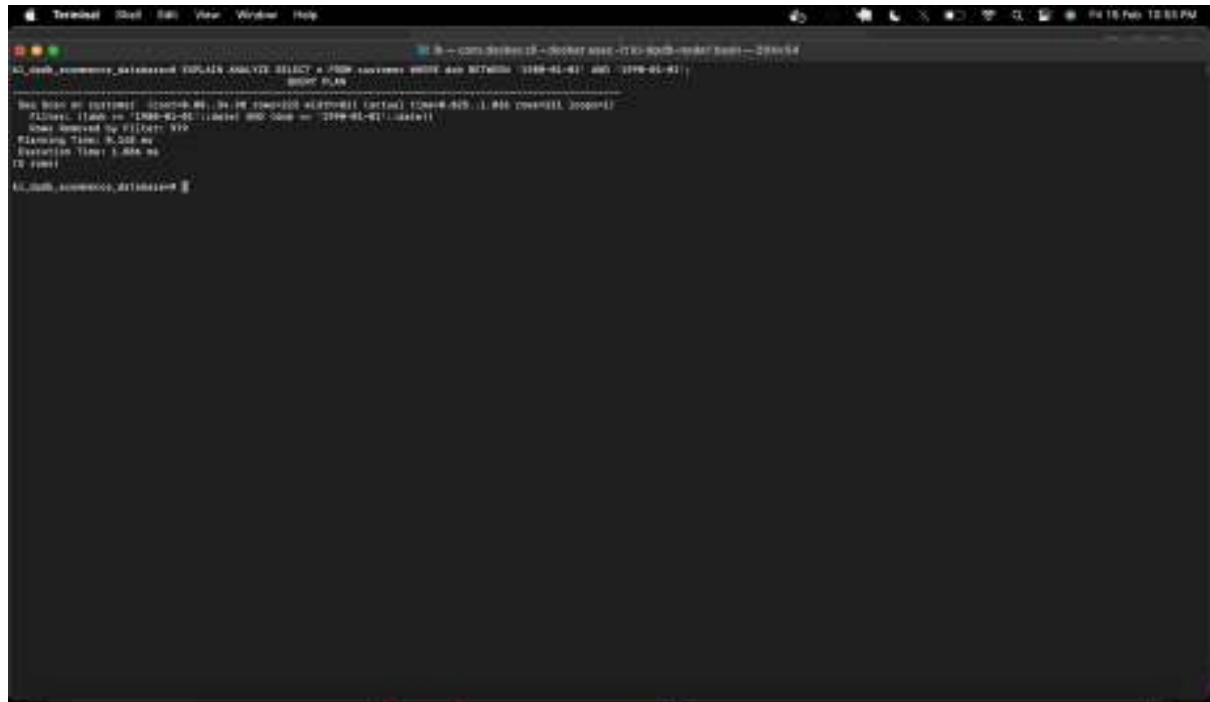
Explanation: There are total 1200 number of records in the customer table.

Fragmentation:

As working only on the single table used **Primary Horizontal Fragmentation**.

query:

```
EXPLAIN ANALYZE SELECT * FROM customer WHERE dob BETWEEN '1980-01-01' AND '1990-01-01';
```



```
al_cool_koonen@al_cool:~$ EXPLAIN ANALYZE SELECT * FROM customer WHERE dob BETWEEN '1980-01-01' AND '1990-01-01';
EXPLAIN
QUERY PLAN
Seq Scan on customer (cost=0.00..34.16 rows=200 estimated actual rows=100..1.000 rows=100000)
Planning Time: 0.148 ms
Execution Time: 1.086 ms
(2 rows)
```

Fig-31: Analyze the performance of customer table.

Explanation: To perform the fragmentation initially we need to analyze the performance of query on the respective table. Observed that, while query is executing it applied the filter based on the given condition i.e., on dob. It took the **planning time of 0.148** milli seconds and **Execution time of 1.086** milli seconds.

query:

```
select version() -- to find the current version of postgres  
create extension pgstattuple -- to create pgstattuple extension
```

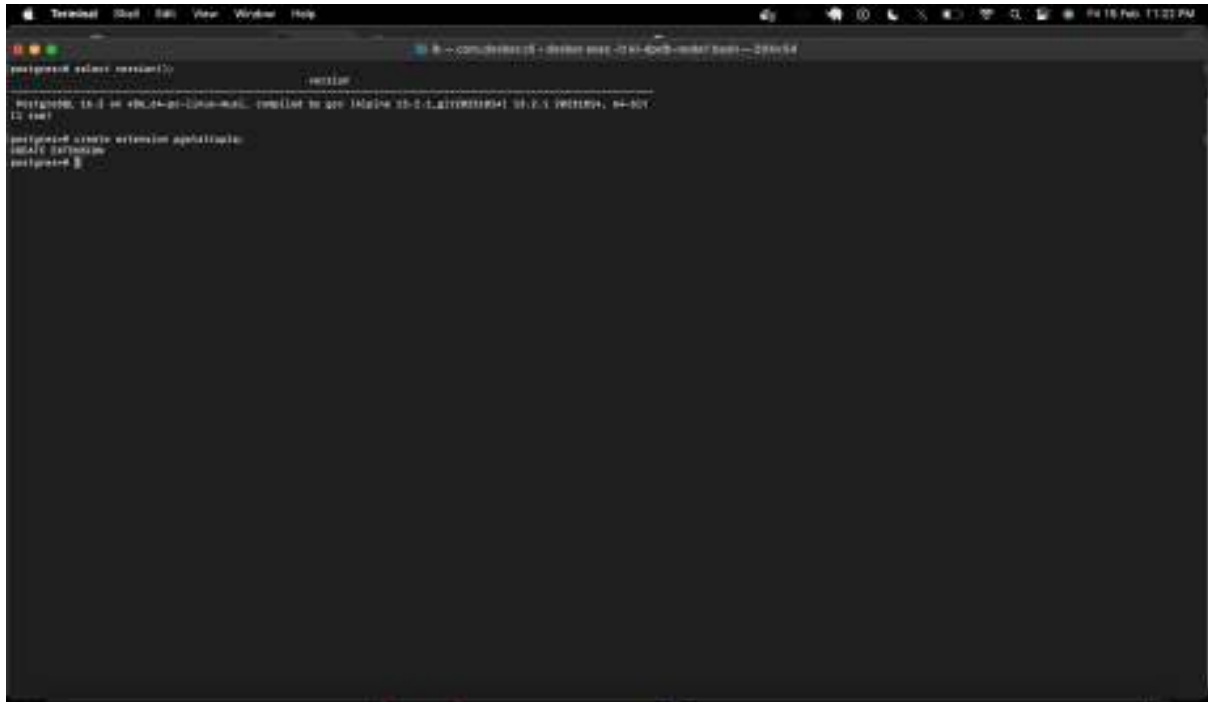
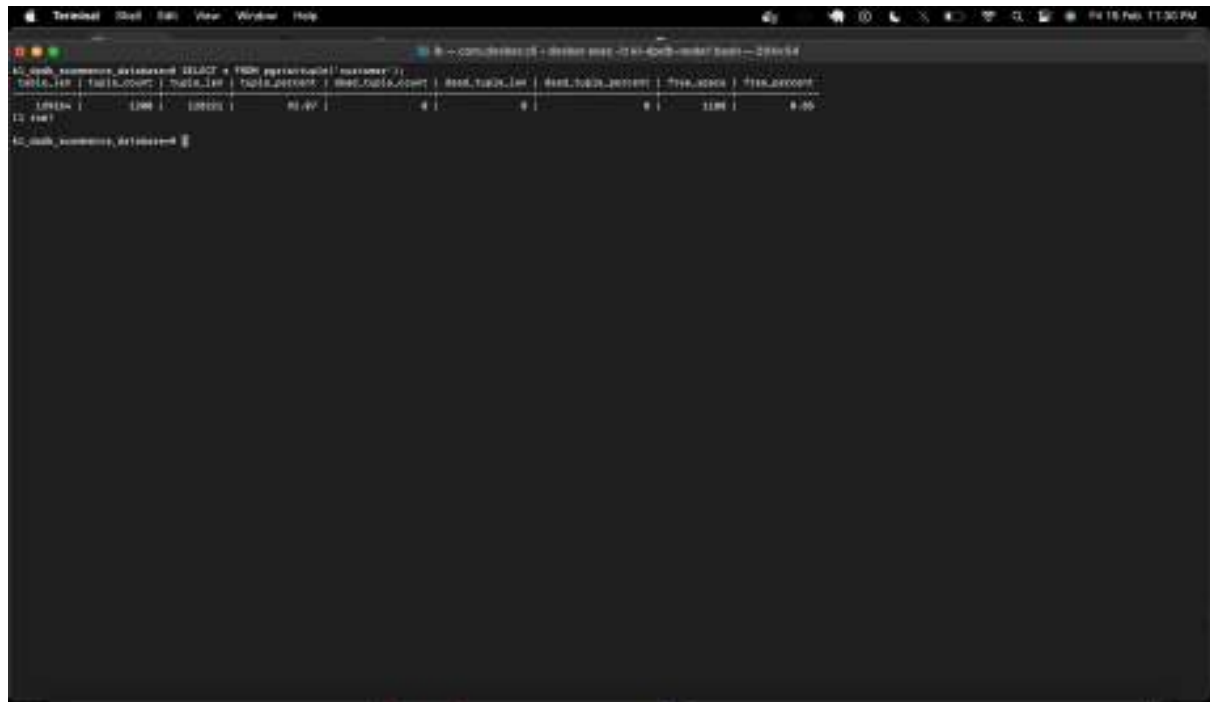
A screenshot of a terminal window with a dark background. The terminal shows the execution of two PostgreSQL commands. The first command, 'select version();', returns the PostgreSQL version '15.3 (Ubuntu 15.3-1ubuntu1)' and the architecture 'aarch64-linux-gnu'. The second command, 'create extension pgstattuple;', returns 'CREATE EXTENSION'. The terminal window has a title bar with standard macOS window controls and a menu bar with options like 'Terminal', 'Shell', 'File', 'View', 'Window', and 'Help'. The system status bar at the top right shows the date and time as '14 Feb 11:21 PM'.

Fig-32: Verify the postgres version and create pgstattuple extension.

Explanation: when we are working on fragmentation, we need to be aware of tuple information. This can be done using “pgstattuple”. This was available from postgres-9.4 and later versions. To make use of pgstattuple we need to create the extension for this, so that we can evaluate the statistics of the table.

query:

```
select * from pgstattuple('customer');
```



The screenshot shows a terminal window with the following output:

```
12_000 customer, diskbased SELECT * FROM pgstattuple('customer')
table_len | table_count | tuple_len | tuple_percent | dead_tuple_count | dead_tuple_len | dead_tuple_percent | free_space | free_percent
-----+-----+-----+-----+-----+-----+-----+-----+-----
139264 | 1200 | 128221 | 92.07 | 0 | 0 | 0 | 1188 | 0.85
```

Fig-33: Get the statistical information on the customer table.

Explanation: using pgstattuple extracted the statistics of customer table tuples. The pgstattuple talks about the

- Total length of the table in the disk i.e., table_len is 139264.
- Total records in the table i.e., tuple_count of 1200.
- Total length utilized by all the records in the table i.e., tuple_length of 128221.
- The percentage of space used by the records in the table from total space i.e., tuple_percent of 92.07.
- The dead tuples are the tuple which need to delete. In the customer table, dead_tuple_count - 0, its length is dead_tuple_length – 0, dead_tuple_percent – 0.
- The total amount of free space i.e., free_space is 1188 and its percent i.e., free_percent is 0.85.

As the dead tuples are 0, which define table has no tuples to delete. As the tables has low amount of free space, much of the space is occupied by the tuples. To have the better space utilization fragmentation is performed (Horizontal Fragmentation) over the tuples.

query:

```
select pg_size_pretty(pg_total_relation_size('customer')) "Table_Size", count(*)  
from customer;
```




Fig-34: Check the space utilized in customer table using pg_size_pretty.

Explanation: The pg_size_pretty talks about the total space occupied by the customer table. The query also helps to find the records in it as well. So, the space occupied by the table is 376KB.

query:

```
CREATE TABLE customer_1950_1970 AS SELECT * FROM customer WHERE dob BETWEEN '1950-01-01' AND '1970-12-31';
CREATE TABLE customer_1971_1990 AS SELECT * FROM customer WHERE dob BETWEEN '1971-01-01' AND '1990-12-31';
CREATE TABLE customer_1991_2000 AS SELECT * FROM customer WHERE dob BETWEEN '1991-01-01' AND '2000-12-31';
```

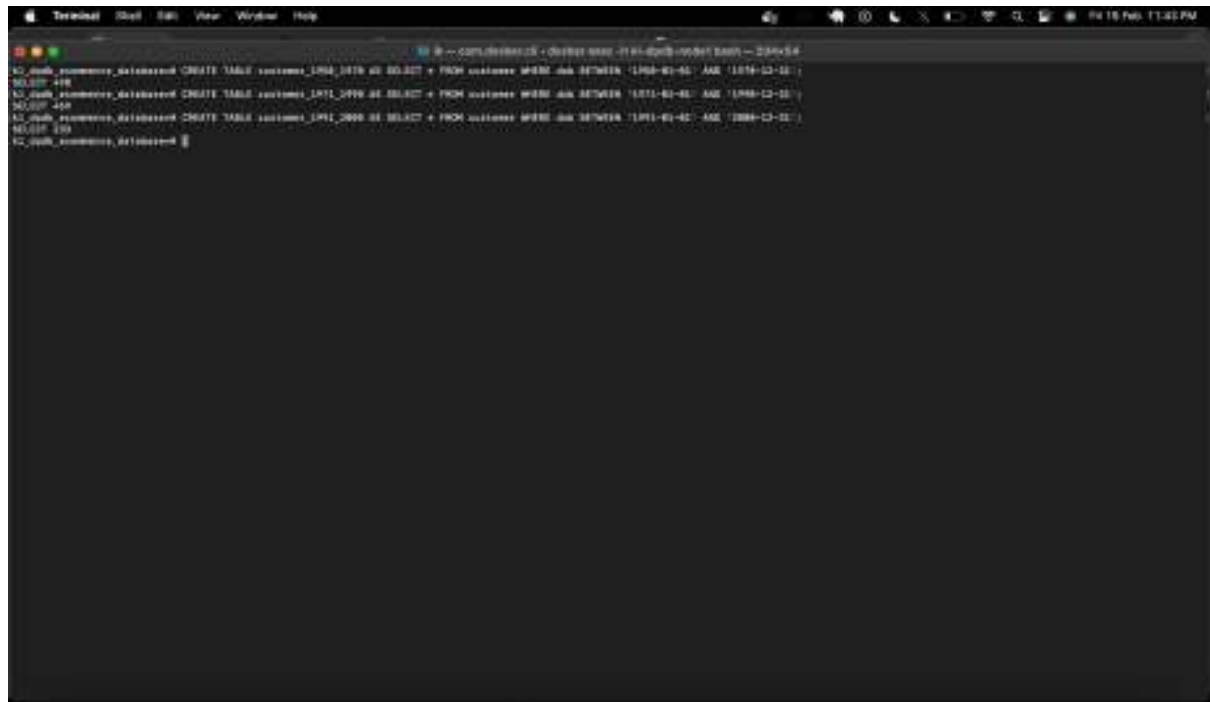
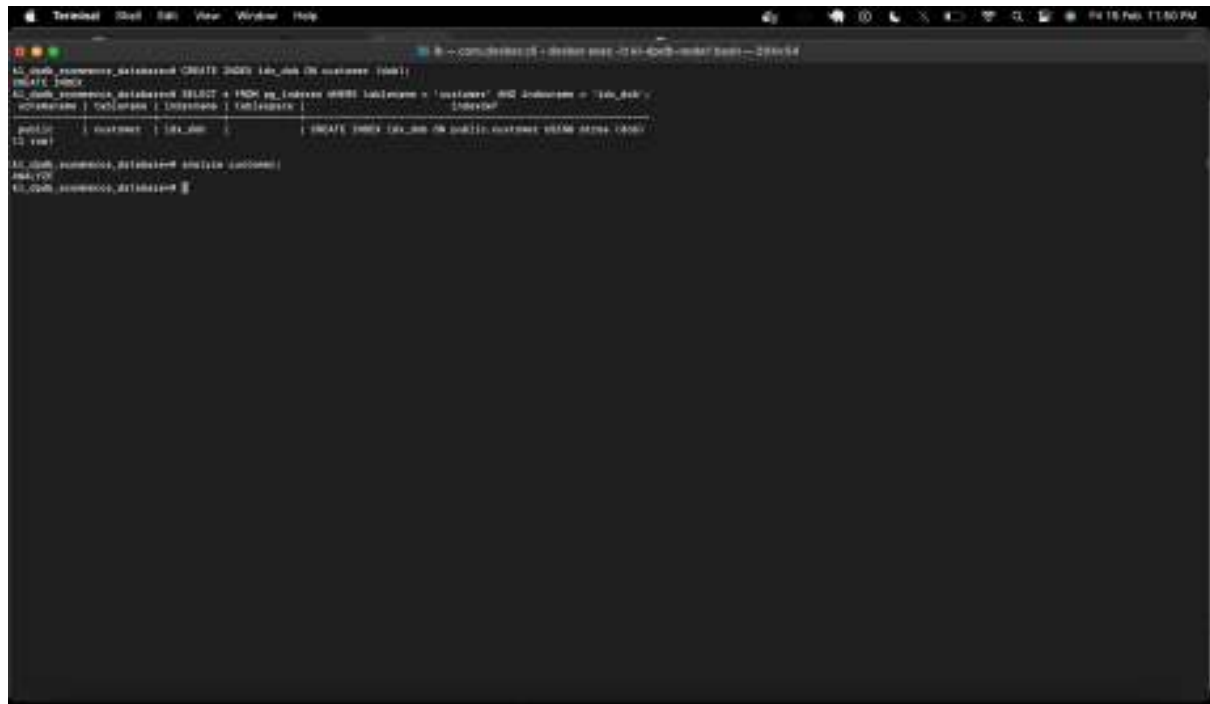
A screenshot of a terminal window with a dark background. The terminal shows three SQL commands being executed in a database. The first command creates a table 'customer_1950_1970' from the 'customer' table, filtering for dates between '1950-01-01' and '1970-12-31'. The output shows 498 records. The second command creates 'customer_1971_1990' with a date range of '1971-01-01' to '1990-12-31', resulting in 469 records. The third command creates 'customer_1991_2000' with a date range of '1991-01-01' to '2000-12-31', resulting in 233 records. The terminal window has a title bar with 'Terminal' and various icons. The top of the terminal shows the shell prompt and the database name 'customer_1950_1970'.

Fig-35: Create fragmented tables from customer table.

Explanation: As we discussed above, the table has low amount of free space, performed fragmentation over the tuples. The whole tuples have classified into 3 fragments i.e., **customer_1950_1970, customer_1971_1990, customer_1991_2000** with **498, 469 and 233 records**. The fragmentation is done based on min and max values of the dob.

query:

```
CREATE INDEX idx_dob ON customer (dob);  
SELECT * FROM pg_indexes WHERE tablename = 'customer' AND indexname = 'idx_dob'
```



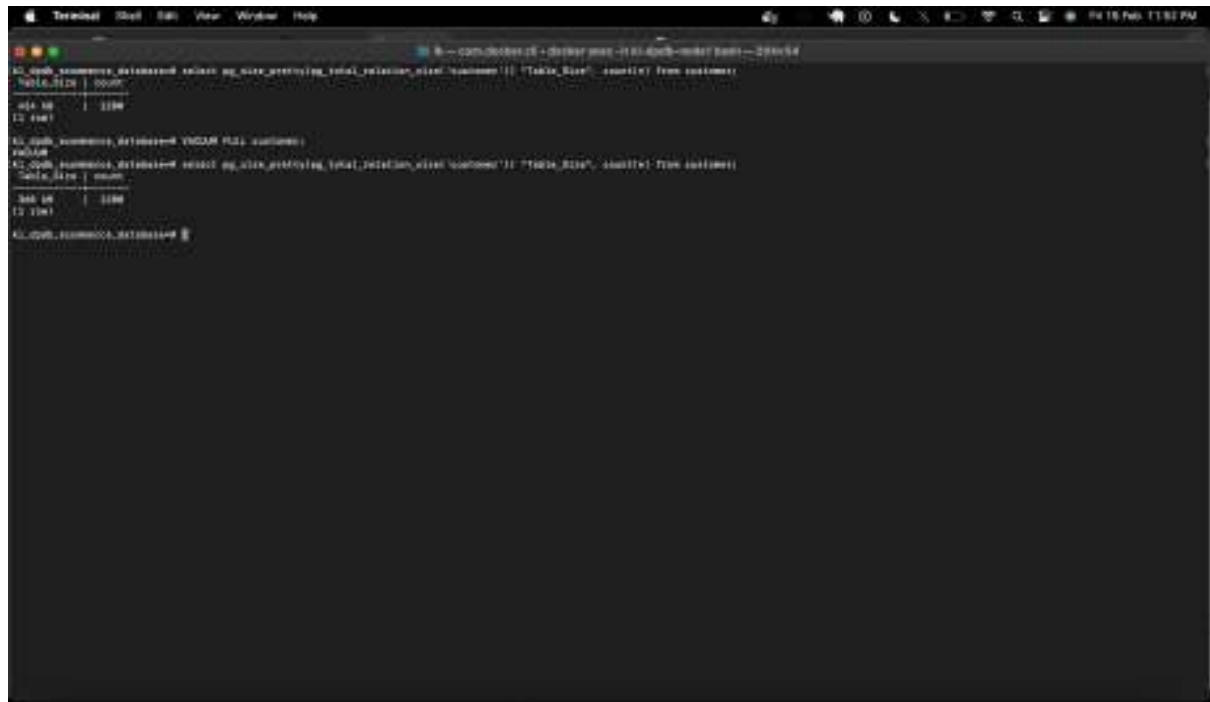
```
1: postgres@postgres:~$ CREATE INDEX idx_dob ON customer (dob);  
CREATE INDEX  
2: postgres@postgres:~$ SELECT * FROM pg_indexes WHERE tablename = 'customer' AND indexname = 'idx_dob';  
tablename | idxname | indexname |  
-----  
public    | customer | idx_dob    | CREATE INDEX idx_dob ON public.customer USING btree (dob)  
1 row  
3: postgres@postgres:~$ ANALYZE customer;  
ANALYZE  
4: postgres@postgres:~$
```

Fig-36: Create the index on the specific column(dob).

Explanation: As the fragmentation is performed using the dob, created the index for the dob column with idx_dob.

query:

```
select pg_size_pretty(pg_total_relation_size('customer')) "Table_Size", count(*)
from customer;
VACCU full customer;
select pg_size_pretty(pg_total_relation_size('customer')) "Table_Size", count(*)
from customer;
```

A screenshot of a PostgreSQL terminal window. The terminal shows the execution of three SQL queries. The first query checks the size of the 'customer' table before defragmentation, returning '360KB' and a count of 1. The second query runs 'VACCU full customer;'. The third query checks the size of the 'customer' table after defragmentation, returning '360KB' and a count of 1. The terminal window has a dark background and a title bar at the top.

```
postgres=# select pg_size_pretty(pg_total_relation_size('customer')) "Table_Size", count(*)
from customer;
Table_Size | count
-----
360KB      | 1
(1 row)

postgres=# VACCU full customer;
VACCU
postgres=# select pg_size_pretty(pg_total_relation_size('customer')) "Table_Size", count(*) from customer;
Table_Size | count
-----
360KB      | 1
(1 row)

postgres=#
```

Fig-37: Space after fragmentation.

Explanation: Once the fragmentation is performed and the created the index on the specific column, the size of the table doesn't decrease. It's still increased as we created index on the dob attribute. Ideally VACCU is used to perform the defragmentation or the optimization. Here also to optimize the size after fragmentation performed VACCU on the customer table by that we can observe the size of the table is 360KB. The results look better after the fragmentation.

query:

```
EXPLAIN ANALYZE SELECT * FROM customer WHERE dob BETWEEN '1980-01-01' AND '1990-01-01';
```

```
EXPLAIN ANALYZE SELECT * FROM customer_1950_1970 WHERE dob BETWEEN '1980-01-01' AND '1990-01-01';
```

```
EXPLAIN ANALYZE SELECT * FROM customer_1971_1990 WHERE dob BETWEEN '1980-01-01' AND '1990-01-01';
```

```
EXPLAIN ANALYZE SELECT * FROM customer_1991_2000 WHERE dob BETWEEN '1980-01-01' AND '1990-01-01';
```

```
Terminal [Sat 17 Feb 12:00 AM]
SQL> EXPLAIN ANALYZE SELECT * FROM customer WHERE dob BETWEEN '1980-01-01' AND '1990-01-01';
QUERY PLAN
--
Bitmap Heap Scan on customer  (cost=0.00..10.00 rows=234 width=41) (actual time=0.343 ms) (rows=234)
  Bitmap Cond: (dob >= '1980-01-01'::date) AND (dob <= '1990-01-01'::date)
  Heap Blocks: exact=0
  --
  Bitmap Index Scan on dob_idx  (cost=0.00..18.42 rows=234 width=41) (actual time=0.000 ms) (rows=234)
    Index Cond: (dob >= '1980-01-01'::date) AND (dob <= '1990-01-01'::date)
Planning Time: 0.321 ms
Execution Time: 0.343 ms
(7 rows)

SQL> EXPLAIN ANALYZE SELECT * FROM customer_1950_1970 WHERE dob BETWEEN '1980-01-01' AND '1990-01-01';
QUERY PLAN
--
Seq Scan on customer_1950_1970  (cost=0.00..18.42 rows=234 width=41) (actual time=0.321 ms) (rows=234)
  Filter: (dob >= '1980-01-01'::date) AND (dob <= '1990-01-01'::date)
  Rows Removed by Filter: 0
Planning Time: 0.034 ms
Execution Time: 0.321 ms
(9 rows)

SQL> EXPLAIN ANALYZE SELECT * FROM customer_1971_1990 WHERE dob BETWEEN '1980-01-01' AND '1990-01-01';
QUERY PLAN
--
Seq Scan on customer_1971_1990  (cost=0.00..18.42 rows=234 width=41) (actual time=0.308 ms) (rows=234)
  Filter: (dob >= '1980-01-01'::date) AND (dob <= '1990-01-01'::date)
  Rows Removed by Filter: 0
Planning Time: 0.011 ms
Execution Time: 0.308 ms
(9 rows)

SQL> EXPLAIN ANALYZE SELECT * FROM customer_1991_2000 WHERE dob BETWEEN '1980-01-01' AND '1990-01-01';
QUERY PLAN
--
Seq Scan on customer_1991_2000  (cost=0.00..18.42 rows=234 width=41) (actual time=0.057 ms) (rows=234)
  Filter: (dob >= '1980-01-01'::date) AND (dob <= '1990-01-01'::date)
  Rows Removed by Filter: 0
Planning Time: 0.035 ms
Execution Time: 0.057 ms
(9 rows)

SQL>
```

Fig-38: Overall performance

Explanation: The above screenshot talks about the overall performance of the table when we perform querying for multiple times. As we are querying for multiple times performance is getting changed with good efficiency and observed the data is retrieving faster after fragmentation. The space also allocated properly, so that we can insert more data now. The Fragmentation plays the vital role in the space and performance of the data.

Correctness of Fragmentation:

After fragmentation is done, we need to evaluate the correctness of the fragmentation. This can be verified using completeness, reconstruction and disjointness.

Completeness of the fragmentation: The completeness of the data talks about the schema, data distribution and records count of the fragmented tables.

query:

Reviewing the fragment tables-

```
\d customer_1950_1970
```

```
\d customer_1971_1990
```

```
\d customer_1991_2000
```

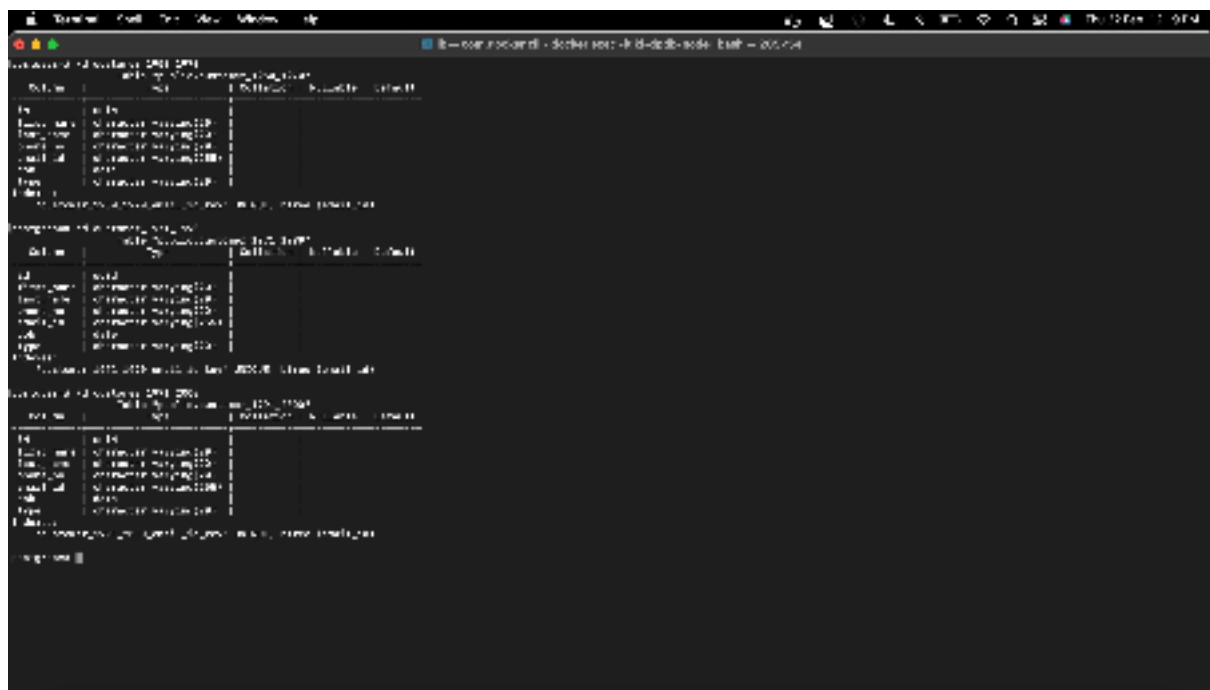


Fig-39: Fragmented tables schema.

Explanation: All the fragmented tables have the same schema which satisfies the data consistency. In the above screen shot all the fragmented tables have same structure.

query:

```
SELECT MIN(dob), MAX(dob) FROM customer_1950_1970;  
SELECT MIN(dob), MAX(dob) FROM customer_1971_1990;  
SELECT MIN(dob), MAX(dob) FROM customer_1991_2000;
```


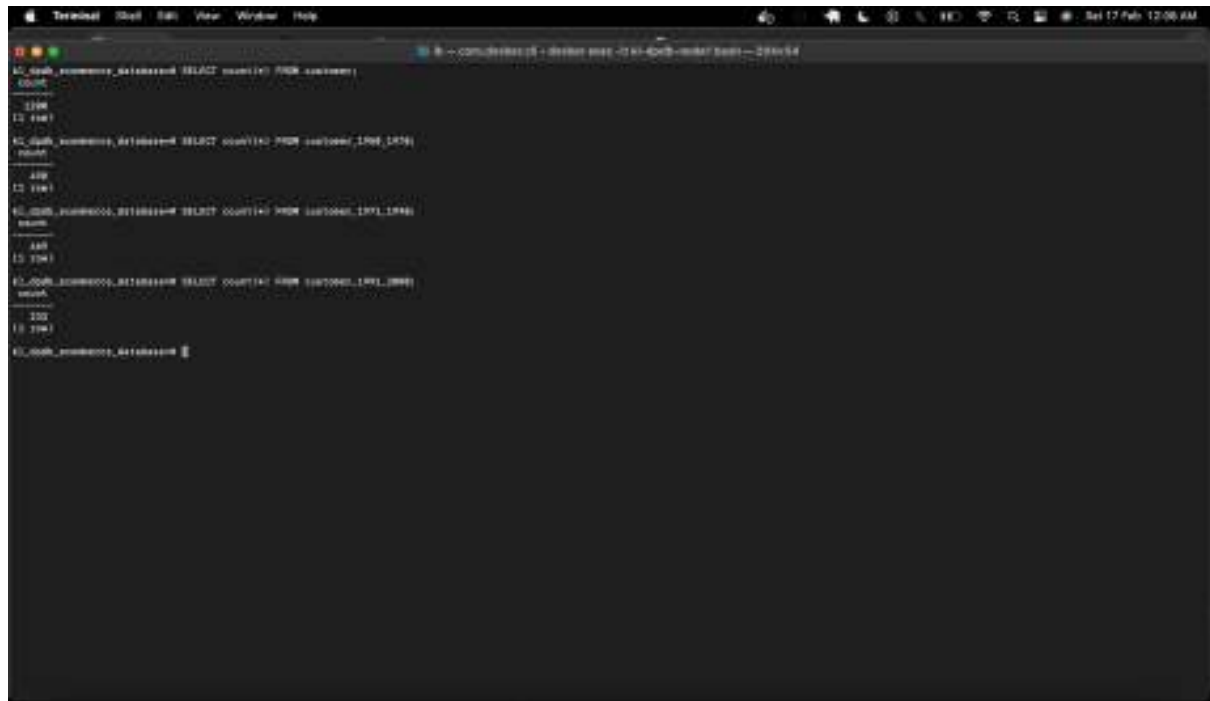


Fig-40: Data Distribution on fragmented tables.

Explanation: The total data in the customer table is distributed in to fragmented tables based on dob. Based on the min and max of value of the fragmented data observed that the data is distributed in the organized way.

query:

```
SELECT count(*) FROM customer;  
SELECT count(*) FROM customer_1950_1970;  
SELECT count(*) FROM customer_1971_1990;  
SELECT count(*) FROM customer_1991_2000;
```



The screenshot shows a terminal window with the following SQL queries and their results:

```
62_csh@ecampus_sandbox:~$ SELECT count(*) FROM customer;  
count  
-----  
1200  
(1 row)  
  
62_csh@ecampus_sandbox:~$ SELECT count(*) FROM customer_1950_1970;  
count  
-----  
498  
(1 row)  
  
62_csh@ecampus_sandbox:~$ SELECT count(*) FROM customer_1971_1990;  
count  
-----  
469  
(1 row)  
  
62_csh@ecampus_sandbox:~$ SELECT count(*) FROM customer_1991_2000;  
count  
-----  
233  
(1 row)  
  
62_csh@ecampus_sandbox:~$
```

Fig-41: Count of records in fragmented tables.

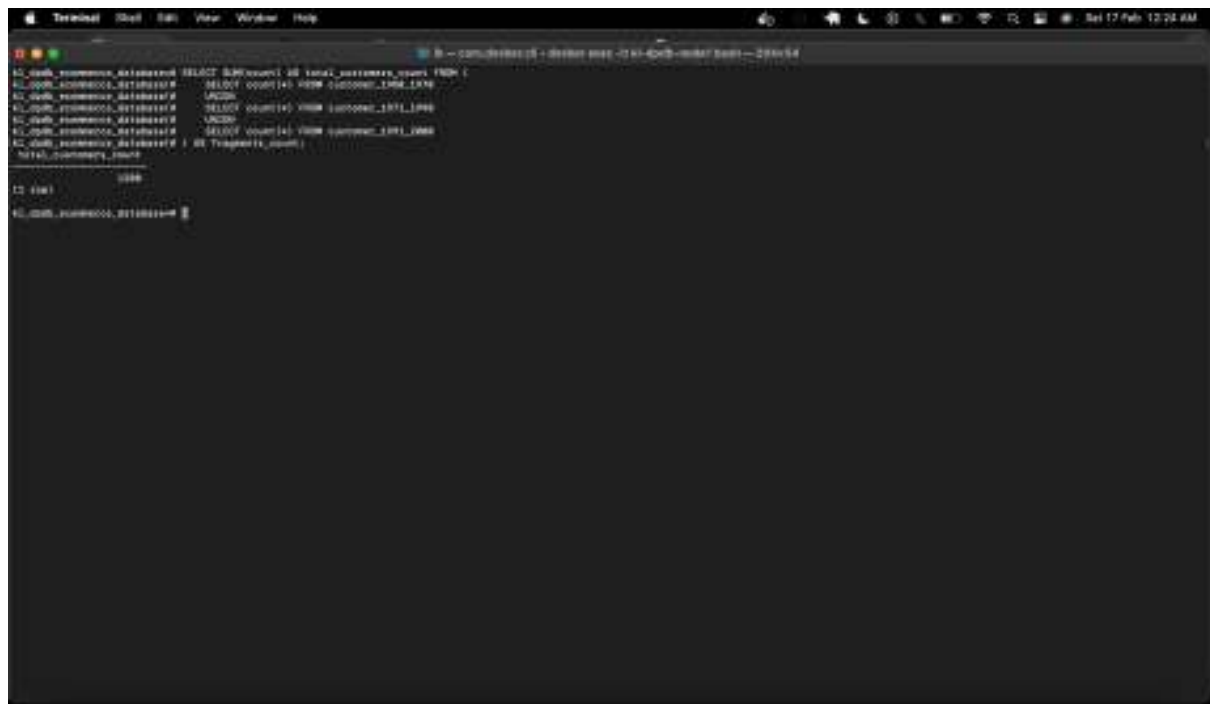
Explanation: The total count of all the fragmented tables is equal to the count of records in the customer table. The customer table has 1200 records, where as customer_1950_1970 has 498, customer_1971_1990 has 469 and customer_1991_2000 has 233 which satisfies the count of records.

As schema of the table, Data distribution and count of records in the fragmented tables says that the Fragmentation of the customer table achieves the completeness.

Reconstruction: When the original table is fragmented in to multiple tables, the union of all the fragmented table defines the original table.

query:

```
-- Reconstruction:
SELECT SUM(count) AS total_customers_count FROM (
    SELECT count(*) FROM customer_1950_1970
    UNION
    SELECT count(*) FROM customer_1971_1990
    UNION
    SELECT count(*) FROM customer_1991_2000
) AS fragments_count;
```



```
Terminal (Shell) View Window Help
-- com.mysql.jdbc.Driver - jdbc:mysql://127.0.0.1:3306/employees - 2016/12/17 12:24 AM

mysql> SELECT SUM(count) AS total_customers_count FROM (
mysql>     SELECT count(*) FROM customer_1950_1970
mysql>     UNION
mysql>     SELECT count(*) FROM customer_1971_1990
mysql>     UNION
mysql>     SELECT count(*) FROM customer_1991_2000
mysql> ) AS fragments_count;
+-----+
| total_customers_count |
+-----+
| 12                |
+-----+
```

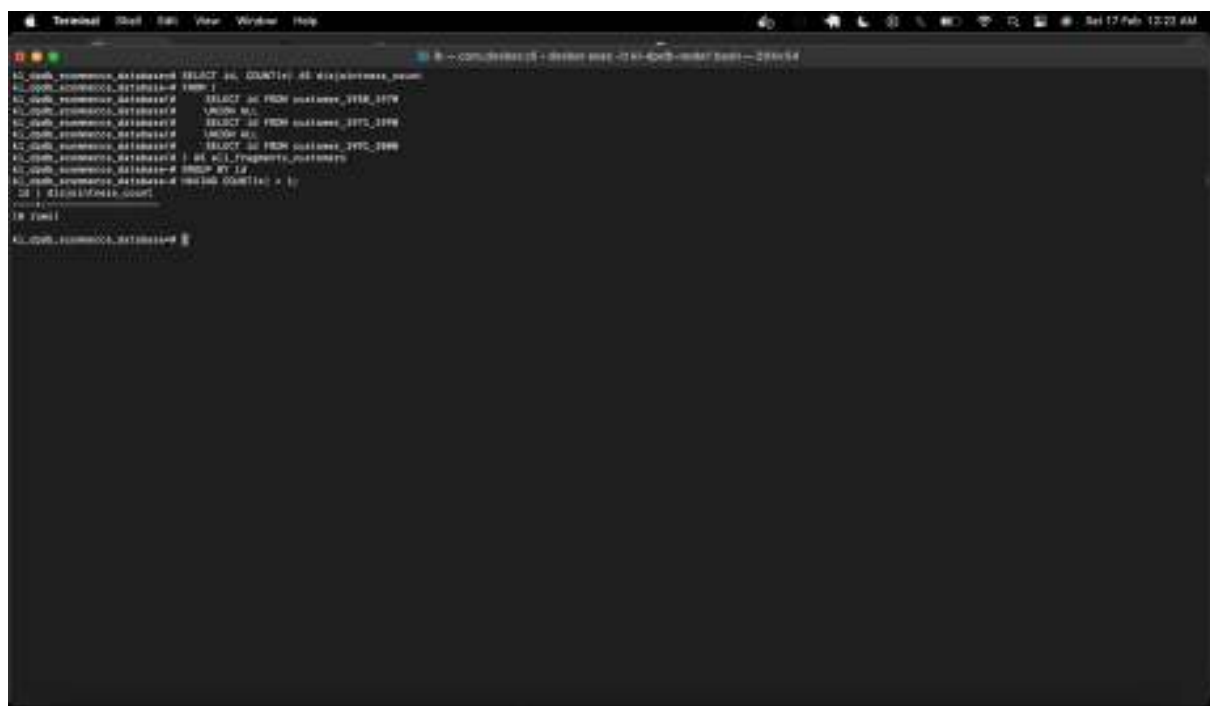
Fig-42: Reconstruction of fragmented tables.

Explanation: The union of all the fragmented tables count is equal to the count of the customer table. By this it achieves the reconstruction.

Disjointness: The disjointness talks about the duplicate of data. There shouldn't be any duplicate data after the fragmentation.

query:

```
-- Disjointness
SELECT id, COUNT(*) AS disjointness_count
FROM (
    SELECT id FROM customer_1950_1970
    UNION ALL
    SELECT id FROM customer_1971_1990
    UNION ALL
    SELECT id FROM customer_1991_2000
) AS all_fragments_customers
GROUP BY id
HAVING COUNT(*) > 1;
```



```
kl_@db: ~$ mysql -u root -p --silent --skip-grant-tables --database=kl_
mysql> SELECT id, COUNT(*) AS disjointness_count
FROM (
    SELECT id FROM customer_1950_1970
    UNION ALL
    SELECT id FROM customer_1971_1990
    UNION ALL
    SELECT id FROM customer_1991_2000
) AS all_fragments_customers
GROUP BY id
HAVING COUNT(*) > 1;
+----+
0 rows in set (0.00 sec)
```

Fig-43: Duplicate data in the fragmented tables.

Explanation: The duplicate records observed in the fragmented tables is 0. By this it says that the fragmented tables achieve the disjointness as well.

Requirement-4:

To perform the concurrency control strategy, PostgreSQL provides a few techniques such as optimistic, pessimistic and MVCC(multi-version concurrency control). In that the MVCC is the most advanced technique to perform concurrent actions securely in the database.

Implementing the MVCC helps to achieve data consistency, it also helps to perform the reading the data and writing the data into the database smoothly without any conflicts.

Use case 1:

The product should only be assigned to one consumer at a time if two distinct customers are attempting to access it with quantity 1. This indicates that the data and the concurrent actions are consistent.

Available Customers:

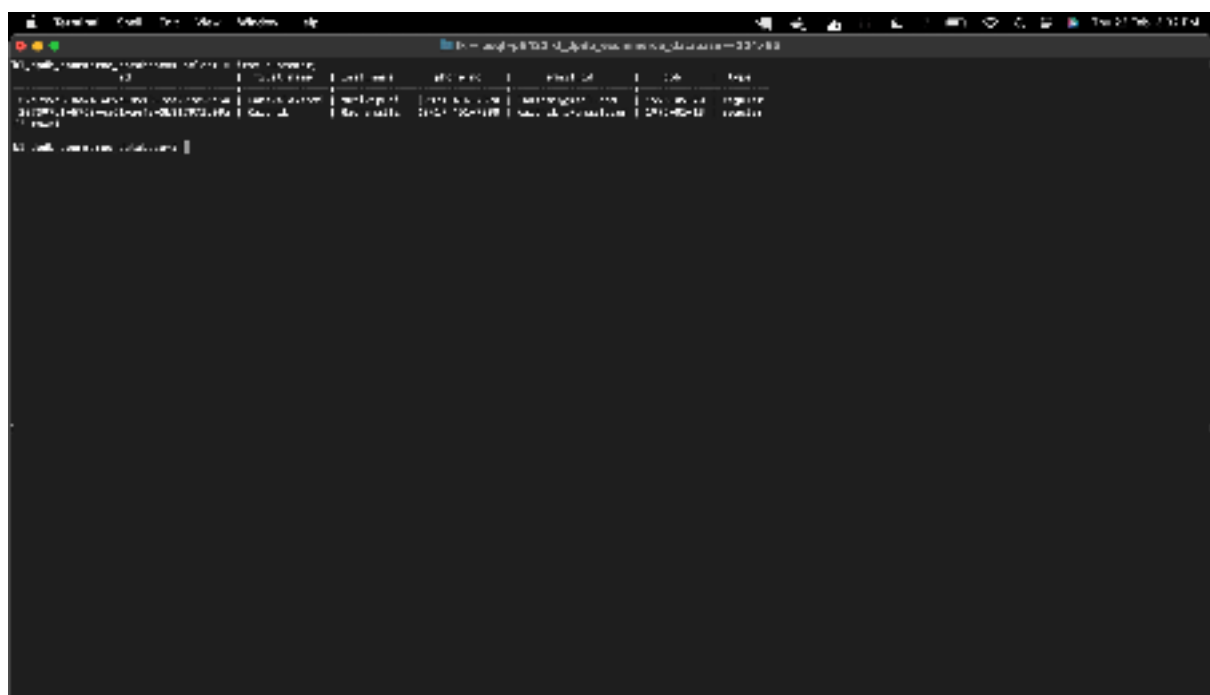
query:

```
select * from customer;
```

available customers-

52d999b8-de44-489d-8913-7ce999e26c5a

207397bf-5759-4a51-a44a-3b317971b09a



The screenshot shows a PostgreSQL database window with the 'customer' table selected. The table has two columns: 'id' and 'name'. The first row has the ID '52d999b8-de44-489d-8913-7ce999e26c5a' and the name 'customer'. The second row has the ID '207397bf-5759-4a51-a44a-3b317971b09a' and the name 'customer'.

id	name
52d999b8-de44-489d-8913-7ce999e26c5a	customer
207397bf-5759-4a51-a44a-3b317971b09a	customer

Fig-44: List of customers.

Explanation: Displayed the list of customers in the customer table. There are two users with id 52d999b8-de44-489d-8913-7ce999e26c5a and 207397bf-5759-4a51-a44a-3b317971b09a.

Available Products:

query:

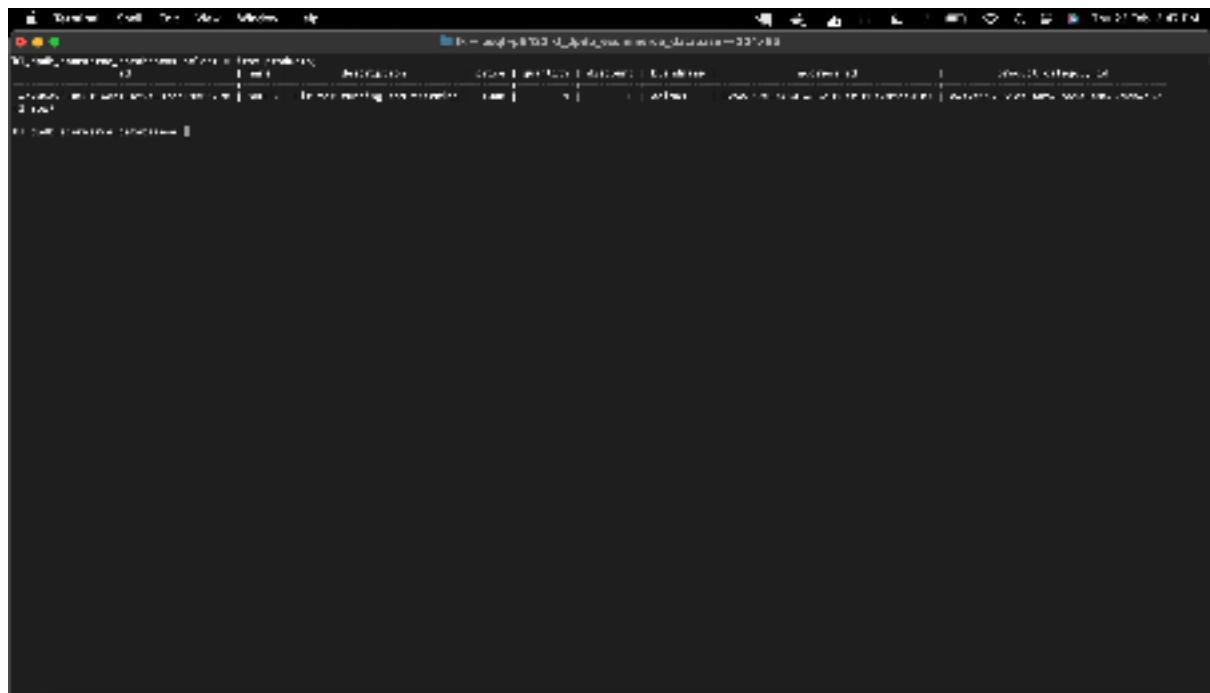
```
select * from products;
```

available products :

4d205a90-053f-4db1-a22a-37cc18355798

name - SHOES

quantity - 1



The screenshot shows a web browser window with a table of products. The table has columns for product ID, name, description, price, and quantity. The first row shows a product with ID 4d205a90-053f-4db1-a22a-37cc18355798, named 'SHOES', with a quantity of 1.

product_id	name	description	price	quantity
4d205a90-053f-4db1-a22a-37cc18355798	SHOES			1

Fig-45: List of products.

Explanation: Displayed the list of available products and its details. Here the product_id 4d205a90-053f-4db1-a22a-37cc18355798 has quantity-1.

PL/pgSQL Function to handle the concurrent actions:

query:

```
CREATE OR REPLACE FUNCTION place_order(
    p_product_id UUID,
    p_customer_id UUID,
    p_quantity INTEGER,
    p_delivery_partner_name VARCHAR(50),
    p_delivery_partner_phone_no VARCHAR(20),
    p_delivery_partner_email VARCHAR(255)
) RETURNS VOID AS $$
DECLARE
    v_order_id UUID;
    v_delivery_partner_id UUID;
    v_product_quantity INTEGER;
BEGIN
    BEGIN
        -- Check if the product is available
        SELECT quantity INTO v_product_quantity FROM products WHERE id =
p_product_id FOR UPDATE;
        -- step-1: Check if the requested quantity is available
        IF v_product_quantity < p_quantity THEN
            RAISE EXCEPTION 'Error in place_order: Product not available. Please
try again later.';
        END IF;
        -- step 2: Insert into orders table
        INSERT INTO orders (status, order_date, address_id, customer_id)
        VALUES ('Processing', CURRENT_DATE, (SELECT address_id FROM
customer_address WHERE customer_id = p_customer_id AND default_address = true),
p_customer_id)
        RETURNING id INTO v_order_id;
        -- Step 3: Update products table
        UPDATE products SET quantity = quantity - p_quantity WHERE id =
p_product_id;
        -- Step 4: Insert into transaction_summary table
        INSERT INTO transaction_summary (total_amount_paid, payment_type,
date_of_payment, order_id)
        VALUES ((SELECT price * p_quantity FROM products WHERE id = p_product_id),
'Credit Card', CURRENT_DATE, v_order_id);
        -- Step 5: Insert into orders_products table
        INSERT INTO orders_products (order_id, product_id, quantity) VALUES
(v_order_id, p_product_id, p_quantity);
        -- Step 6: Insert into delivery_partner table
        INSERT INTO delivery_partner (name, phone_no, email, order_id)
        VALUES (p_delivery_partner_name, p_delivery_partner_phone_no,
p_delivery_partner_email, v_order_id)
        RETURNING id INTO v_delivery_partner_id;
        -- Step 7: Insert into customer_delivery_partner table
        INSERT INTO customer_delivery_partner (customer_id, delivery_partner_id)
VALUES (p_customer_id, v_delivery_partner_id);
```


The function also includes the error handling, which raises the exception and roll backs the data in case of any issue during the transaction.

The above PL/pgSQL internally handles the MVCC, The MVCC creates the snapshot while performing the transactions(one transaction does not affect the other transaction)and provides the locking mechanism to handle the concurrent transactions, dead lock situations, and inconsistent transactions. It majorly follow “ The WRITER never blocks the READER” and “THE READER never blocks the WRITER”.

Python Script to perform concurrent orders by the customer:

query:

```
import psycopg2
from concurrent.futures import ThreadPoolExecutor

# connection details
host = "localhost"
port = 5432
database = "kl_dpdb_ecommerce_database"
user = "postgres"
password = "spaceman1236"

def place_order(product_id, customer_id, quantity, delivery_partner_name, phone_no,
email):
    try:
        place_order_connection = psycopg2.connect(
            host=host,
            port=port,
            database=database,
            user=user,
            password=password
        )
        place_order_cursor = place_order_connection.cursor()
        place_order_cursor.execute(
            """
            SELECT place_order(
                %s, %s, %s, %s, %s, %s
            );
            """
            ,
            (product_id, customer_id, quantity, delivery_partner_name, phone_no,
email)
        )
        place_order_connection.commit()
        print(f"order Placed for {customer_id}!")

    except psycopg2.Error as e:
        error_message = str(e)
        if "product not available" in error_message.lower():
            print(f"Product not available for {customer_id}! please try again")
        else:
            print(f"Error connecting to PostgreSQL: {e}")
```



```

finally:
    if place_order_cursor:
        place_order_cursor.close()
    if place_order_connection:
        place_order_connection.close()

if __name__ == "__main__":
    # customer-1
    customer_1_details = ('4d205a90-053f-4db1-a22a-37cc18355798', '52d999b8-de44-489d-8913-7ce999e26c5a', 1, 'robin', '(817) 777-4089', 'robin@example.com')
    # customer-2
    customer_2_details = ('4d205a90-053f-4db1-a22a-37cc18355798', '207397bf-5759-4a51-a44a-3b317971b09a', 1, 'josey', '(978) 717-4389', 'josey1@example.com')

    with ThreadPoolExecutor(max_workers=2) as executor:
        # Trigger the function concurrently with different user inputs
        executor.submit(place_order, *customer_1_details)
        executor.submit(place_order, *customer_2_details)

```

```

$ python3 script.py
[{"id": "4d205a90-053f-4db1-a22a-37cc18355798", "email": "robin@example.com", "phone": "(817) 777-4089", "status": "PENDING"}, {"id": "207397bf-5759-4a51-a44a-3b317971b09a", "email": "josey1@example.com", "phone": "(978) 717-4389", "status": "PENDING"}]

```

Fig-47: Script to execute concurrent orders.

query:

python kl_dpdb_req-4.py


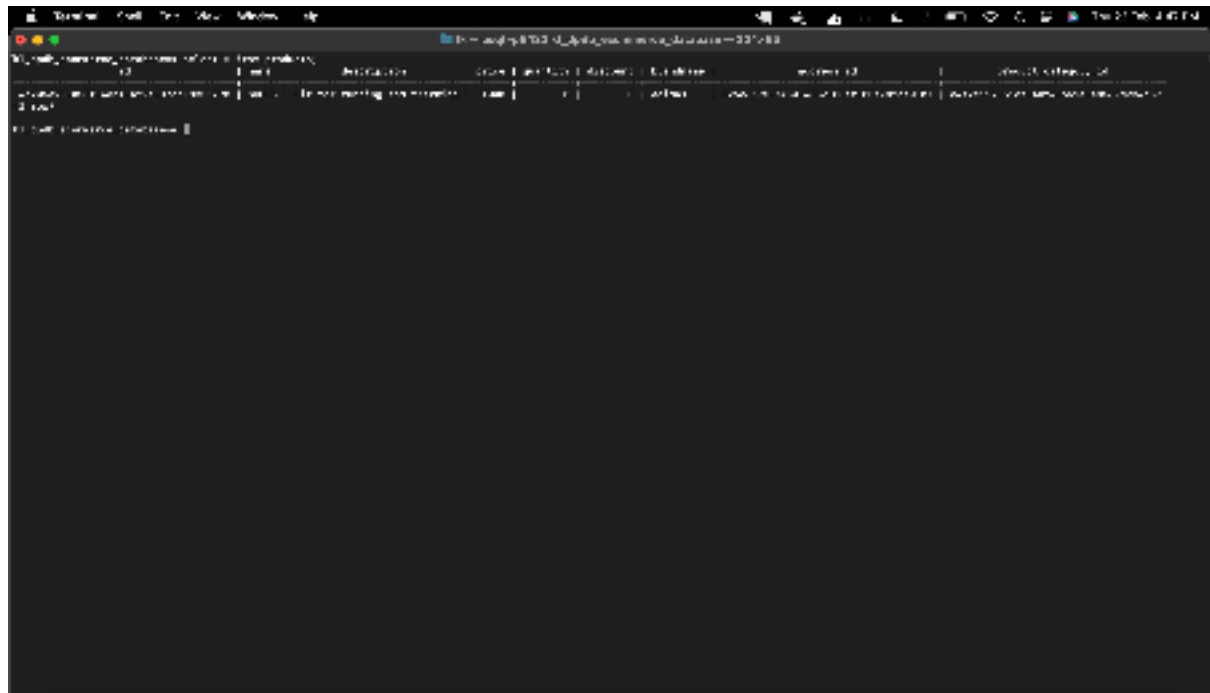


Fig-49: Run the python script.

Explanation: The python script is executed. Observed that the order got placed by only one customer(52d999b8-de44-489d-8913-7ce999e26c5a) as there is only product. For other user the order got rejected and displayed "Product not available for '207397bf-5759-4a51-a44a-3b317971b09a! please try again.'" This ensures is data consistent in the database and handles the concurrent orders.

query:

```
select * from products;
```



product_id	product_name	quantity	unit	price	category
1	Apple	0	kg	100	Fruit

Fig-50: Quantity is 0 in products table.

Explanation: once after order is placed, observed that the quantity is '0' for respective product. This is that data is maintaining the consistency.

```
select * from orders;
```

[illegible]

Explanation: order placed for customer- 52d999b8-de44-489d-8913-7ce999e26c5a and displayed the respective details.

query:

```
select * from transaction_summary;  
select * from orders_products;  
select * from delivery_partner;  
select * from customer_delivery_partner;
```

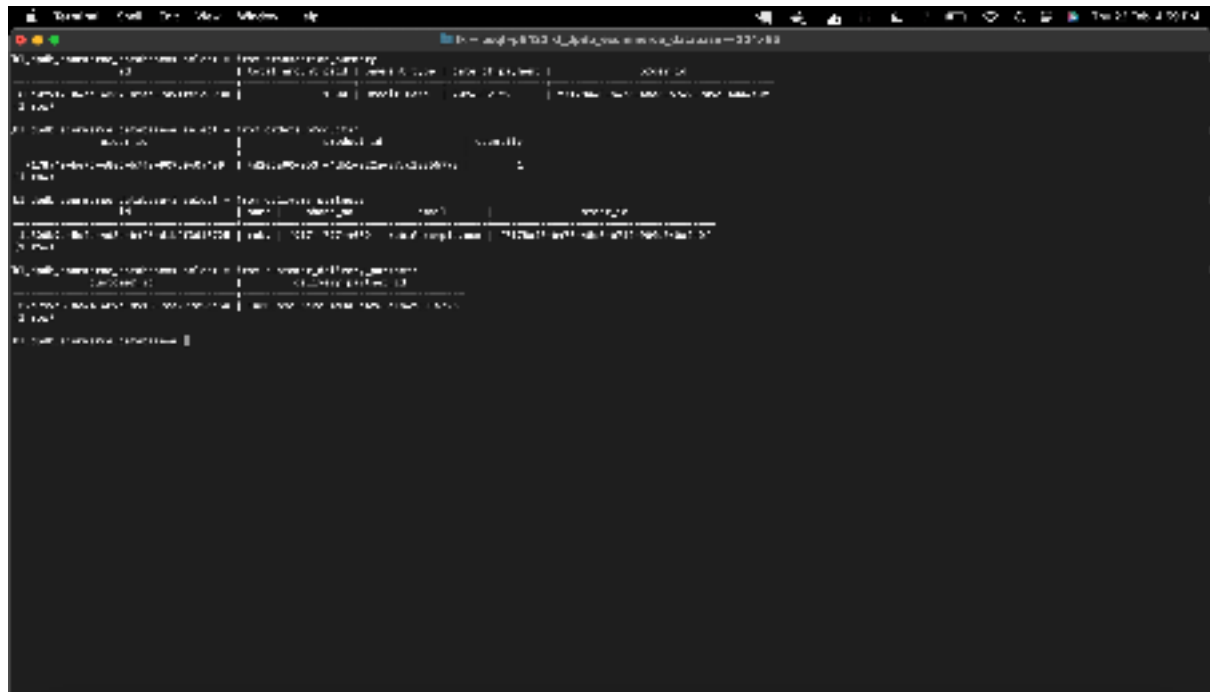


Fig-52: Updated data in all the required tables.

Explanation: updated data in all the required tables (transaction_summary, orders_products, delivery_partner, customer_delivery_partner). Which says data is consistent.

Use case 2:

when there are two product quantities, and two distinct customers are attempting to place order for the same product. Both clients ought to be able to place the order in this instance.

query:

```
select * from products;  
UPDATE products SET quantity = 2 WHERE id='4d205a90-053f-4db1-a22a-37cc18355798';  
select * from products;
```

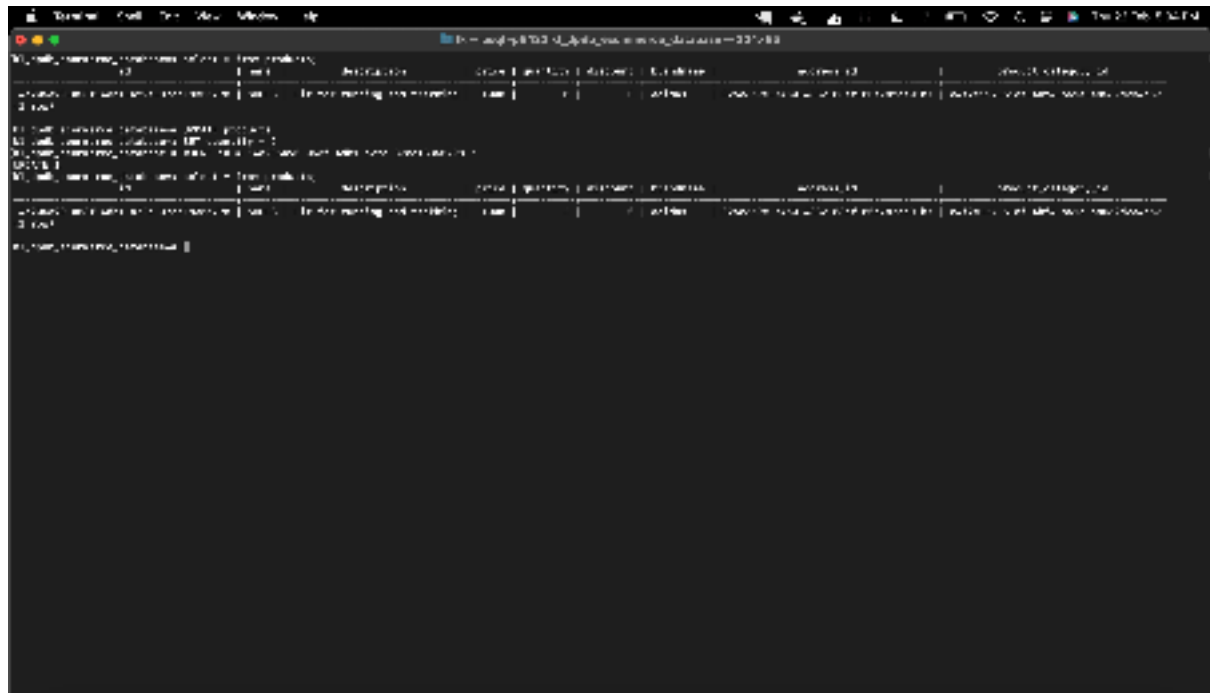
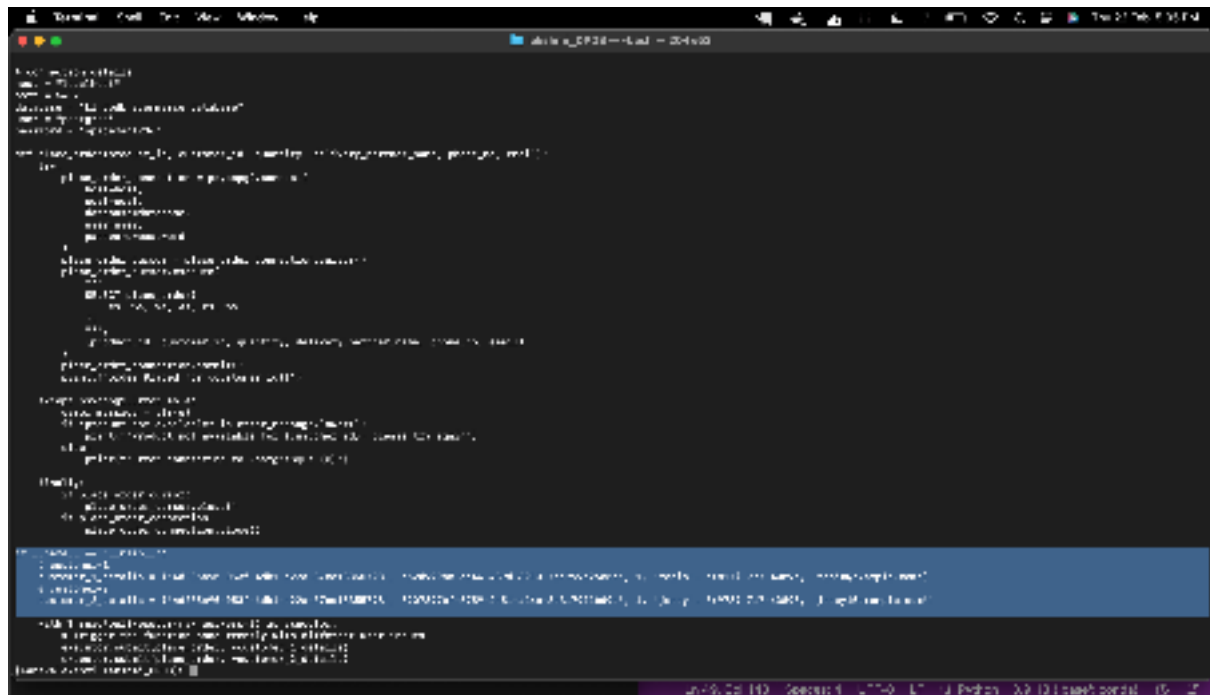


Fig-53: Update the product quantity to 2.

query:

```
if __name__ == "__main__":
    # customer-1
    customer_1_details = ('4d205a90-053f-4db1-a22a-37cc18355798', '52d999b8-de44-489d-8913-7ce999e26c5a', 1, 'robin', '(817) 777-4089', 'robin@example.com')
    # customer-2
    customer_2_details = ('4d205a90-053f-4db1-a22a-37cc18355798', '207397bf-5759-4a51-a44a-3b317971b09a', 1, 'josey', '(978) 717-4389', 'josey1@example.com')
```



```
def update_inputs(customer_id, name, phone, email, address, city, state, zip_code):
    """Update the inputs for a customer in the database.

    Parameters:
        customer_id (str): The customer ID.
        name (str): The customer name.
        phone (str): The customer phone number.
        email (str): The customer email address.
        address (str): The customer address.
        city (str): The customer city.
        state (str): The customer state.
        zip_code (str): The customer zip code.

    Returns:
        bool: True if the update was successful, False otherwise.
    """
    # Connect to the database
    conn = sqlite3.connect('customer.db')
    cursor = conn.cursor()

    # Update the customer information
    cursor.execute(
        'UPDATE customers SET name=?, phone=?, email=?, address=?, city=?, state=?, zip_code=? WHERE customer_id=?'
    )

    # Commit the changes
    conn.commit()

    # Close the connection
    conn.close()

    return True


# Test the function
customer_1_details = ('4d205a90-053f-4db1-a22a-37cc18355798', '52d999b8-de44-489d-8913-7ce999e26c5a', 1, 'robin', '(817) 777-4089', 'robin@example.com')
customer_2_details = ('4d205a90-053f-4db1-a22a-37cc18355798', '207397bf-5759-4a51-a44a-3b317971b09a', 1, 'josey', '(978) 717-4389', 'josey1@example.com')
```

Fig-54: Update the inputs as per the use case-2.

query:

python kl_dpdb_req-4.py


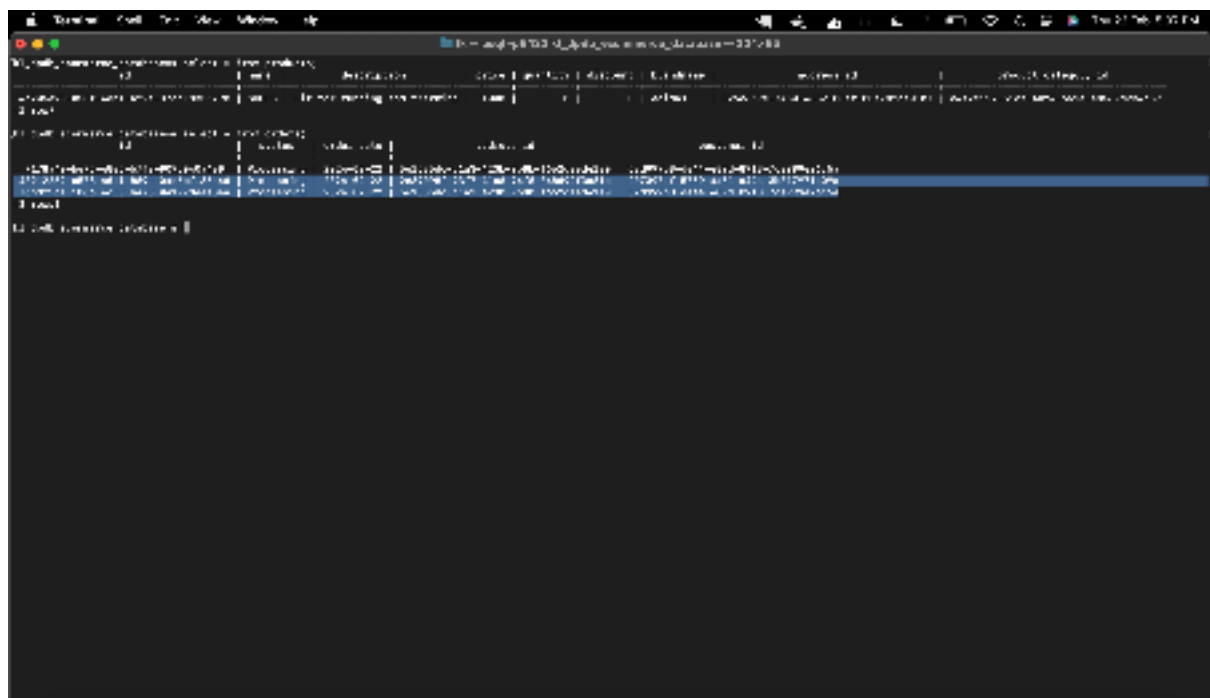


Fig-55: order successfully placed by two customers.

Explanation: Two different users wish to place an order with quantity one, and there are two products available in the products table. So, the order was successfully placed for the two customers.

query:

```
select * from products;  
select * from orders;
```



The screenshot shows a terminal window with a SQL query result. The query is `select * from orders;`. The result is a table with 10 columns: `id`, `customer_id`, `product_id`, `quantity`, `order_date`, `status`, `address`, `city`, `state`, and `zip`. There are two rows of data. The first row has `id` 1, `customer_id` 1, `product_id` 1, `quantity` 1, `order_date` 2023-01-01, `status` processing, `address` 123 Main St, `city` New York, `state` NY, and `zip` 10001. The second row has `id` 2, `customer_id` 2, `product_id` 2, `quantity` 2, `order_date` 2023-01-02, `status` processing, `address` 456 Elm St, `city` Los Angeles, `state` CA, and `zip` 90001.

id	customer_id	product_id	quantity	order_date	status	address	city	state	zip
1	1	1	1	2023-01-01	processing	123 Main St	New York	NY	10001
2	2	2	2	2023-01-02	processing	456 Elm St	Los Angeles	CA	90001

Fig-56: Two orders placed by two different customers.

Explanation: The order details got successfully inserted into the orders table with status as processing, order_date with current date, address associated with the customer.

```
select * from transaction_summary;
```

[illegible]

Fig-57: Updated transaction summary for two orders.

query:

```
select * from orders_products;
select * from delivery_partner;
select * from customer_delivery_partner;
```

[illegible]

Fig-58: Data updated in all the dependency tables.

use case 3: when two clients attempt to place an order repeatedly. As an example, let's say that a product has quantity "10," and customers 1 and customer 2 each placed three orders.

query:

```
UPDATE products SET quantity = 10 WHERE id='4d205a90-053f-4db1-a22a-37cc18355798';
```

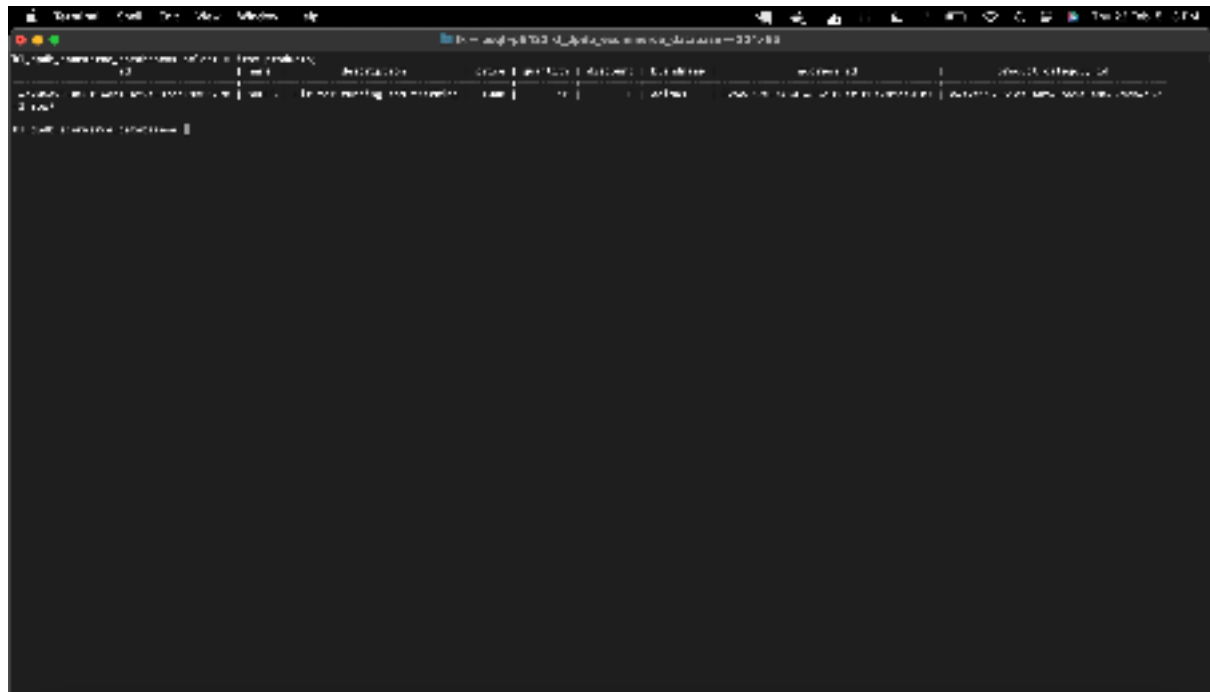


Fig-59: Updated product quantity to 10

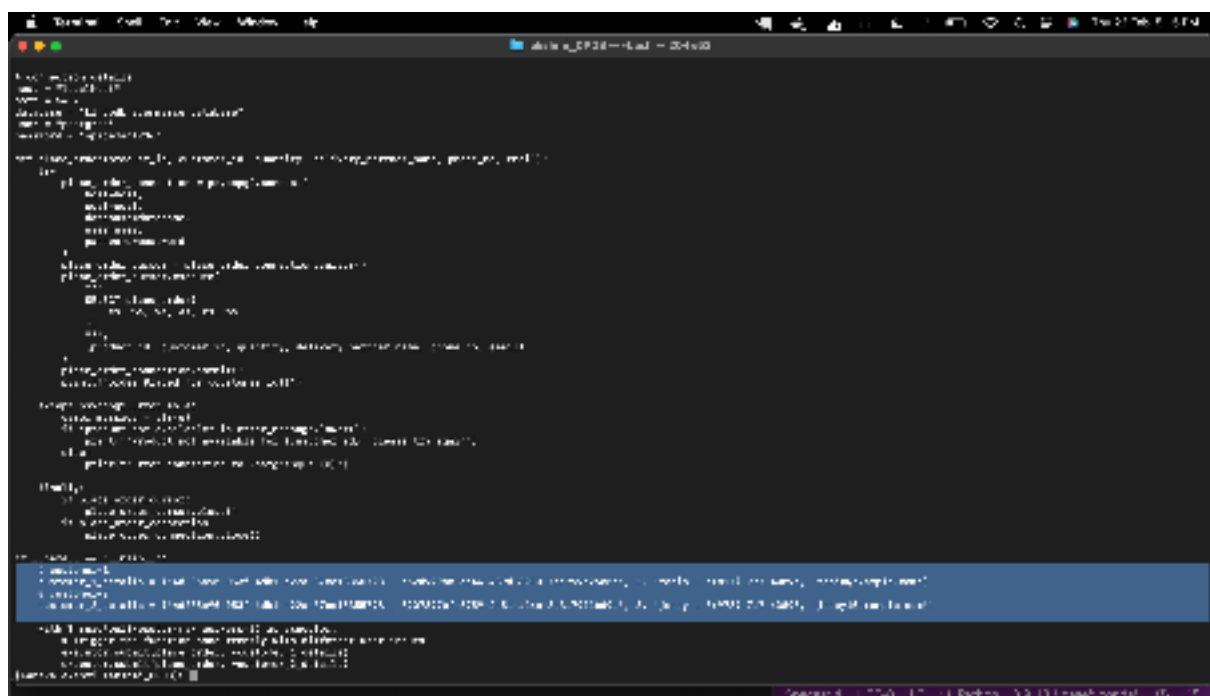


Fig-60: Updated the script as per the use case 3.

```
python kl_dpdb_req-4.py
```

Fig-61: Order placed multiple times with two customers.

```
select * from products;
```

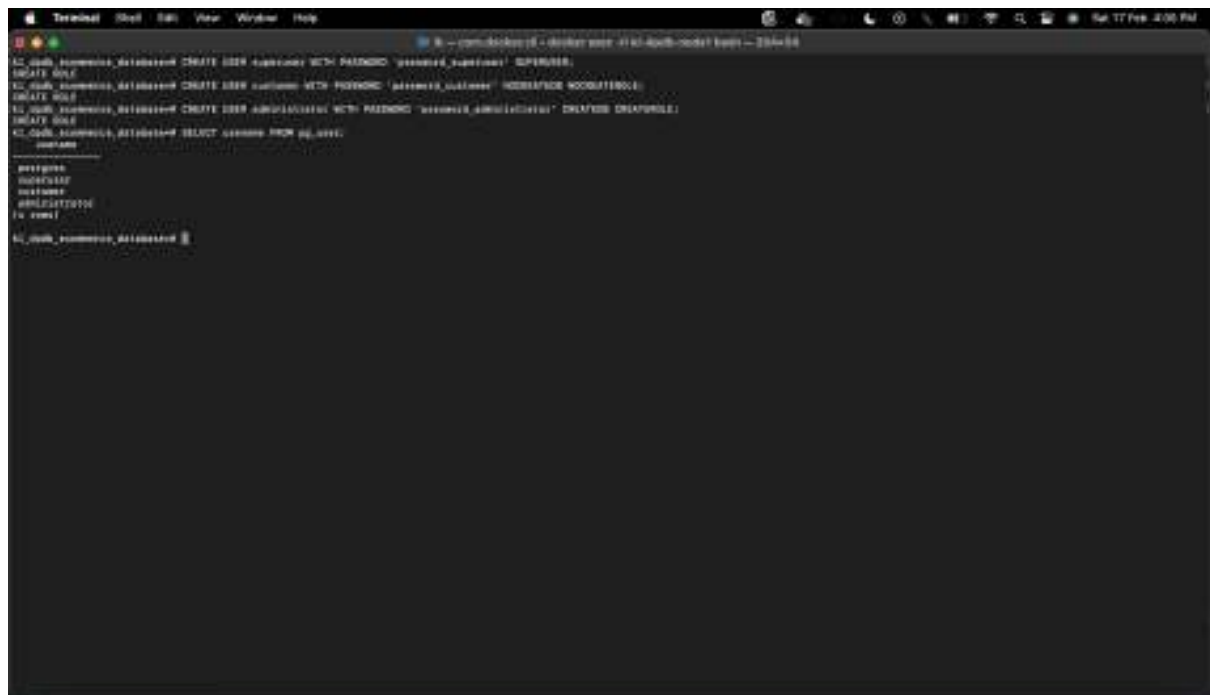
Fig-62: Quantity of the product is 0.

Requirement-5

In the database system they will be multiple users who access to the database. But all the users will not have all privileges to all tables in the database. Users have the privileges based on their responsibilities.

query:

```
-- Create Users:
-- Create superuser
CREATE USER superuser WITH PASSWORD 'password_superuser' SUPERUSER;
-- Create customers user
CREATE USER customer WITH PASSWORD 'password_customer' NOCREATEDB NOCREATEROLE;
-- Create administrators user
CREATE USER administrators WITH PASSWORD 'password_administrators' CREATEDB
CREATEROLE;
-- list of users
SELECT username FROM pg_user;
```

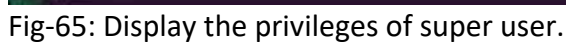


```
Terminal [Shell] [File] [View] [Window] [Help]
k:/_sql_howto/pg - docker exec -it sql-howto-postgres /bin/bash
k/_sql_howto/pg: CREATE USER superuser WITH PASSWORD 'password_superuser' SUPERUSER;
CREATE ROLE
k/_sql_howto/pg: CREATE USER customer WITH PASSWORD 'password_customer' NOCREATEDB NOCREATEROLE;
CREATE ROLE
k/_sql_howto/pg: CREATE USER administrators WITH PASSWORD 'password_administrators' CREATEDB
CREATE ROLE;
k/_sql_howto/pg: SELECT username FROM pg_user;
username
superuser
customer
administrators
(3 rows)
k/_sql_howto/pg: 
```

Fig-63: Create users to give access control.

Explanation: Created three different users – superuser, customer and administrator. Once after users are created, displayed the list of users.

Fig-64: Grant access to superuser.



Explanation: Ideally super user has access on all tables in the database system with all privileges so, provided all privileges to super user and displayed the complete data.

Grant access to administrator:

query:

```
CREATE OR REPLACE FUNCTION grant_privileges_to_administrator() RETURNS VOID AS $$
BEGIN
-- Granting privileges on the 'Products' table
EXECUTE 'GRANT SELECT, INSERT, UPDATE, DELETE ON TABLE Products TO administrator';
-- Granting privileges on the 'Customer' table
EXECUTE 'GRANT SELECT, INSERT, UPDATE, DELETE ON TABLE Customer TO administrator';
-- Granting privileges on the 'cart' table
EXECUTE 'GRANT SELECT, UPDATE, DELETE ON TABLE cart TO administrator';
-- Granting privileges on the 'orders' table
EXECUTE 'GRANT SELECT, INSERT, UPDATE, DELETE ON TABLE orders TO administrator';
-- Granting privileges on the 'Transaction summary' table
EXECUTE 'GRANT SELECT ON TABLE "transaction_summary" TO administrator';
-- Granting privileges on the 'supplier' table
EXECUTE 'GRANT SELECT, INSERT, UPDATE, DELETE ON TABLE supplier TO administrator';
-- Granting privileges on the 'Address' table
EXECUTE 'GRANT SELECT, INSERT, UPDATE, DELETE ON TABLE Address TO administrator';
-- Granting privileges on the 'product_category' table
EXECUTE 'GRANT SELECT, INSERT, UPDATE, DELETE ON TABLE product_category TO
administrator';
-- Granting privileges on the 'Reviews' table
EXECUTE 'GRANT SELECT, DELETE ON TABLE Reviews TO administrator';
-- Granting privileges on the 'gift_vouchers' table
EXECUTE 'GRANT SELECT, INSERT, UPDATE, DELETE ON TABLE gift_vouchers TO
administrator';
-- Granting privileges on the 'delivery_partner' table
EXECUTE 'GRANT SELECT, INSERT, UPDATE, DELETE ON TABLE delivery_partner TO
administrator';
    END;
    $$ LANGUAGE plpgsql;

-- Execute the function to grant privileges
SELECT grant_privileges_to_administrator();

SELECT table_name, grantee, privilege_type
FROM information_schema.table_privileges
WHERE table_catalog = 'kl_dpdb_ecommerce_database' AND grantee = 'administrator';
```

```

mysql> GRANT SELECT, INSERT, UPDATE, DELETE ON TABLE Products TO administrator;
mysql> GRANT SELECT, INSERT, UPDATE, DELETE ON TABLE Customers TO administrator;
mysql> GRANT SELECT, UPDATE, DELETE ON TABLE Invt TO administrator;
mysql> GRANT SELECT, INSERT, UPDATE, DELETE ON TABLE orders TO administrator;
mysql> GRANT SELECT ON TABLE 'SalesOrderSummary' TO administrator;
mysql> GRANT SELECT, INSERT, UPDATE, DELETE ON TABLE similar TO administrator;
mysql> GRANT SELECT, INSERT, UPDATE, DELETE ON TABLE Address TO administrator;
mysql> GRANT SELECT, INSERT, UPDATE, DELETE ON TABLE inventory TO administrator;
mysql> GRANT SELECT, INSERT, UPDATE, DELETE ON TABLE gift_vouchers TO administrator;
mysql> GRANT SELECT, INSERT, UPDATE, DELETE ON TABLE delivery_order TO administrator;

-- IN LANGUAGE SQL;
CREATE FUNCTION
  kl_sdb_administrator_privileges
  kl_sdb_administrator_privileges
  kl_sdb_administrator_privileges SELECT table_name, grantee, privilege_type
  kl_sdb_administrator_privileges FROM information_schema.table_privileges
  kl_sdb_administrator_privileges WHERE table_name = 'kl_sdb_administrator' AND grantee = 'administrator';

table_name | grantee | privilege_type
-----
products | administrator | INSERT
products | administrator | SELECT
products | administrator | UPDATE
products | administrator | DELETE
orders | administrator | INSERT
orders | administrator | SELECT
orders | administrator | UPDATE
products_category | administrator | SELECT
products_category | administrator | UPDATE
products_category | administrator | DELETE
invt | administrator | SELECT
invt | administrator | UPDATE
invt | administrator | INSERT
orders | administrator | INSERT
orders | administrator | SELECT
orders | administrator | UPDATE
SalesOrderSummary | administrator | INSERT
customers | administrator | INSERT
customers | administrator | UPDATE
customers | administrator | SELECT
suppliers | administrator | INSERT
suppliers | administrator | SELECT
suppliers | administrator | UPDATE
inventory | administrator | INSERT
inventory | administrator | SELECT
inventory | administrator | UPDATE
gift_vouchers | administrator | INSERT
gift_vouchers | administrator | SELECT

```

Fig-66: Grant access to administrator.

```

mysql> GRANT SELECT, DELETE ON TABLE Address TO administrator;
mysql> GRANT SELECT, INSERT, UPDATE, DELETE ON TABLE gift_vouchers TO administrator;
mysql> GRANT SELECT, INSERT, UPDATE, DELETE ON TABLE delivery_order TO administrator;

-- IN LANGUAGE SQL;
CREATE FUNCTION
  kl_sdb_administrator_privileges
  kl_sdb_administrator_privileges
  kl_sdb_administrator_privileges SELECT table_name, grantee, privilege_type
  kl_sdb_administrator_privileges FROM information_schema.table_privileges
  kl_sdb_administrator_privileges WHERE table_name = 'kl_sdb_administrator' AND grantee = 'administrator';

table_name | grantee | privilege_type
-----
products | administrator | INSERT
products | administrator | SELECT
products | administrator | UPDATE
products | administrator | DELETE
orders | administrator | INSERT
orders | administrator | SELECT
orders | administrator | UPDATE
products_category | administrator | SELECT
products_category | administrator | UPDATE
products_category | administrator | DELETE
invt | administrator | SELECT
invt | administrator | UPDATE
invt | administrator | INSERT
orders | administrator | INSERT
orders | administrator | SELECT
orders | administrator | UPDATE
SalesOrderSummary | administrator | INSERT
customers | administrator | INSERT
customers | administrator | UPDATE
customers | administrator | SELECT
suppliers | administrator | INSERT
suppliers | administrator | SELECT
suppliers | administrator | UPDATE
inventory | administrator | INSERT
inventory | administrator | SELECT
inventory | administrator | UPDATE
gift_vouchers | administrator | INSERT
gift_vouchers | administrator | SELECT
gift_vouchers | administrator | UPDATE
delivery_order | administrator | INSERT
delivery_order | administrator | SELECT
delivery_order | administrator | UPDATE

```

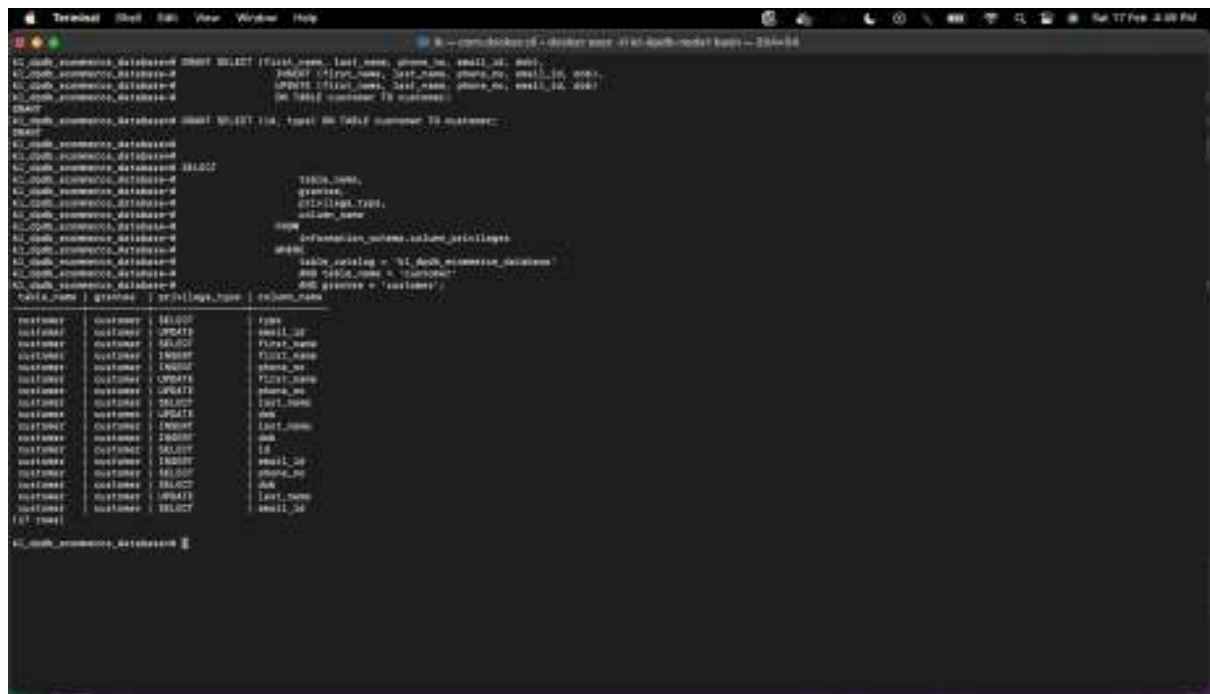
Fig-67: list of complete privileges for administrator.

Explanation: Granted privileges to administrator on the respective tables. Displayed the list of privileges given to the administrator.

Grant access to customer user on customer table:

query:

```
-- Granting SELECT, INSERT, UPDATE privileges on specified columns
GRANT SELECT (first_name, last_name, phone_no, email_id, dob),
INSERT (first_name, last_name, phone_no, email_id, dob),
UPDATE (first_name, last_name, phone_no, email_id, dob)
ON TABLE customer TO customer;
-- Granting SELECT, UPDATE privileges on specified columns
GRANT SELECT (id, type) ON TABLE customer TO customer;
-- cross check the preveligies
SELECT table_name, grantee, privilege_type, column_name
FROM information_schema.column_privileges
WHERE table_catalog = 'kl_dpdb_ecommerce_database' AND table_name = 'customer' AND
grantee = 'customer';
```



The screenshot shows a SQL Developer terminal window with the following content:

```
kl_dpdb_ecommerce_database> GRANT SELECT (first_name, last_name, phone_no, email_id, dob),
kl_dpdb_ecommerce_database> INSERT (first_name, last_name, phone_no, email_id, dob),
kl_dpdb_ecommerce_database> UPDATE (first_name, last_name, phone_no, email_id, dob)
kl_dpdb_ecommerce_database> ON TABLE customer TO customer;
GRANT
kl_dpdb_ecommerce_database> GRANT SELECT (id, type) ON TABLE customer TO customer;
GRANT
kl_dpdb_ecommerce_database> SELECT
kl_dpdb_ecommerce_database> table_name, grantee, privilege_type, column_name
kl_dpdb_ecommerce_database> FROM information_schema.column_privileges
kl_dpdb_ecommerce_database> WHERE table_catalog = 'kl_dpdb_ecommerce_database' AND table_name = 'customer' AND
kl_dpdb_ecommerce_database> grantee = 'customer';
```

table_name	grantee	privilege_type	column_name
customer	customer	SELECT	first_name
customer	customer	SELECT	last_name
customer	customer	SELECT	phone_no
customer	customer	SELECT	email_id
customer	customer	SELECT	dob
customer	customer	INSERT	first_name
customer	customer	INSERT	last_name
customer	customer	INSERT	phone_no
customer	customer	INSERT	email_id
customer	customer	INSERT	dob
customer	customer	UPDATE	first_name
customer	customer	UPDATE	last_name
customer	customer	UPDATE	phone_no
customer	customer	UPDATE	email_id
customer	customer	UPDATE	dob

Fig-68: Grant access to customer user on customer table.

Explanation: Provided the privileges to customer on the customer table and displayed it.

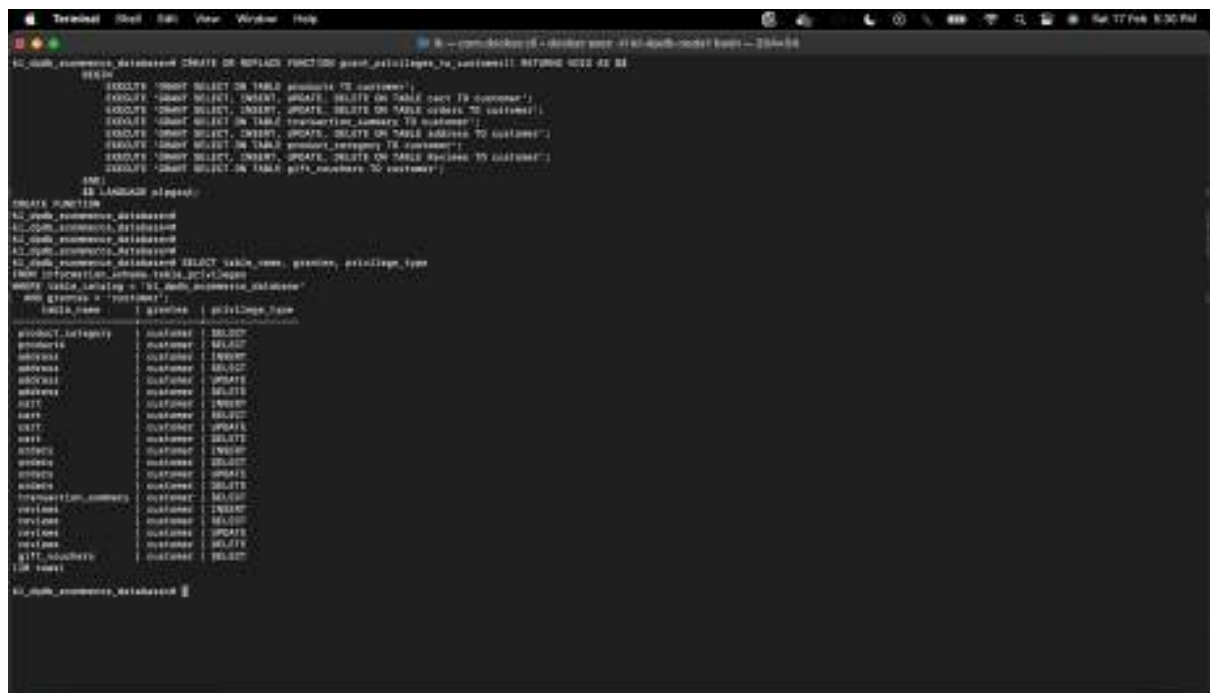
Grant access to customer user on remaining tables:

query:

```
CREATE OR REPLACE FUNCTION grant_privileges_to_customer() RETURNS VOID AS $$
BEGIN
-- Granting privileges on the 'Products' table
EXECUTE 'GRANT SELECT ON TABLE Products TO customer';
-- Granting privileges on the 'cart' table
EXECUTE 'GRANT SELECT, INSERT, UPDATE, DELETE ON TABLE cart TO customer';
-- Granting privileges on the 'orders' table
EXECUTE 'GRANT SELECT, INSERT, UPDATE, DELETE ON TABLE orders TO customer';
-- Granting privileges on the 'transaction_summary' table
EXECUTE 'GRANT SELECT ON TABLE transaction_summary TO customer';
-- Granting privileges on the 'address' table
EXECUTE 'GRANT SELECT, INSERT, UPDATE, DELETE ON TABLE address TO customer';
-- Granting privileges on the 'product_category' table
EXECUTE 'GRANT SELECT ON TABLE product_category TO customer';
-- Granting privileges on the 'Reviews' table
EXECUTE 'GRANT SELECT, INSERT, UPDATE, DELETE ON TABLE Reviews TO customer';
-- Granting privileges on the 'gift_vouchers' table
EXECUTE 'GRANT SELECT ON TABLE gift_vouchers TO customer';
END;
$$ LANGUAGE plpgsql;
```

-- Execute the function to grant privileges

```
SELECT grant_privileges_to_customer();
SELECT table_name, grantee, privilege_type
FROM information_schema.table_privileges
WHERE table_catalog = 'kl_dpdb_ecommerce_database' AND grantee = 'customer';
```



```
kl_dpdb_ecommerce_database@ORA11G-RAC1:~$ sqlplus 'scott/tiger@//localhost:1521/orcl'
SQL> EXECUTE grant_privileges_to_customer();
EXECUTE grant_privileges_to_customer()
SQL> SELECT table_name, grantee, privilege_type
FROM information_schema.table_privileges
WHERE table_catalog = 'kl_dpdb_ecommerce_database'
AND grantee = 'customer';
TABLE_NAME          GRANTEE            PRIVILEGE_TYPE
-----
PRODUCT_CATEGORY    customer           SELECT
PRODUCTS            customer           SELECT
CART                 customer           SELECT, INSERT, UPDATE, DELETE
ORDERS               customer           SELECT, INSERT, UPDATE, DELETE
TRANSACTION_SUMMARY customer           SELECT
ADDRESS              customer           SELECT, INSERT, UPDATE, DELETE
PRODUCT_CATEGORY     customer           SELECT
REVIEWS              customer           SELECT, INSERT, UPDATE, DELETE
GIFT_VOUCHERS        customer           SELECT
```

Fig-69: Grant access to customer user on remaining tables and display respective privileges.

query:

```
psql -U customer -d kl_dpdb_ecommerce_database
```

```
select * from products;
```

```
delete from products where id='b136c1cf-d1e8-483f-9064-1c866f25195f';
```

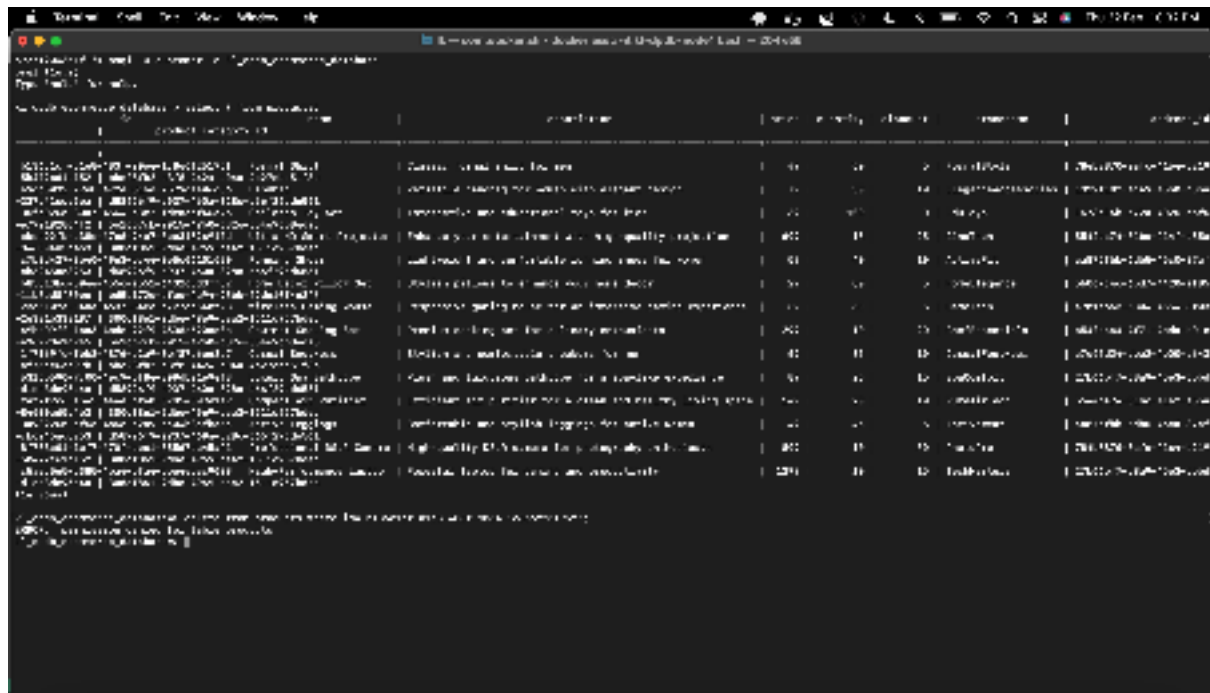


Fig-70: Delete Product from customer user.

Explanation: After providing the grant access to customer, tried to delete the data from products table got an error message “ permission denied for table products”.

Requirement – 6:

The process of copying the data from one server to another server in the postgres is known as the postgres data replication. The postgres supports the replication strategy, which achieves fault tolerance, data migration, parallel execution with good performance. It has the single-master Replication, multi-master replication architecture. The replication can implement in uni-directional/bi-directional.

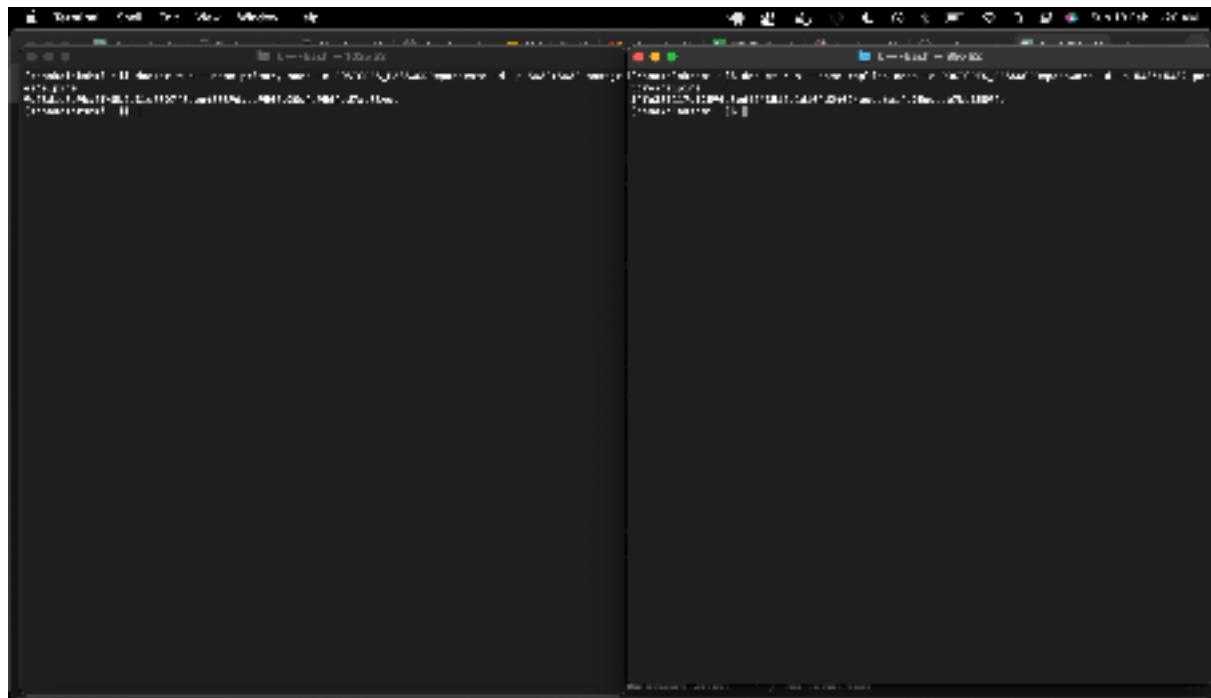


Fig-71: create two postgres instances.

```
query:
which initdb
initdb -D /tmp/db_primary
```

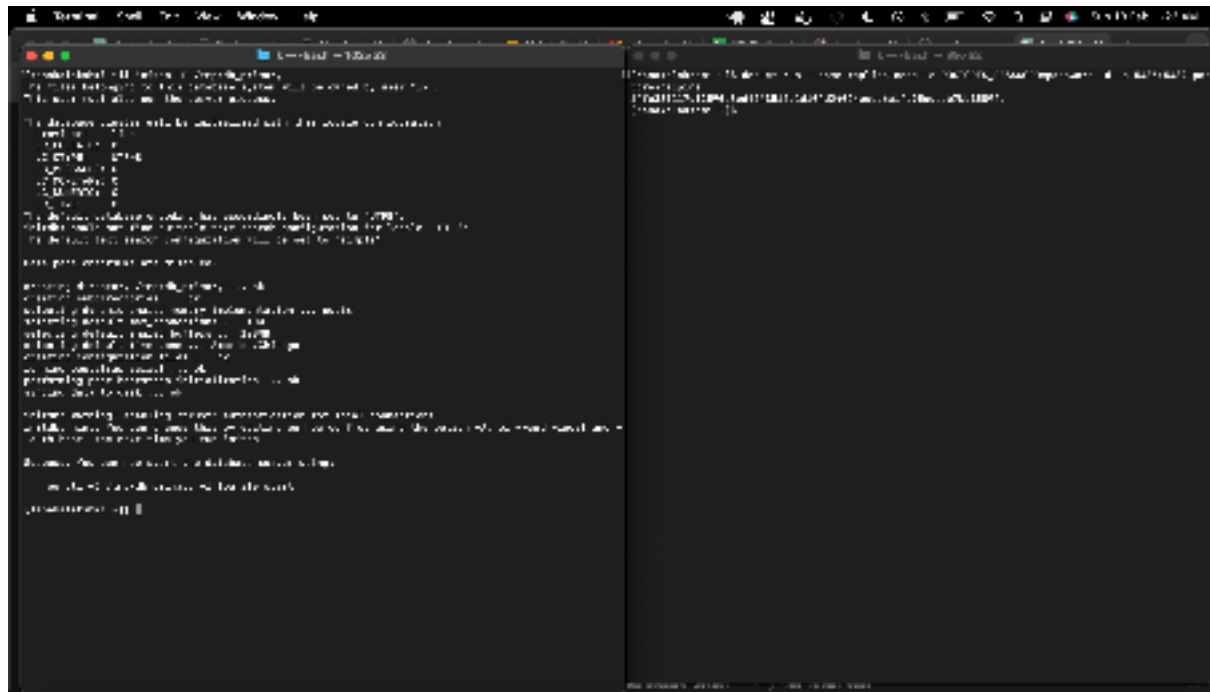


Fig-72: created Directory for the primary DB.

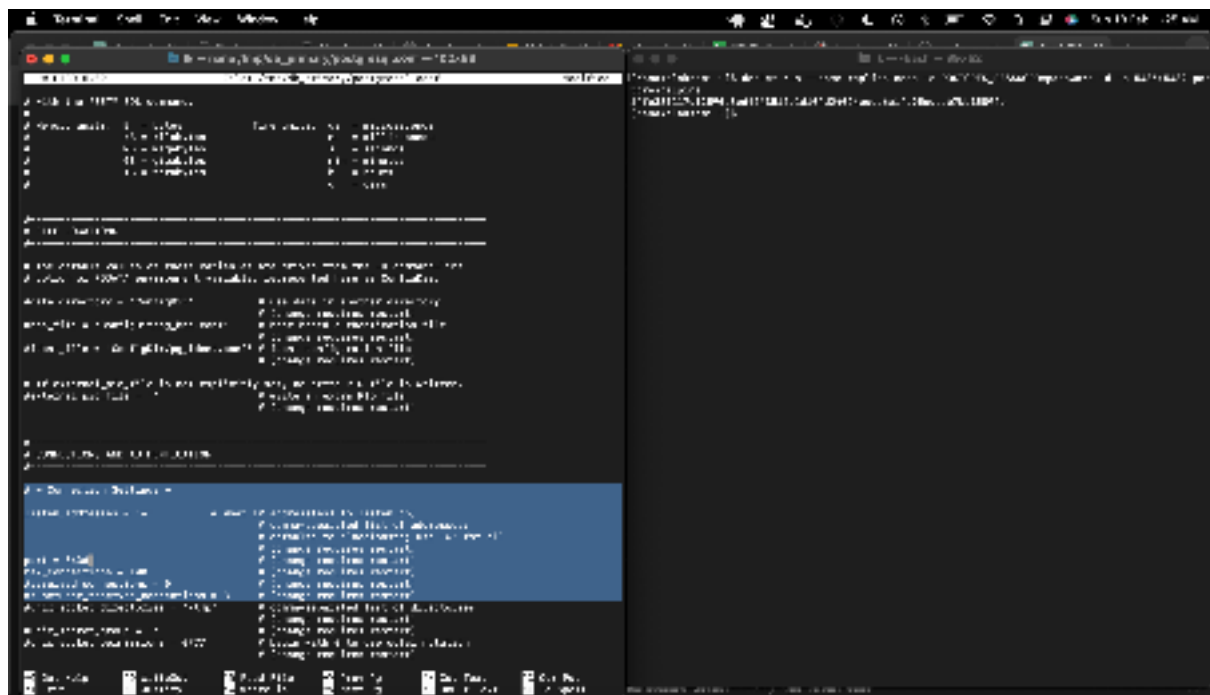




Fig-73: update configs in postgresql.conf

```
nano /tmp/primary_db/postgresql.conf
pg_ctl -D /tmp/primary_db start
```

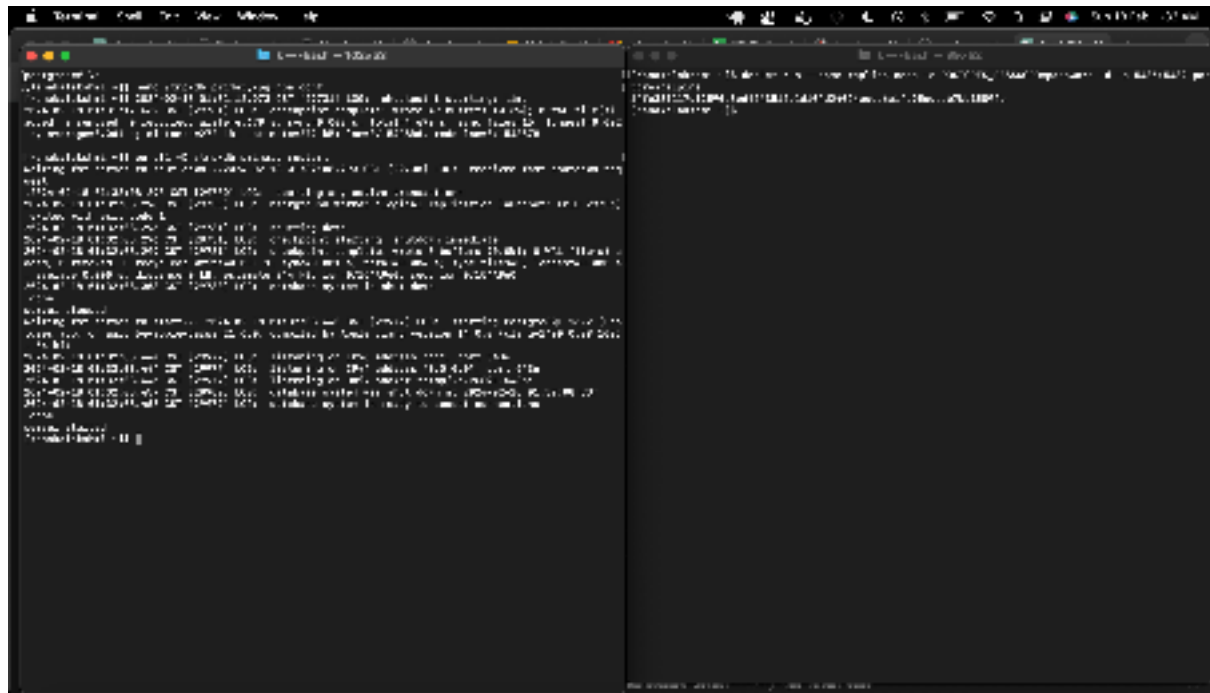



```
psql --port=5436 postgres
create user replica_user replication;
```



query:

pg_ctl -D /tmp/primary_db restart

A screenshot of a terminal window with a dark background. The terminal shows the execution of the command 'pg_ctl -D /tmp/primary_db restart'. The output indicates that the PostgreSQL server is being restarted. The terminal text is as follows:

```
postgres@pg:~$ pg_ctl -D /tmp/primary_db restart
pg_ctl: waiting for server to shut down... done
pg_ctl: waiting for server to start... done
pg_ctl: server started
```

Fig-77: Restart the postgres instance.

Explanation: Once after all the configurations are done, restarted postgres instance to reflect the latest config changes.

query:

```
pg_basebackup -h localhost -U replica_user --checkpoint=fast -D /tmp/replica_db/ -R  
--slot=data_node -C --port=5436
```




Fig-78: Copy the configurations to be used on the replica node.

Explanation: created the backup file with all the configuration, which makes easier in configuring the replica server.

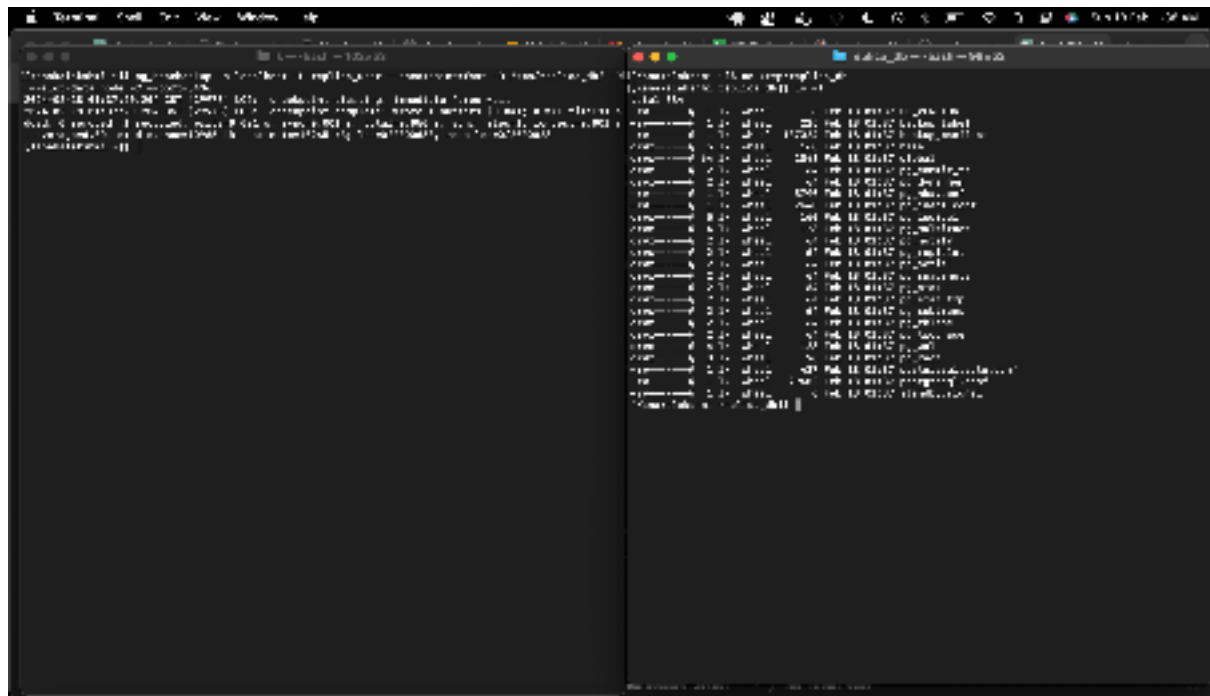


Fig-79: update the configs in replica node.

Explanation: The **standby.signal** is used for streaming replication. It always listens to the primary node to maintain the data synchronized. In case of any node failure, it detects it and communicate us by providing efficient logs.

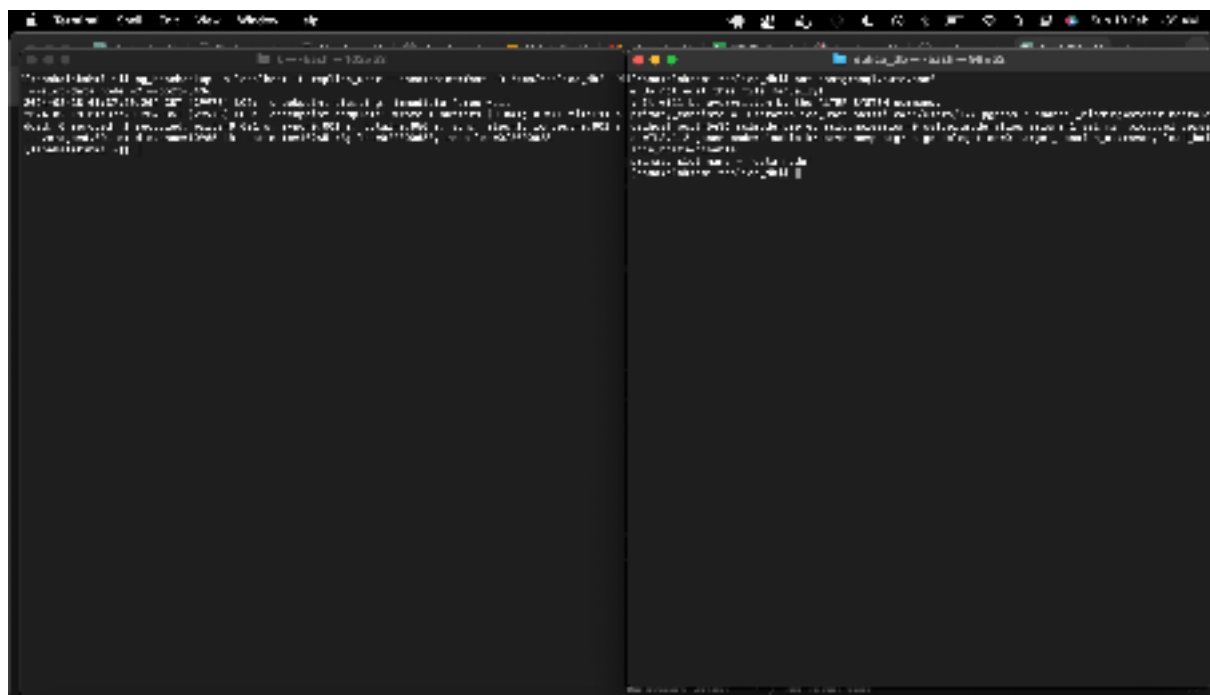


Fig-80: update the configs in replica node.

Now, Let's verify the replication strategy:

query:

instance-1:

```
CREATE TABLE customer (  
id UUID PRIMARY KEY DEFAULT uuid_generate_v4(),  
first_name VARCHAR(20) NOT NULL,  
last_name VARCHAR(20) NOT NULL,  
phone_no VARCHAR(20) NOT NULL,  
email_id VARCHAR(255) NOT NULL,  
dob DATE NOT NULL,  
type VARCHAR(20) NOT NULL,  
CONSTRAINT invalid_customer_phone CHECK (phone_no ~ '^\(\d{3}\) \d{3}-\d{4}$'),  
CONSTRAINT invalid_customer_email CHECK (email_id ~ '^[A-Za-z0-9._%+-]+@[A-Za-z0-9.-]+\.[A-Za-z]{2,}$'),  
UNIQUE (phone_no),  
UNIQUE (email_id));
```

Instance-2:

```
select * from customer;
```

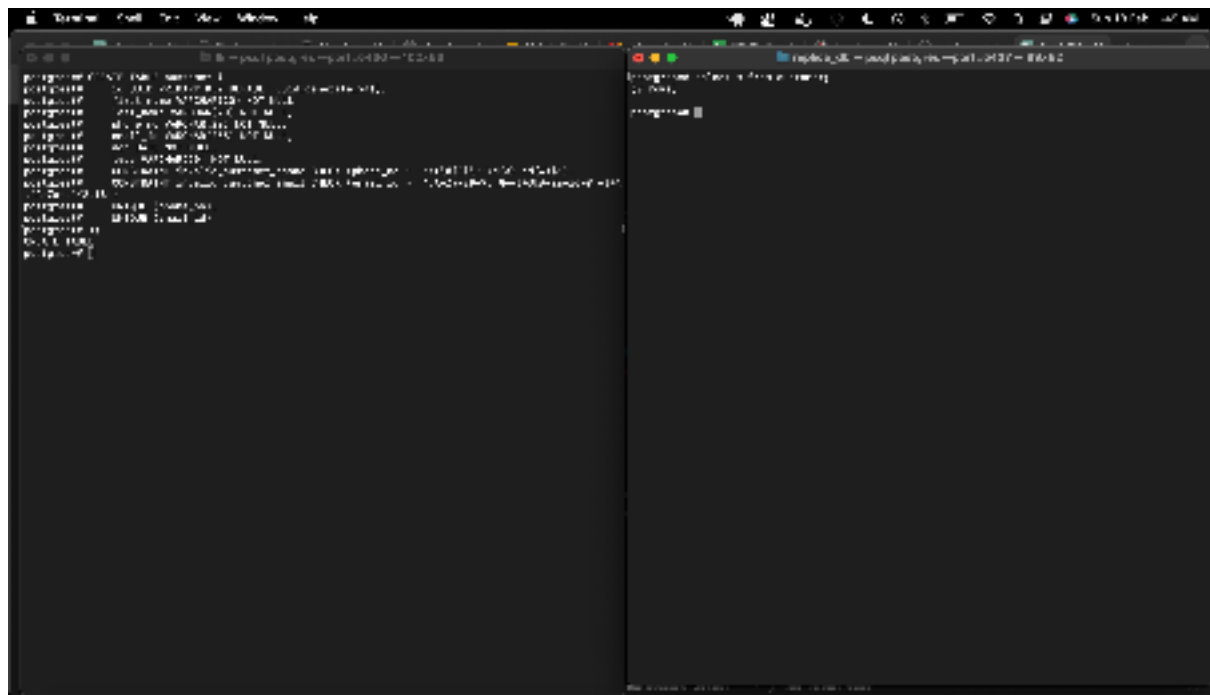


Fig-84: Create the table in instance-1 and verify data in instance-2.

Explanation: created the customer table in primary node(port:5436) and verified the data in the replica node(port:5437).

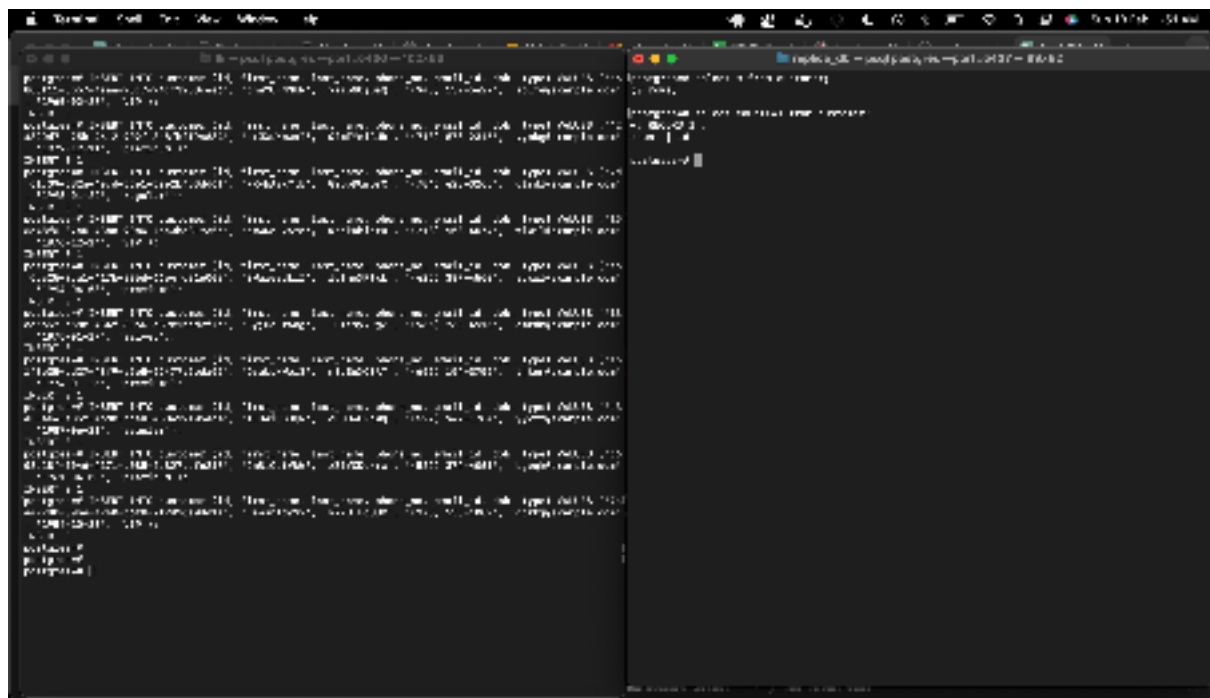


Fig-85: Insert records in instance-1 and verify the count of records from instance-2.

Explanation: inserted data into customer table in primary node(port:5436) and read the count of records in the replica node(port:5437).

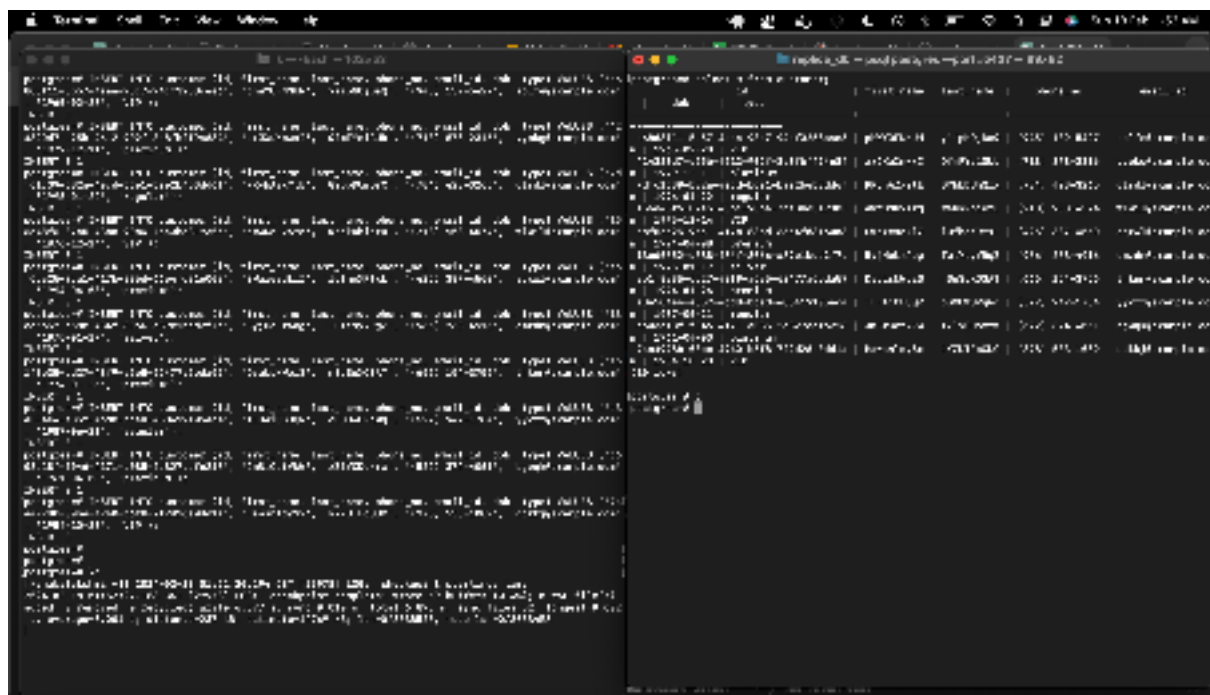


Fig-86: Display the data in instance-2 and exit instance-1.

query:

pg_ctl -D /tmp/primary_db stop --stopping the node1

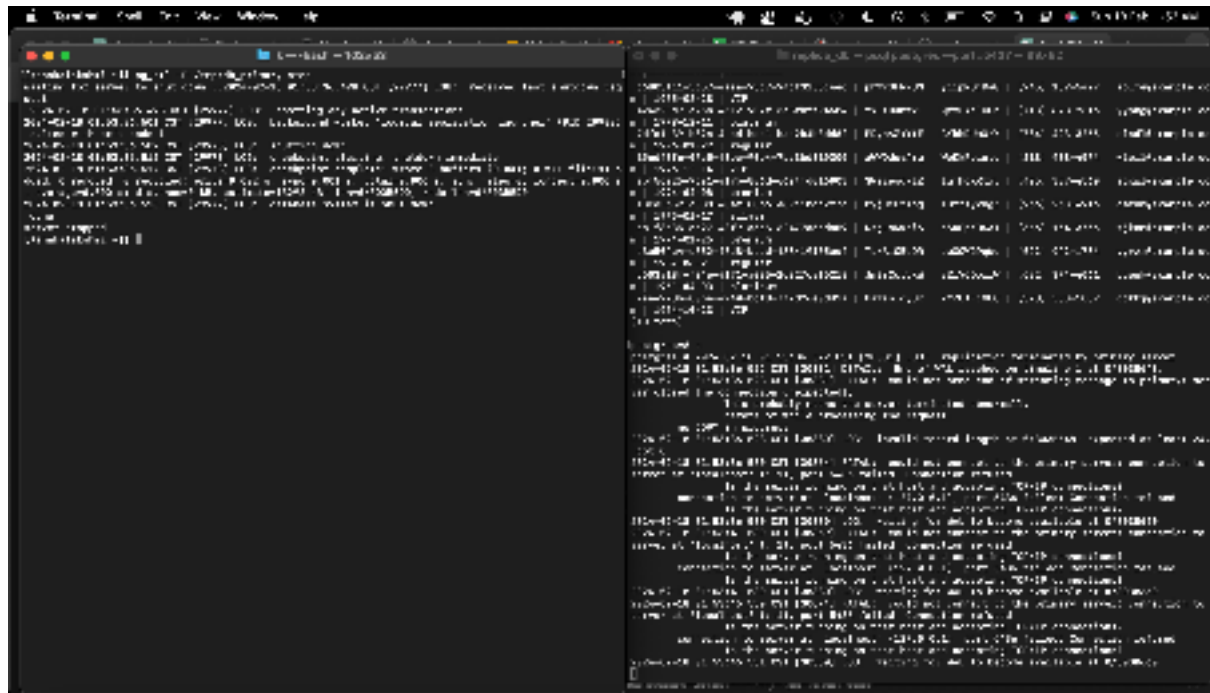


Fig-87: Bring down the instance-1.

Explanation: To verify the node failure, brought the primary node(port:5436) and verified the logs in replica node(port:5437).

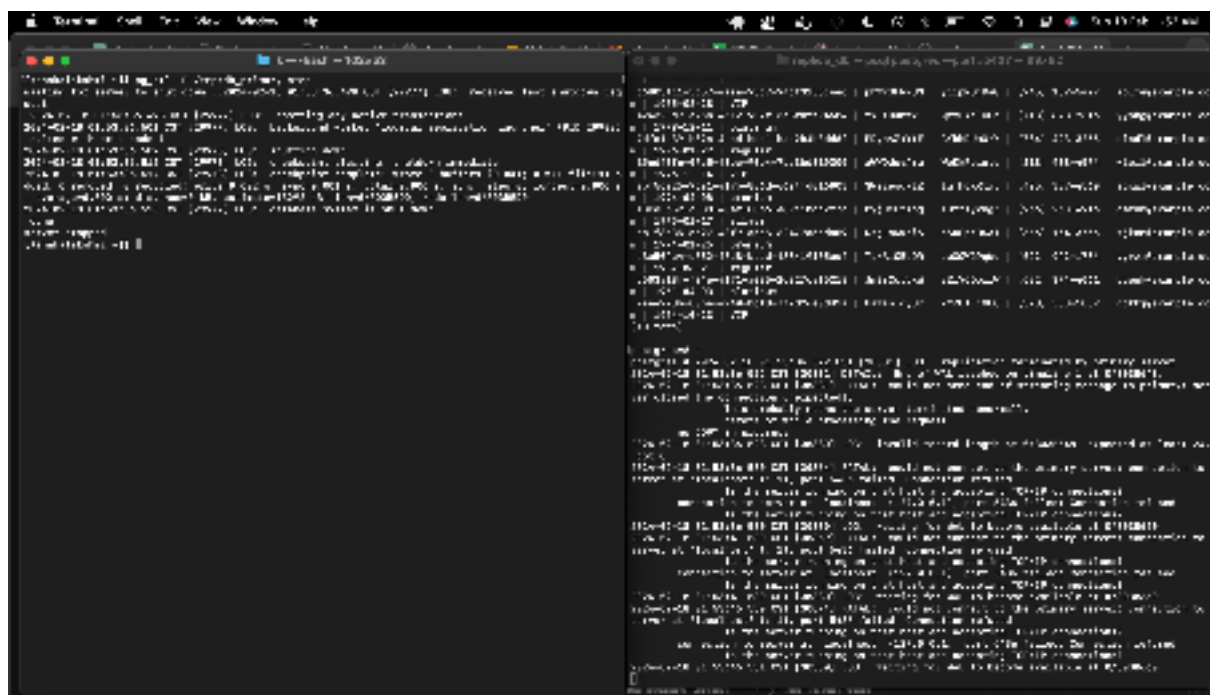


Fig-88: Observe logs in instance-2

The screenshot shows a terminal window with two panes. The left pane displays Redis logs for instance-2, showing it is a slave of instance-1 and is currently in a 'standby' state. The right pane shows a table of Redis instances, including instance-1 (master) and instance-2 (slave). The table has columns for instance_id, first_name, last_name, phone_no, and email_id.

instance_id	first_name	last_name	phone_no	email_id
1	John	Doe	123-4567	john.doe@redis.com
2	Jane	Doe	123-4567	jane.doe@redis.com
3	John	Doe	123-4567	john.doe@redis.com
4	Jane	Doe	123-4567	jane.doe@redis.com
5	John	Doe	123-4567	john.doe@redis.com
6	Jane	Doe	123-4567	jane.doe@redis.com
7	John	Doe	123-4567	john.doe@redis.com
8	Jane	Doe	123-4567	jane.doe@redis.com
9	John	Doe	123-4567	john.doe@redis.com
10	Jane	Doe	123-4567	jane.doe@redis.com

Fig-89: Display the data in instance-2.

Explanation: Even though the primary node is down, can still read the data from replica node. Which achieves the **fault tolerance**.

Requirement – 7:

Parallel Query Execution:

It speeds up the execution of queries. It adjusts the machine usage according to workloads. queries can scan massive datasets, such as join statements.

query:

```
CREATE OR REPLACE FUNCTION create_customer(  
    p_first_name VARCHAR(20),  
    p_last_name  VARCHAR(20),  
    p_phone_no   VARCHAR(20),  
    p_email_id   VARCHAR(255),  
    p_dob        DATE,  
    p_flat_no    INTEGER,  
    p_street     VARCHAR(50),  
    p_city       VARCHAR(50),  
    p_state      VARCHAR(50),  
    p_country    VARCHAR(50),  
    p_zip_code   VARCHAR(10)  
) RETURNS VOID AS $$  
DECLARE  
    v_customer_id UUID;  
    v_address_id  UUID;  
BEGIN  
    -- Insert into customer table  
    INSERT INTO customer (id, first_name, last_name, phone_no, email_id, dob, type)  
    VALUES (uuid_generate_v4(), p_first_name, p_last_name, p_phone_no, p_email_id,  
p_dob, 'regular')  
    RETURNING id INTO v_customer_id;  
  
    -- Insert into address table  
    INSERT INTO address (id, flat_no, street, city, state, country, zip_code)  
    VALUES (uuid_generate_v4(), p_flat_no, p_street, p_city, p_state, p_country,  
p_zip_code)  
    RETURNING id INTO v_address_id;  
  
    -- Insert into customer_address table  
    INSERT INTO customer_address (customer_id, address_id, default_address)  
    VALUES (v_customer_id, v_address_id, true);  
END;  
$$ LANGUAGE plpgsql PARALLEL SAFE;
```



Fig-90: Customer signup function

Explanation: Created create_customer function, which is used as the customer sign-up, which reads the input details of the customer and stores the data in the database. It input details are first_name, last_name, phone_no, emailId, dob, flat_no, street, city, state, country, zip-code. It stores the data in the customer, address, and customer_address table. As it's a new sign-up customer type will be regular, based on the orders of the customer, the type will be upgraded from backend and default address will be true as it's the first address. Implemented the function with **PARALLEL SAFE**.

query:

```
SELECT create_customer(  
    'Jack',  
    'Teresa',  
    '(910) 436-7890',  
    'jack.teresa@example.com',  
    '1990-01-01',  
    401,  
    'Main Street',  
    'Tampa',  
    'Florida',  
    'USA',  
    '65341'  
);
```

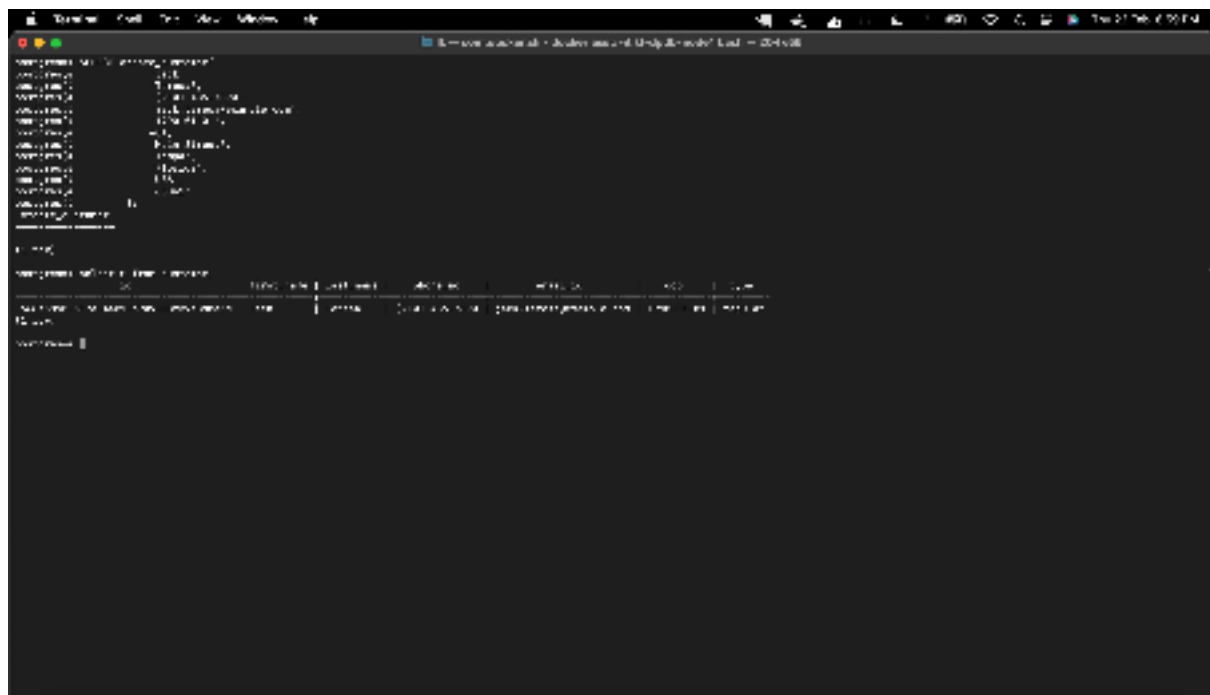


Fig-91: Customer sign-up details

Explanation: This is the sign-up details provided by the customer and triggering the create_customer function. Where will provide all the customer details such as first_name, last_name, phone_no, email_id, dob and type, and all the other details will update to address and customer_address tables.

PL/pgSQL code to place order by customer:

query:

```
CREATE OR REPLACE FUNCTION place_order_by_customer(
p_customer_id UUID,
p_product_id UUID,
p_product_quantity INTEGER,
p_delivery_partner_name VARCHAR(50),
p_delivery_partner_phone_no VARCHAR(20),
p_delivery_partner_email VARCHAR(255)
) RETURNS VOID AS $$
DECLARE
v_order_id UUID;
v_address_id UUID;
v_delivery_partner_id UUID;
BEGIN
-- Step 1: Check if the product and quantity are available
PERFORM 1 FROM products WHERE id = p_product_id AND quantity >= p_product_quantity;
IF NOT FOUND THEN
RAISE EXCEPTION 'Customer %, sorry, the requested product is not available. Please
try again later.', p_customer_id;
END IF;
-- Step 2: Insert data into orders
INSERT INTO orders (status, address_id, customer_id)
VALUES ('Processing', (SELECT address_id FROM customer_address WHERE customer_id =
p_customer_id AND default_address = true), p_customer_id)
RETURNING id INTO v_order_id;
-- Step 3: Insert data into transaction_summary
INSERT INTO transaction_summary (total_amount_paid, payment_type, order_id)
VALUES ((SELECT price * p_product_quantity FROM products WHERE id = p_product_id),
'Credit Card', v_order_id);
-- Step 4: Insert data into delivery_partner
INSERT INTO delivery_partner (name, phone_no, email, order_id)
VALUES (p_delivery_partner_name, p_delivery_partner_phone_no,
p_delivery_partner_email, v_order_id)
RETURNING id INTO v_delivery_partner_id;
-- Step 5: Insert data into customer_delivery_partner
INSERT INTO customer_delivery_partner (customer_id, delivery_partner_id)
VALUES (p_customer_id, v_delivery_partner_id);
-- Step 6: Insert data into orders_products
INSERT INTO orders_products (order_id, product_id, quantity)
VALUES (v_order_id, p_product_id, p_product_quantity);
-- Step 7: Update product quantity
UPDATE products SET quantity = quantity - p_product_quantity WHERE id =
p_product_id;
EXCEPTION
WHEN OTHERS THEN
-- Rollback the transaction in case of an exception
RAISE EXCEPTION 'Error in place_order_by_customer: %', SQLERRM;
END;
$$ LANGUAGE plpgsql;
```

```

-- Create a function to place an order by customer
CREATE OR REPLACE FUNCTION place_order(
    customer_id INT,
    product_id INT,
    product_quantity INT,
    delivery_partner VARCHAR(50))
RETURNS VOID AS
$$
BEGIN
    -- Check if the product is available
    IF (SELECT COUNT(*) FROM products WHERE product_id = product_id AND stock < product_quantity) > 0 THEN
        RAISE EXCEPTION 'Product is not available';
    END IF;

    -- Check if the delivery partner is valid
    IF (SELECT COUNT(*) FROM delivery_partners WHERE delivery_partner = delivery_partner) < 1 THEN
        RAISE EXCEPTION 'Invalid delivery partner';
    END IF;

    -- Insert the order into the orders table
    INSERT INTO orders (customer_id, product_id, product_quantity, delivery_partner, order_date)
    VALUES (customer_id, product_id, product_quantity, delivery_partner, NOW());

    -- Update the stock of the product
    UPDATE products SET stock = stock - product_quantity WHERE product_id = product_id;

    -- Commit the transaction
    COMMIT;
END;

```

Fig-92: Create function to place order by customer.

Explanation: This is the function, where the customer can place the order successfully, and even delivery partner is assigned to the order. This function updates the data in all the respective tables and achieves the data consistency. Here the input details are customer_id, product_id, product_quantity, and delivery_partner details. Using these details the order of the customer will be placed and delivery partner will get assigned to it. In case of any error occurs during the transactions, ROLLBACK is implemented which helps to maintain the data consistently.

Input data to place the order by customer:

query:

DO \$\$

BEGIN

PERFORM place_order_by_customer(

'668e5890-517a-4ae1-b909-03f39a3d8e6d', -- customer_id

'f5148b50-d8f4-45a7-a5aa-03cf6e047627', -- product_id

2, -- product_quantity

'John Doe', -- delivery_partner_name

'(973) 456-7890', -- delivery_partner_phone_no

'john.doe@example.com' -- delivery_partner_email

);

END \$\$;

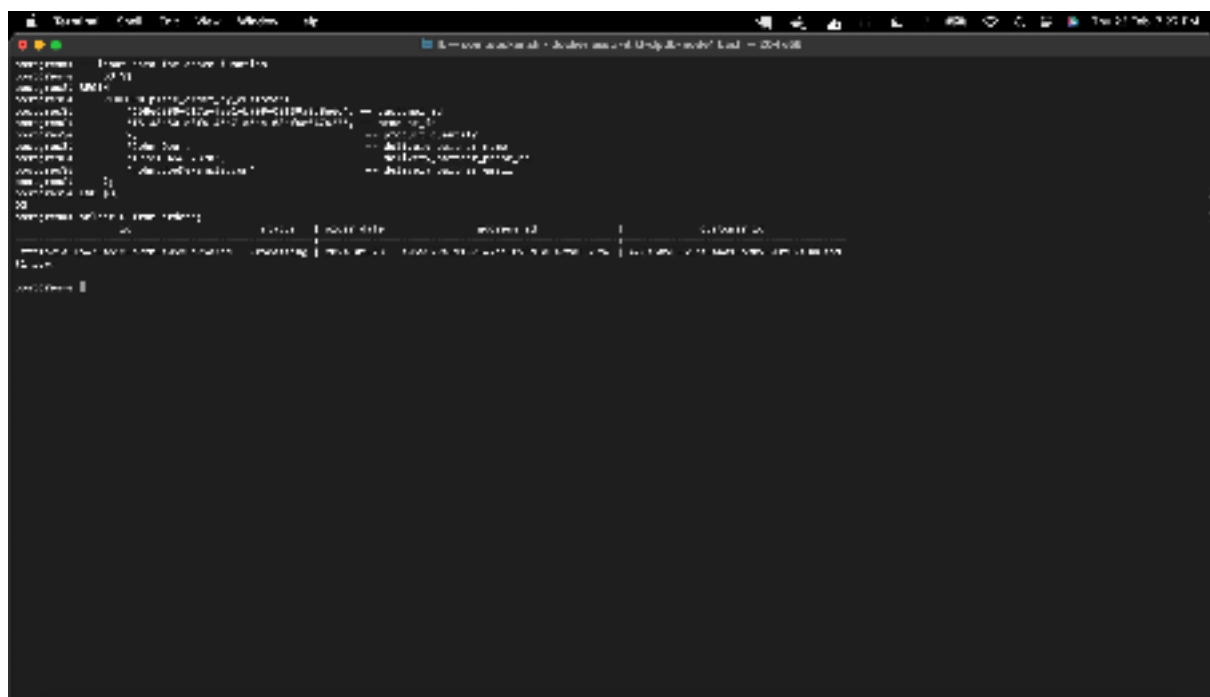


Fig-93: Input data to trigger place_order_by_customer and display the orders.

Explanation: This is the required input data to place the order by the customer. Few details will be read from the customer, and few details will get from our existing database. The delivery partner details are stores in database. Use those delivery partner details to assign an order to delivery partner.

use case: use multiple tables such as customer products, orders, transaction_summary and get the successfully ordered products paid more than their average total.

query:

```
SELECT prod.id AS product_id, prod.name AS product_name, ord.id AS order_id,
ord.status AS order_status,
ts.total_amount_paid, ts.payment_type, ts.date_of_payment
FROM products prod
JOIN orders_products ord_prod ON prod.id = ord_prod.product_id -- Corrected alias
from 'ord' to 'ord_prod'
JOIN orders ord ON ord_prod.order_id = ord.id -- Corrected alias from 'o' to 'ord'
JOIN transaction_summary ts ON ord.id = ts.order_id
JOIN product_category pc ON prod.product_category_id = pc.id
WHERE ord.customer_id = '668e5890-517a-4ae1-b909-03f39a3d8e6d' AND pc.section =
'Electronics' AND ts.total_amount_paid > (
    SELECT AVG(total_amount_paid)
    FROM transaction_summary
    WHERE order_id IN (
        SELECT id
        FROM orders
        WHERE customer_id = '668e5890-517a-4ae1-b909-03f39a3d8e6d')
);
```

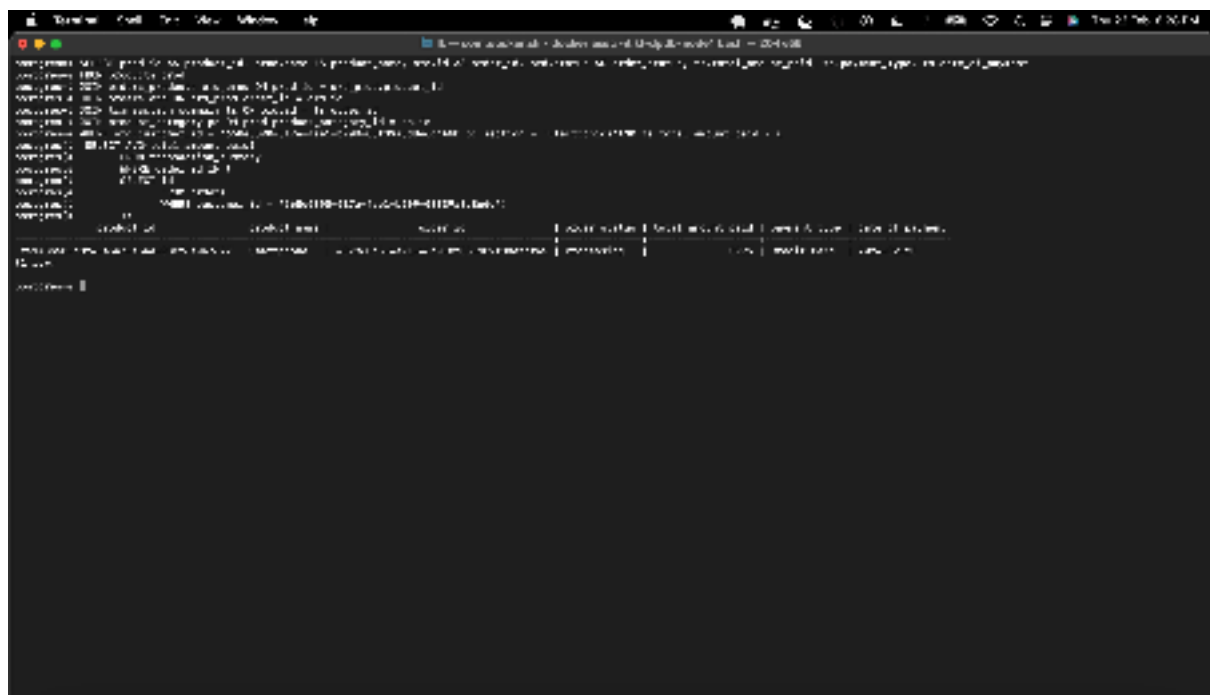


Fig-94: Get the products which has price more than the average of total products price ordered by the customer.

Explanation: customer placed multiple orders, but out of them the product which has the highest than the average price is “smartphone”. Using multiple tables such as products, orders, transaction summary, customers fetched the required details from the database. Without performing the parallel execution mechanism.

query:

EXPLAIN

```
SELECT prod.id AS product_id, prod.name AS product_name, ord.id AS order_id,  
ord.status AS order_status, ts.total_amount_paid, ts.payment_type,  
ts.date_of_payment  
FROM products prod  
JOIN orders_products ord_prod ON prod.id = ord_prod.product_id  
JOIN orders ord ON ord_prod.order_id = ord.id  
JOIN transaction_summary ts ON ord.id = ts.order_id  
JOIN product_category pc ON prod.product_category_id = pc.id  
WHERE ord.customer_id = '668e5890-517a-4ae1-b909-03f39a3d8e6d' AND pc.section =  
'Electronics' AND ts.total_amount_paid > (  
    SELECT AVG(total_amount_paid)  
    FROM transaction_summary  
    WHERE order_id IN (  
        SELECT id  
        FROM orders  
        WHERE customer_id = '668e5890-517a-4ae1-b909-03f39a3d8e6d')  
    );
```

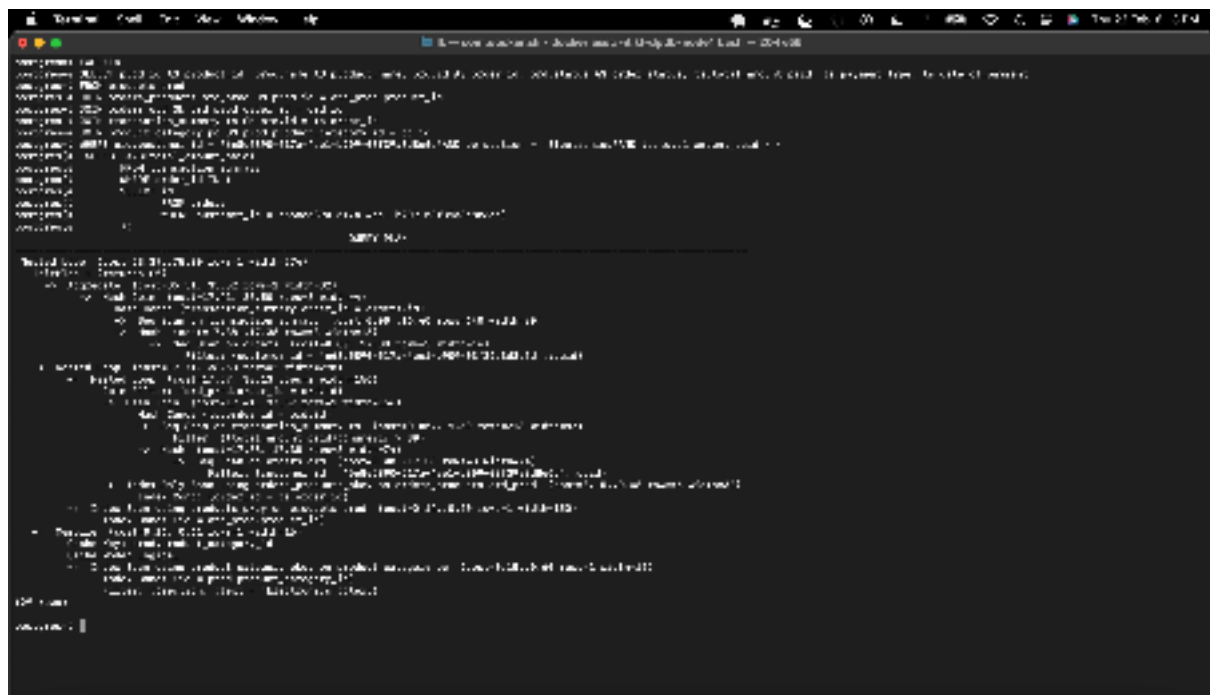


Fig-95: Analyzed the query performance.

Explanation: Without using parallel execution mechanism, analyzed the performance of the query. Observed that fetched the data using multiple nested loops, scanned sequentially to retrieve the required data. Applied filters, aggregation functions to retrieve the data from the database. Depending on the size of the tables, applied filters for better performance.

Validate Parallel execution of query:

query:

```
SET max_parallel_workers = 4;
```

EXPLAIN

```
SELECT prod.id AS product_id, prod.name AS product_name, ord.id AS order_id,  
ord.status AS order_status,  
ts.total_amount_paid, ts.payment_type, ts.date_of_payment  
FROM products prod  
JOIN orders_products ord_prod ON prod.id = ord_prod.product_id -- Corrected alias  
from 'ord' to 'ord_prod'  
JOIN orders ord ON ord_prod.order_id = ord.id -- Corrected alias from 'o' to 'ord'  
JOIN transaction_summary ts ON ord.id = ts.order_id  
JOIN product_category pc ON prod.product_category_id = pc.id  
WHERE ord.customer_id = '668e5890-517a-4ae1-b909-03f39a3d8e6d' AND pc.section =  
'Electronics' AND ts.total_amount_paid > (  
    SELECT AVG(total_amount_paid)  
    FROM transaction_summary  
    WHERE order_id IN (  
        SELECT id  
        FROM orders  
        WHERE customer_id = '668e5890-517a-4ae1-b909-03f39a3d8e6d'  
    )  
);
```

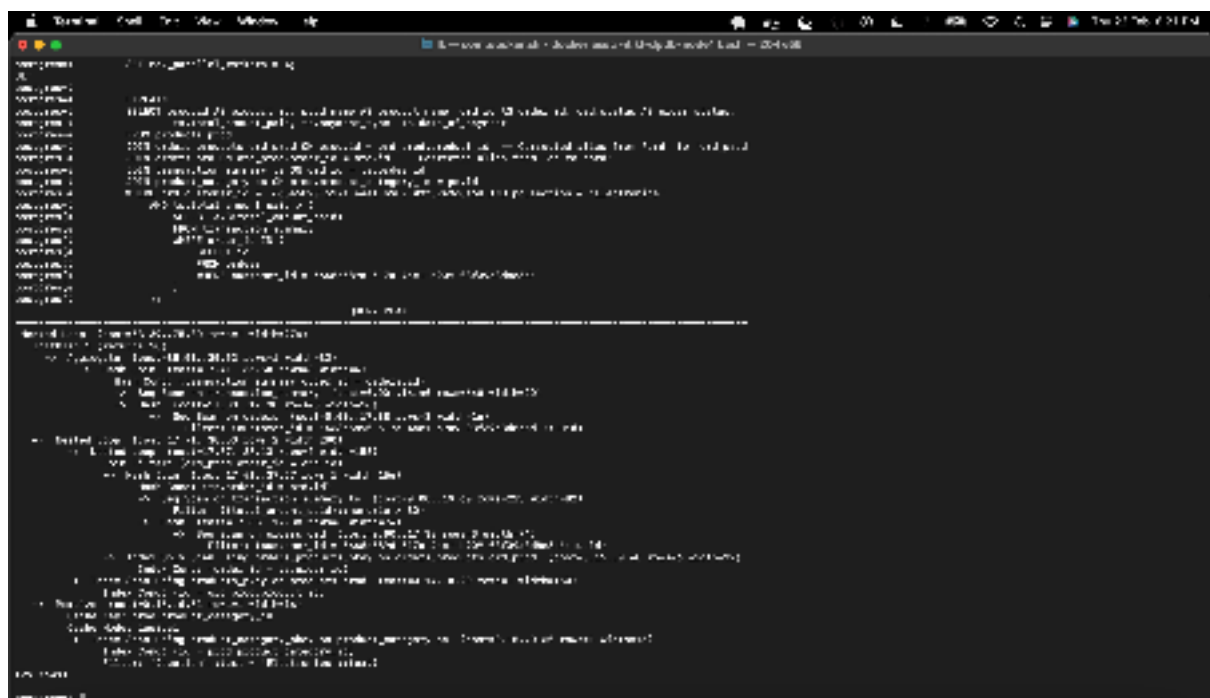


Fig-96: Analyzed and Executed query with 4 parallel worker nodes.

Explanation: Provide the max_parallel_worker nodes a 4. The postgres internally uses used the parallel worker nodes based on the requirement. Depending on the situation, the database uses the parallel worker nodes. The Execution of above query not effected by assigning the parallel worker nodes. Ideally the execution of parallel worker nodes depends

on various factors such as size of the data, query complexity, resources available and configurations. To execute this query the parallel query execution is not applicable.

Requirement-8:

use case 1: A customer can store multiple addresses and choose any one out it as default address. Customer can place order to any address as per choice. There also flexibility for the customer to update the address, incase customer adds the same existing address then database will through an error message "Address already exists. Please make use of it."

code:

DO \$\$

DECLARE

-- Input parameters

input_customer_id UUID := '668e5890-517a-4ae1-b909-03f39a3d8e6d';

input_flat_no INTEGER := 101;

input_street VARCHAR(50) := 'Maple lane, Apartment 3C';

input_city VARCHAR(50) := 'River City';

input_state VARCHAR(50) := 'Texas';

input_country VARCHAR(50) := 'USA';

input_zip_code VARCHAR(10) := '75001';

-- Variables for data

get_address_id UUID;

get_existing_address RECORD;

BEGIN

BEGIN

-- Insert data into the address table

INSERT INTO address (flat_no, street, city, state, country, zip_code)

VALUES (input_flat_no, input_street, input_city, input_state,

input_country, input_zip_code)

RETURNING id INTO get_address_id;

-- Insert data into the customer_address table

INSERT INTO customer_address (customer_id, address_id, default_address)

VALUES (input_customer_id, get_address_id, false);

-- Check if any other address exists for the customer

FOR get_existing_address IN

SELECT addr.*

FROM address addr

JOIN customer_address cust_addr ON addr.id = cust_addr.address_id

WHERE cust_addr.customer_id = input_customer_id AND addr.id <>

get_address_id

LOOP

-- Validate with the input

IF get_existing_address.flat_no = input_flat_no AND

get_existing_address.street = input_street AND

get_existing_address.city = input_city AND

get_existing_address.state = input_state AND

get_existing_address.country = input_country AND

```

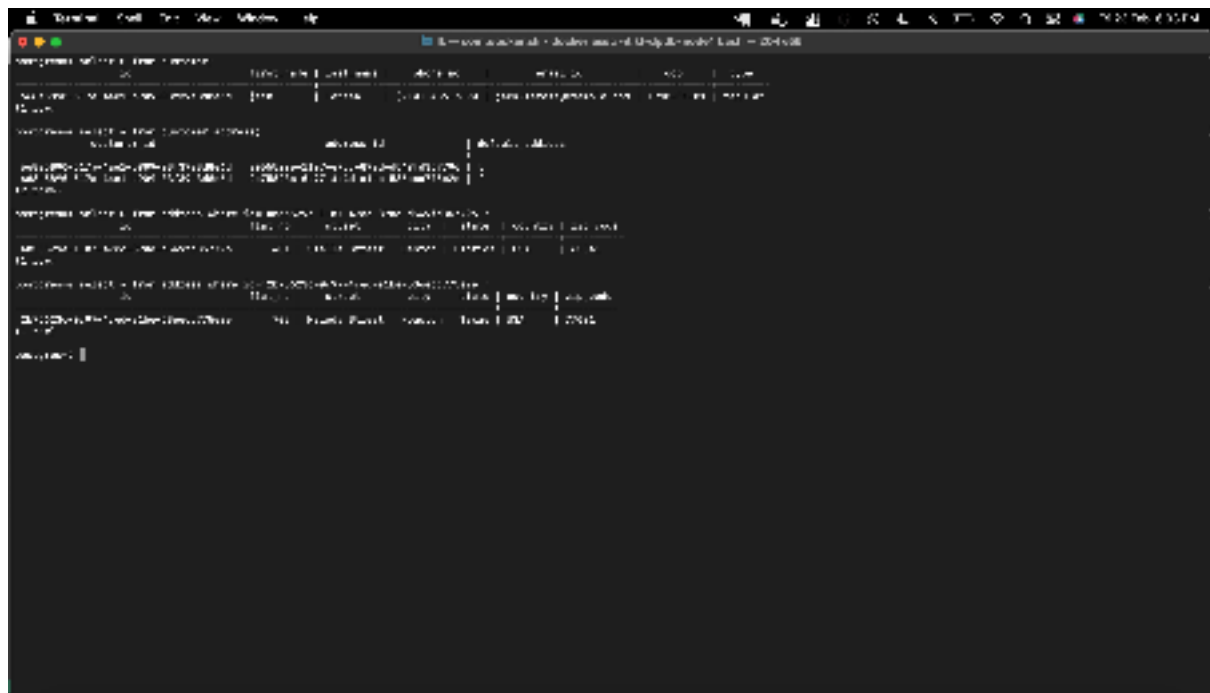
        get_existing_address.zip_code = input_zip_code THEN
            RAISE EXCEPTION 'Address already exists. Please make use of it.';
        END IF;
    END LOOP;

EXCEPTION
    WHEN OTHERS THEN
        -- An error occurred, roll back the transaction
        RAISE NOTICE 'Triggered Error: %', SQLERRM;
        ROLLBACK;
        RETURN;
    END;

    -- Everything went well, commit the transaction
    COMMIT;
    RAISE NOTICE 'Address saved successfully';

END $$;

```



The screenshot shows a SQL query result in a terminal window. The query is: `SELECT * FROM addresses WHERE customer_id = '668e5890-517a-4ae1-b909-03f39a3d8e6d';`. The result is a table with 7 columns: address_id, flat_no, street, city, state, country, and zip_code. There are 2 rows of data.

address_id	flat_no	street	city	state	country	zip_code
a055529e-1107-47cc-890d-d549fd357796	401	Stella Street	Denton	Florida	USA	65431
2b75623c-0c97-4c4d-a1be-536ecc775e2e	702	Melody Street	Houston	Texas	USA	77001

Fig-97: Display all the addresses for a customer.

Explanation: customer_id 668e5890-517a-4ae1-b909-03f39a3d8e6d has total 2 addresses with address id's a055529e-1107-47cc-890d-d549fd357796 and 2b75623c-0c97-4c4d-a1be-536ecc775e2e.

The address for a055529e-1107-47cc-890d-d549fd357796 has flat_no. '401', street 'Stella Street', city 'Denton', state 'Florida', country 'USA' and zip_code '65431'.

The address for 2b75623c-0c97-4c4d-a1be-536ecc775e2e has flat_no. '702', street 'Melody Street', city 'Houston', state 'Texas', country 'USA' and zip_code '77001'.

Now Let's try to insert a new address to 668e5890-517a-4ae1-b909-03f39a3d8e6d

```

1  # Import the necessary modules
2  import pandas as pd
3  import numpy as np
4  from sklearn.preprocessing import StandardScaler
5  from sklearn.model_selection import train_test_split
6  from sklearn.metrics import accuracy_score, confusion_matrix, classification_report
7  from sklearn.svm import SVC
8  from sklearn.metrics import roc_auc_score
9  from sklearn.metrics import roc_curve
10 from sklearn.metrics import auc
11
12 # Load the dataset
13 data = pd.read_csv('data.csv')
14
15 # Split the data into training and testing sets
16 X_train, X_test, y_train, y_test = train_test_split(data, data['target'],
17                                                    test_size=0.2,
18                                                    random_state=42)
19
20 # Scale the features
21 scaler = StandardScaler()
22 X_train = scaler.fit_transform(X_train)
23 X_test = scaler.transform(X_test)
24
25 # Train the SVM model
26 svm = SVC(kernel='rbf')
27 svm.fit(X_train, y_train)
28
29 # Evaluate the model
30 y_pred = svm.predict(X_test)
31 accuracy = accuracy_score(y_test, y_pred)
32 conf_matrix = confusion_matrix(y_test, y_pred)
33 class_report = classification_report(y_test, y_pred)
34
35 # Print the results
36 print('Accuracy: %.2f' % accuracy)
37 print('Confusion Matrix: \n', conf_matrix)
38 print('Classification Report: \n', class_report)
39
40 # ROC Curve
41 fpr, tpr, _ = roc_curve(y_test, svm.decision_function(X_test))
42 auc_roc = auc(fpr, tpr)
43
44 # Print the AUC
45 print('AUC: %.2f' % auc_roc)

```

Fig-98: Code to add the Customer Address.

[illegible]

Fig-99: Executed code to add the Customer Address.

Explanation: As per the use case, customer inserted the new address which is not there in existing system associated to the customer. Observed that the address saved successfully into the database system.

Let's see the saved address in the customer, customer_address and address tables.

address_id	customer_id	flat_no	street	city	state	country	zip_code
1	668e5890-517a-4ae1-b909-03f39a3d8e6d	101	Maple lane, Apartment 3C	River City	Texas	USA	75001
2	668e5890-517a-4ae1-b909-03f39a3d8e6d	101	Maple lane, Apartment 3C	River City	Texas	USA	75001

Fig-100: Displayed the latest stored address of the customer.

Explanation: once a new address is saved by the customer, we can observe that address is updated into all the required tables. For customer id '668e5890-517a-4ae1-b909-03f39a3d8e6d', a new address is added with flat_no – 101, street – Maple lane, Apartment 3C, city – River City, state – Texas, country- USA and zip-code 75001.

Now, Let's add the same address, for the '668e5890-517a-4ae1-b909-03f39a3d8e6d' and check the **error handling mechanism and rollback**.

code:

DO \$\$

DECLARE

-- Input parameters

input_customer_id UUID := '668e5890-517a-4ae1-b909-03f39a3d8e6d';

input_flat_no INTEGER := 101;

input_street VARCHAR(50) := 'Maple lane, Apartment 3C';

input_city VARCHAR(50) := 'River City';

input_state VARCHAR(50) := 'Texas';

input_country VARCHAR(50) := 'USA';

input_zip_code VARCHAR(10) := '75001';

-- Variables for data

get_address_id UUID;

get_existing_address RECORD;

BEGIN

-- Start the transaction

```

BEGIN
    -- Insert data into the address table
    INSERT INTO address (flat_no, street, city, state, country, zip_code)
    VALUES (input_flat_no, input_street, input_city, input_state,
input_country, input_zip_code)
    RETURNING id INTO get_address_id;
    RAISE NOTICE 'Inserted into address table for address_id: %',
get_address_id;

    -- Insert data into the customer_address table
    INSERT INTO customer_address (customer_id, address_id, default_address)
    VALUES (input_customer_id, get_address_id, false);
    RAISE NOTICE 'Inserted into customer_address table for customer_id: %,
address_id: %', input_customer_id, get_address_id;

    -- Check if any other address exists for the customer
    FOR get_existing_address IN
        SELECT addr.*
        FROM address addr
        JOIN customer_address cust_addr ON addr.id = cust_addr.address_id
        WHERE cust_addr.customer_id = input_customer_id AND addr.id <>
get_address_id
    LOOP
        -- Validate with the input
        IF get_existing_address.flat_no = input_flat_no AND
get_existing_address.street = input_street AND
get_existing_address.city = input_city AND
get_existing_address.state = input_state AND
get_existing_address.country = input_country AND
get_existing_address.zip_code = input_zip_code THEN
            RAISE EXCEPTION 'Address already exists. Please make use of it.';
        END IF;
    END LOOP;

    EXCEPTION
        WHEN OTHERS THEN
            -- An error occurred, roll back the transaction
            RAISE NOTICE 'Error occurred: %', SQLERRM;

            -- Log the rollback
            RAISE NOTICE 'Error occurred, rolling back the data';
            ROLLBACK;
            RETURN;
    END;

    -- Everything went well, commit the transaction
    COMMIT;
    RAISE NOTICE 'Address saved successfully';

END $$;

```

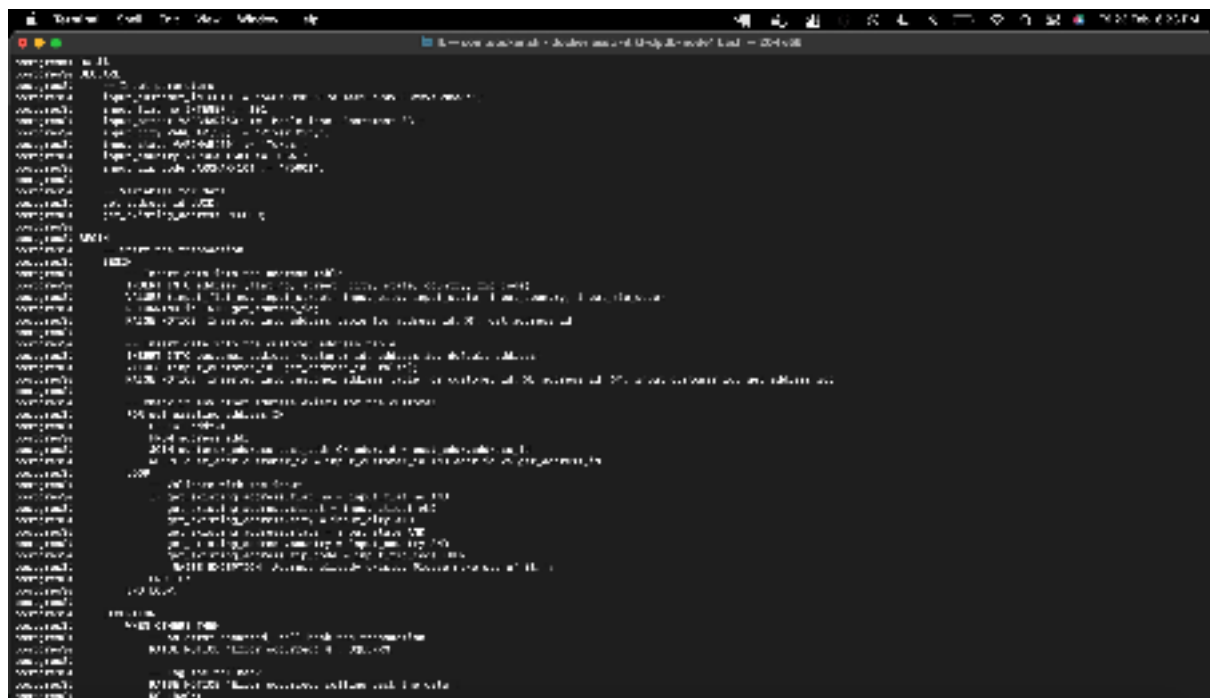



Fig-101: Stored the same address by the "668e5890-517a-4ae1-b909-03f39a3d8e6d".

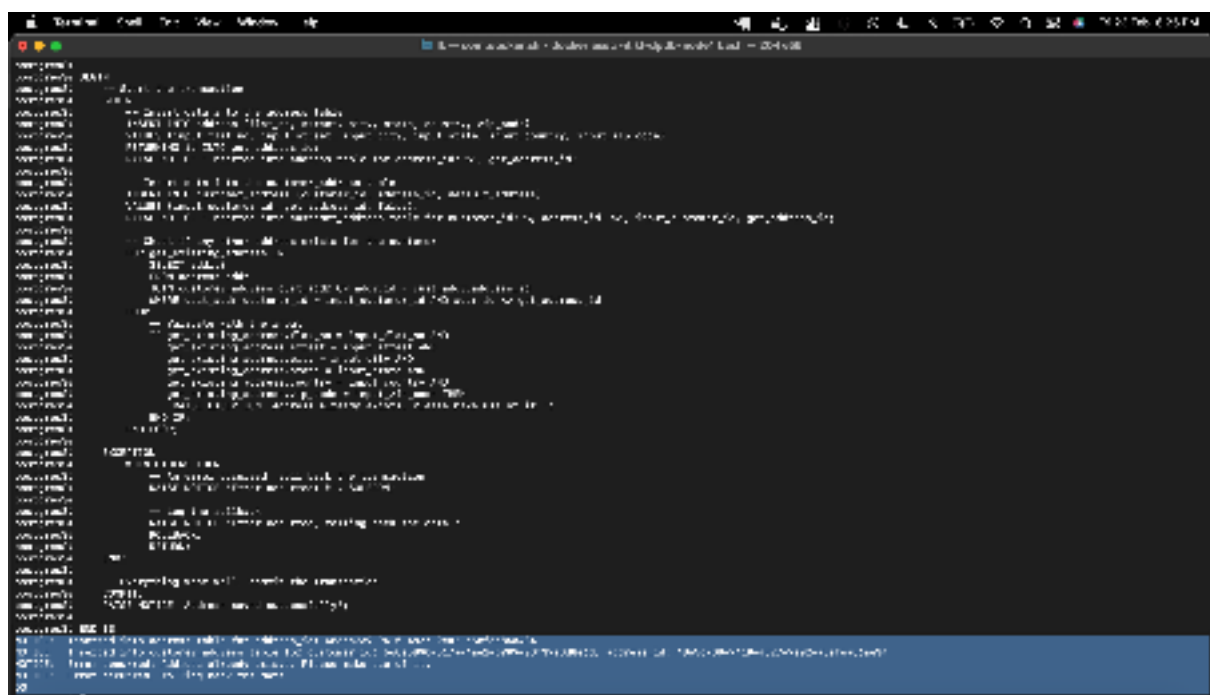


Fig-102: Display the error while saving the same address by the "668e5890-517a-4ae1-b909-03f39a3d8e6d".

Explanation: When user tried to add the same address again, while executing the function we can see that data inserted into address table, customer_address table and upon validating the inserted data, observed that the customer is trying to add the duplicate address into the database system So, as error occurred due to duplicate address, the data got roll backed from the database.

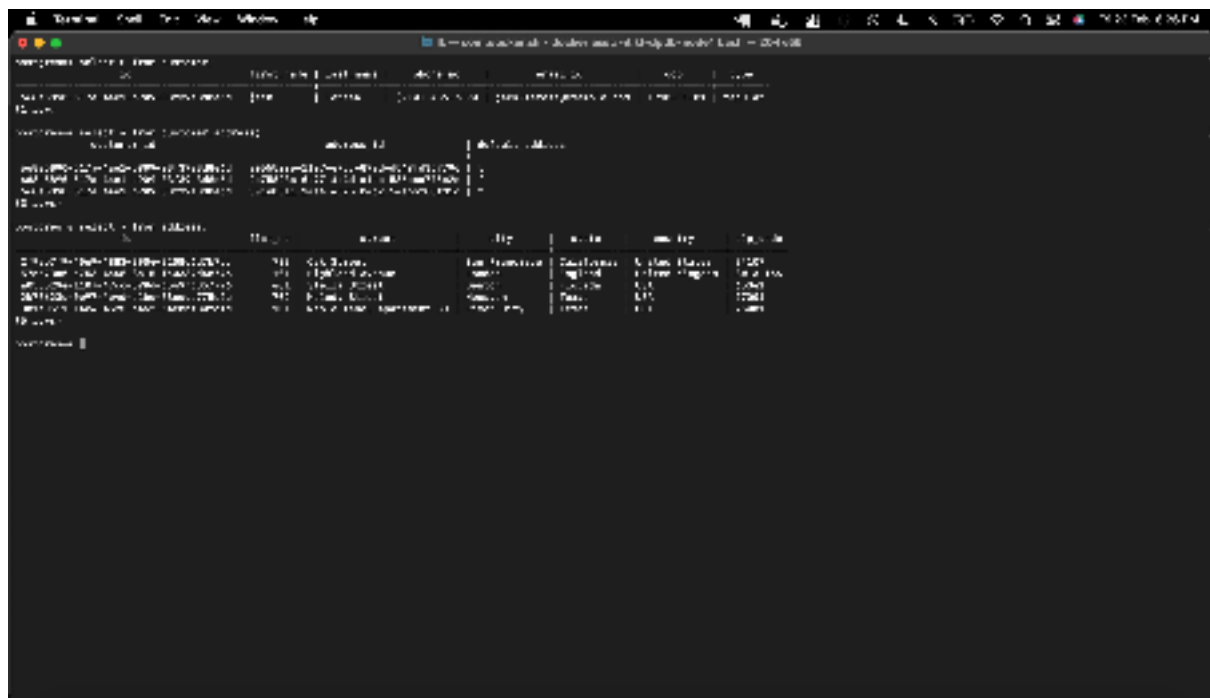


Fig-103: Insertions rolled back, due to error occurred while execution.

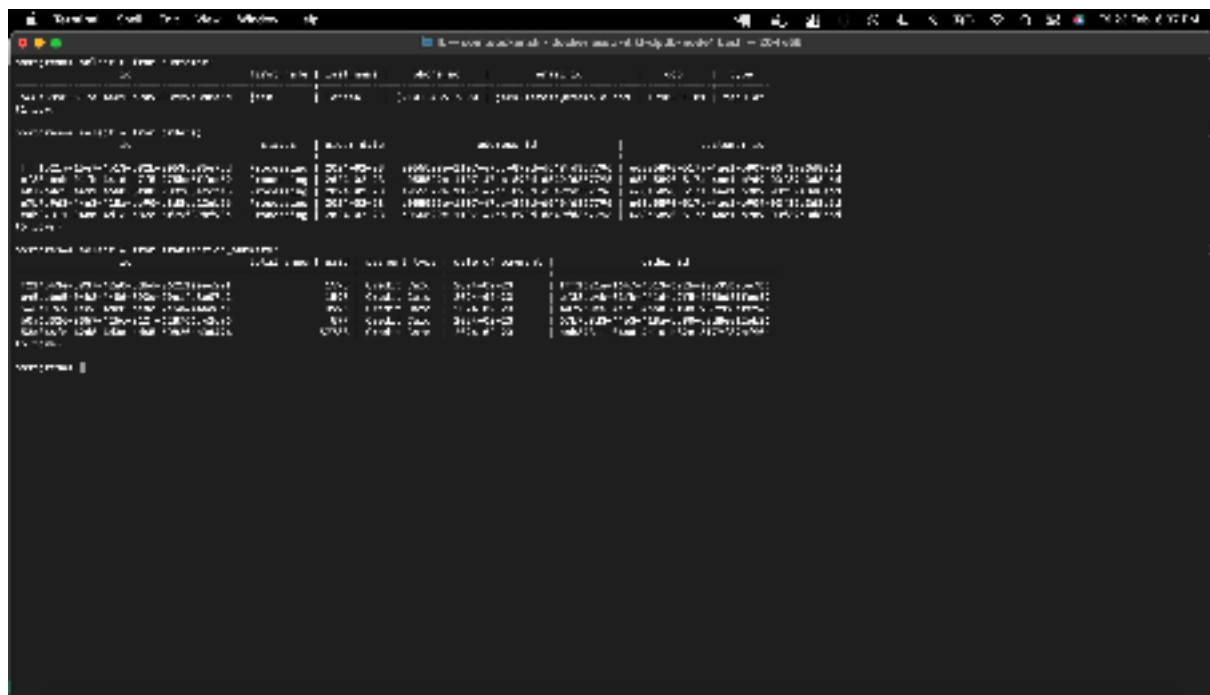
Explanation: In the above executed code, as error occurred during the transactions and rollback is performed, the inserted data is also rolled back in respective tables(inserted data before the error occur).

Usecase-2: A customer may use a credit card to make purchases up to Rs. 1,000,000 every month. He/She is unable to place an order using a credit card after their monthly credit card spending exceeds Rs. 1,00,000. He/She ought to employ different ways to pay.

Let's see the customer spendings on credit card:

query:

```
select * from customer;
select * from orders;
select * from transaction_summary;
```



The screenshot shows a database query result in a terminal window. The query is: `select * from orders; select * from transaction_summary;`. The results are displayed in two tables.

id	customer_id	product_id	quantity	price	total_price	status	created_at	updated_at
1	668e5890-517a-4ae1-b909-03f39a3d8e6d	1	1	15000	15000	Completed	2021-03-11	2021-03-11
2	668e5890-517a-4ae1-b909-03f39a3d8e6d	2	1	15000	15000	Completed	2021-03-11	2021-03-11
3	668e5890-517a-4ae1-b909-03f39a3d8e6d	3	1	15000	15000	Completed	2021-03-11	2021-03-11
4	668e5890-517a-4ae1-b909-03f39a3d8e6d	4	1	15000	15000	Completed	2021-03-11	2021-03-11

id	customer_id	product_id	quantity	price	total_price	status	created_at	updated_at
1	668e5890-517a-4ae1-b909-03f39a3d8e6d	1	1	15000	15000	Completed	2021-03-11	2021-03-11
2	668e5890-517a-4ae1-b909-03f39a3d8e6d	2	1	15000	15000	Completed	2021-03-11	2021-03-11
3	668e5890-517a-4ae1-b909-03f39a3d8e6d	3	1	15000	15000	Completed	2021-03-11	2021-03-11
4	668e5890-517a-4ae1-b909-03f39a3d8e6d	4	1	15000	15000	Completed	2021-03-11	2021-03-11

Fig-104: Customer '668e5890-517a-4ae1-b909-03f39a3d8e6d' orders and transaction_summary.

Explanation: From the above screenshot customer has "668e5890-517a-4ae1-b909-03f39a3d8e6d" has 4 orders with the total amount of 65,578 using credit card.

Now Let's make the one more transaction for same user which exceeds 1,00,000.

query:

```
select * from products;
```



id	name	description	price	quantity	category	brand	manufacturer
8893a2e3-e040-4427-9e60-b89d8e0b4482	Product 1	Product 1 description	50000	10	Electronics	Brand A	Manufacturer A
...

Explanation: In the above screen shot we can observe that that product id “8893a2e3-e040-4427-9e60-b89d8e0b4482” has price 50000

code:

DO \$\$

DECLARE

-- Input parameters

input_payment_type VARCHAR(20) := 'Credit Card';

input_customer_id UUID := '668e5890-517a-4ae1-b909-03f39a3d8e6d';

input_product_id UUID := '8893a2e3-e040-4427-9e60-b89d8e0b4482';

input_product_quantity INTEGER := 1;

input_delivery_partner_name VARCHAR(50) := 'Fedx';

input_delivery_partner_phone_no VARCHAR(20) := '(980) 426-7190';

input_delivery_partner_email VARCHAR(255) := 'Fedx@example.com';

-- Variables for data

v_total_amount_paid INTEGER;

v_credit_limit INTEGER := 100000;

v_remaining_credit INTEGER;

v_order_amount INTEGER;

v_order_id UUID;

v_delivery_partner_id UUID;

BEGIN

-- Start the transaction

BEGIN

-- Step 1: Check if the product and quantity are available

PERFORM 1 FROM products WHERE id = input_product_id AND quantity >= input_product_quantity;

IF NOT FOUND THEN

RAISE EXCEPTION 'Sorry, the requested product is not available. Please try again later.';

END IF;

-- Step 2: Insert data into orders

INSERT INTO orders (status, order_date, address_id, customer_id)

```

VALUES ('Processing', CURRENT_DATE,
        (SELECT address_id FROM customer_address WHERE customer_id =
input_customer_id AND default_address = true),
        input_customer_id)
RETURNING id INTO v_order_id;
RAISE NOTICE 'Inserted into orders table for order_id: %', v_order_id;

-- Step 3: Calculate order amount (replace this with your own logic)
v_order_amount := input_product_quantity * (SELECT price FROM products
WHERE id = input_product_id);

-- Step 4: Check if the customer has enough credit limit
SELECT COALESCE(SUM(total_amount_paid), 0) INTO v_total_amount_paid
FROM transaction_summary
WHERE order_id IN (SELECT id FROM orders WHERE customer_id =
input_customer_id) AND payment_type = input_payment_type;

v_remaining_credit := v_credit_limit - v_total_amount_paid -
v_order_amount;

IF v_remaining_credit >= 0 THEN
    -- Update transaction_summary only if all conditions passed
    INSERT INTO transaction_summary (total_amount_paid, payment_type,
date_of_payment, order_id)
VALUES ((SELECT price * input_product_quantity FROM products WHERE id =
input_product_id), input_payment_type, CURRENT_DATE, v_order_id);
    RAISE NOTICE 'Inserted into transaction_summary table for customer_id:
%', input_customer_id;
ELSE
    RAISE EXCEPTION 'Customer % has exceeded the credit limit. Please use
another payment method.', input_customer_id;
END IF;

-- Step 5: Insert data into delivery_partner
INSERT INTO delivery_partner (name, phone_no, email, order_id)
VALUES (input_delivery_partner_name, input_delivery_partner_phone_no,
input_delivery_partner_email, v_order_id)
RETURNING id INTO v_delivery_partner_id;
RAISE NOTICE 'Inserted into delivery_partner table for delivery_partner_id:
%', v_delivery_partner_id;

-- Step 6: Insert data into customer_delivery_partner
INSERT INTO customer_delivery_partner (customer_id, delivery_partner_id)
VALUES (input_customer_id, v_delivery_partner_id);
RAISE NOTICE 'Inserted into customer_delivery_partner table';

-- Step 7: Insert data into orders_products
INSERT INTO orders_products (order_id, product_id, quantity)
VALUES (v_order_id, input_product_id, input_product_quantity);
RAISE NOTICE 'Inserted into orders_products table';

```

```

-- Step 8: Update product quantity
UPDATE products SET quantity = quantity - input_product_quantity WHERE id =
input_product_id;
RAISE NOTICE 'Updated product quantity for product_id: %',
input_product_id;

EXCEPTION
WHEN OTHERS THEN
    -- An error occurred, roll back the transaction
    RAISE NOTICE 'Error occurred: %', SQLERRM;

    -- Log the rollback
    RAISE NOTICE 'Error occurred, rolling back the data';

    ROLLBACK;
    RETURN;

END;

-- Everything went well, commit the transaction
COMMIT;
RAISE NOTICE 'Order placed successfully';

END $$;

```

```

SQL> SET SERVEROUTPUT ON;
SQL> EXECUTE IMMEDIATE '
    CREATE OR REPLACE FUNCTION place_order_credit_card(
        p_product_id IN NUMBER,
        p_quantity IN NUMBER,
        p_customer_id IN NUMBER,
        p_credit_card IN VARCHAR2)
    RETURN VARCHAR2 IS
        v_result VARCHAR2(100);
    BEGIN
        -- Check if product exists
        IF (SELECT COUNT(*) FROM products WHERE id = p_product_id) = 0 THEN
            v_result := 'Product not found';
            RETURN v_result;
        END IF;

        -- Check if customer exists
        IF (SELECT COUNT(*) FROM customers WHERE id = p_customer_id) = 0 THEN
            v_result := 'Customer not found';
            RETURN v_result;
        END IF;

        -- Check if credit card is valid
        IF (SELECT COUNT(*) FROM credit_cards WHERE customer_id = p_customer_id AND card_number = p_credit_card) = 0 THEN
            v_result := 'Credit card not found';
            RETURN v_result;
        END IF;

        -- Check if product is available
        IF (SELECT quantity FROM products WHERE id = p_product_id) < p_quantity THEN
            v_result := 'Product not available';
            RETURN v_result;
        END IF;

        -- Place order
        INSERT INTO orders (product_id, quantity, customer_id, credit_card)
        VALUES (p_product_id, p_quantity, p_customer_id, p_credit_card);

        -- Update product quantity
        UPDATE products SET quantity = quantity - p_quantity WHERE id = p_product_id;

        -- Commit
        COMMIT;

        v_result := 'Order placed successfully';
        RETURN v_result;
    END;';

SQL> EXECUTE IMMEDIATE '
    DECLARE
        v_result VARCHAR2(100);
    BEGIN
        v_result := place_order_credit_card(1, 1, 1, '1234567890123456');
        DBMS_OUTPUT.PUT_LINE(v_result);
    END;';

SQL> /
Order placed successfully

```

Fig-105: Execute query to place order using credit card.


```

BEGIN
    -- Step 1: Check if the product and quantity are available
    PERFORM 1 FROM products WHERE id = input_product_id AND quantity >=
input_product_quantity;
    IF NOT FOUND THEN
        RAISE EXCEPTION 'Sorry, the requested product is not available. Please
try again later.';
    END IF;

    -- Step 2: Insert data into orders
    INSERT INTO orders (status, order_date, address_id, customer_id)
    VALUES ('Processing', CURRENT_DATE,
        (SELECT address_id FROM customer_address WHERE customer_id =
input_customer_id AND default_address = true),
        input_customer_id)
    RETURNING id INTO v_order_id;
    RAISE NOTICE 'Inserted into orders table for order_id: %', v_order_id;

    -- Step 3: Calculate order amount (replace this with your own logic)
    v_order_amount := input_product_quantity * (SELECT price FROM products
WHERE id = input_product_id);

    -- Step 4: Check if the customer has enough credit limit
    SELECT COALESCE(SUM(total_amount_paid), 0) INTO v_total_amount_paid
    FROM transaction_summary
    WHERE order_id IN (SELECT id FROM orders WHERE customer_id =
input_customer_id) AND payment_type = input_payment_type;

    v_remaining_credit := v_credit_limit - v_total_amount_paid -
v_order_amount;

    IF v_remaining_credit >= 0 THEN
        -- Update transaction_summary only if all conditions passed
        INSERT INTO transaction_summary (total_amount_paid, payment_type,
date_of_payment, order_id)
        VALUES ((SELECT price * input_product_quantity FROM products WHERE id =
input_product_id), input_payment_type, CURRENT_DATE, v_order_id);
        RAISE NOTICE 'Inserted into transaction_summary table for customer_id:
%', input_customer_id;
    ELSE
        RAISE EXCEPTION 'Customer % has exceeded the credit limit. Please use
another payment method.', input_customer_id;
    END IF;

    -- Step 5: Insert data into delivery_partner
    INSERT INTO delivery_partner (name, phone_no, email, order_id)
    VALUES (input_delivery_partner_name, input_delivery_partner_phone_no,
input_delivery_partner_email, v_order_id)
    RETURNING id INTO v_delivery_partner_id;
    RAISE NOTICE 'Inserted into delivery_partner table for delivery_partner_id:
%', v_delivery_partner_id;

```



```

-- Step 6: Insert data into customer_delivery_partner
INSERT INTO customer_delivery_partner (customer_id, delivery_partner_id)
VALUES (input_customer_id, v_delivery_partner_id);
RAISE NOTICE 'Inserted into customer_delivery_partner table';

-- Step 7: Insert data into orders_products
INSERT INTO orders_products (order_id, product_id, quantity)
VALUES (v_order_id, input_product_id, input_product_quantity);
RAISE NOTICE 'Inserted into orders_products table';

-- Step 8: Update product quantity
UPDATE products SET quantity = quantity - input_product_quantity WHERE id =
input_product_id;
RAISE NOTICE 'Updated product quantity for product_id: %',
input_product_id;

EXCEPTION
  WHEN OTHERS THEN
    -- An error occurred, roll back the transaction
    RAISE NOTICE 'Error occurred: %', SQLERRM;


    -- Log the rollback
    RAISE NOTICE 'Error occurred, rolling back the data';

    ROLLBACK;
    RETURN;
END;

-- Everything went well, commit the transaction
COMMIT;
RAISE NOTICE 'Order placed successfully';

END $$;

```



Explanation: Order placed successfully using the debit card. Observed that data is inserted into all the corresponding tables such as orders, transaction_summary, delivery_partner, customer_delivery_partner, orders_products tables and update the products quantity in the products table.

